

Discriminative Training part 2

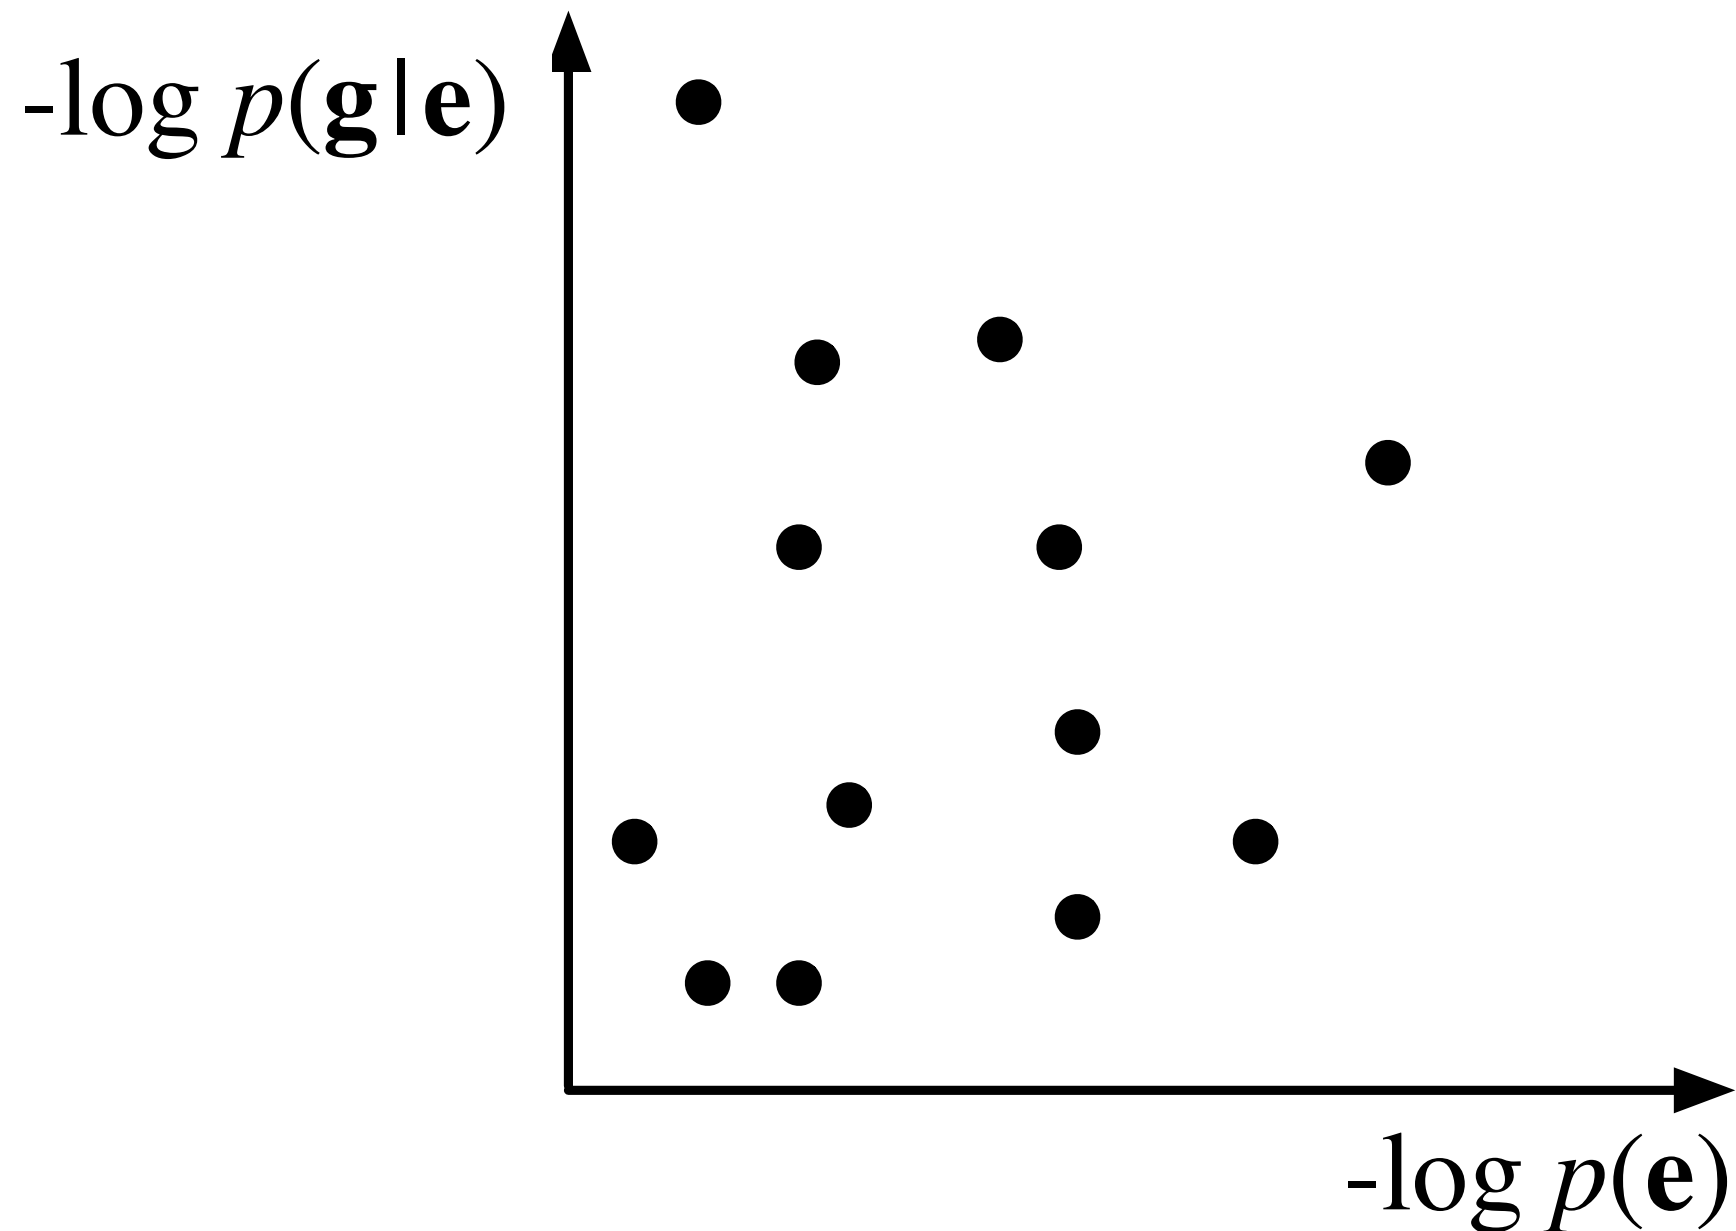
Machine Translation Lecture 12

Instructor: Chris Callison-Burch
TAs: Mitchell Stern, Justin Chiu

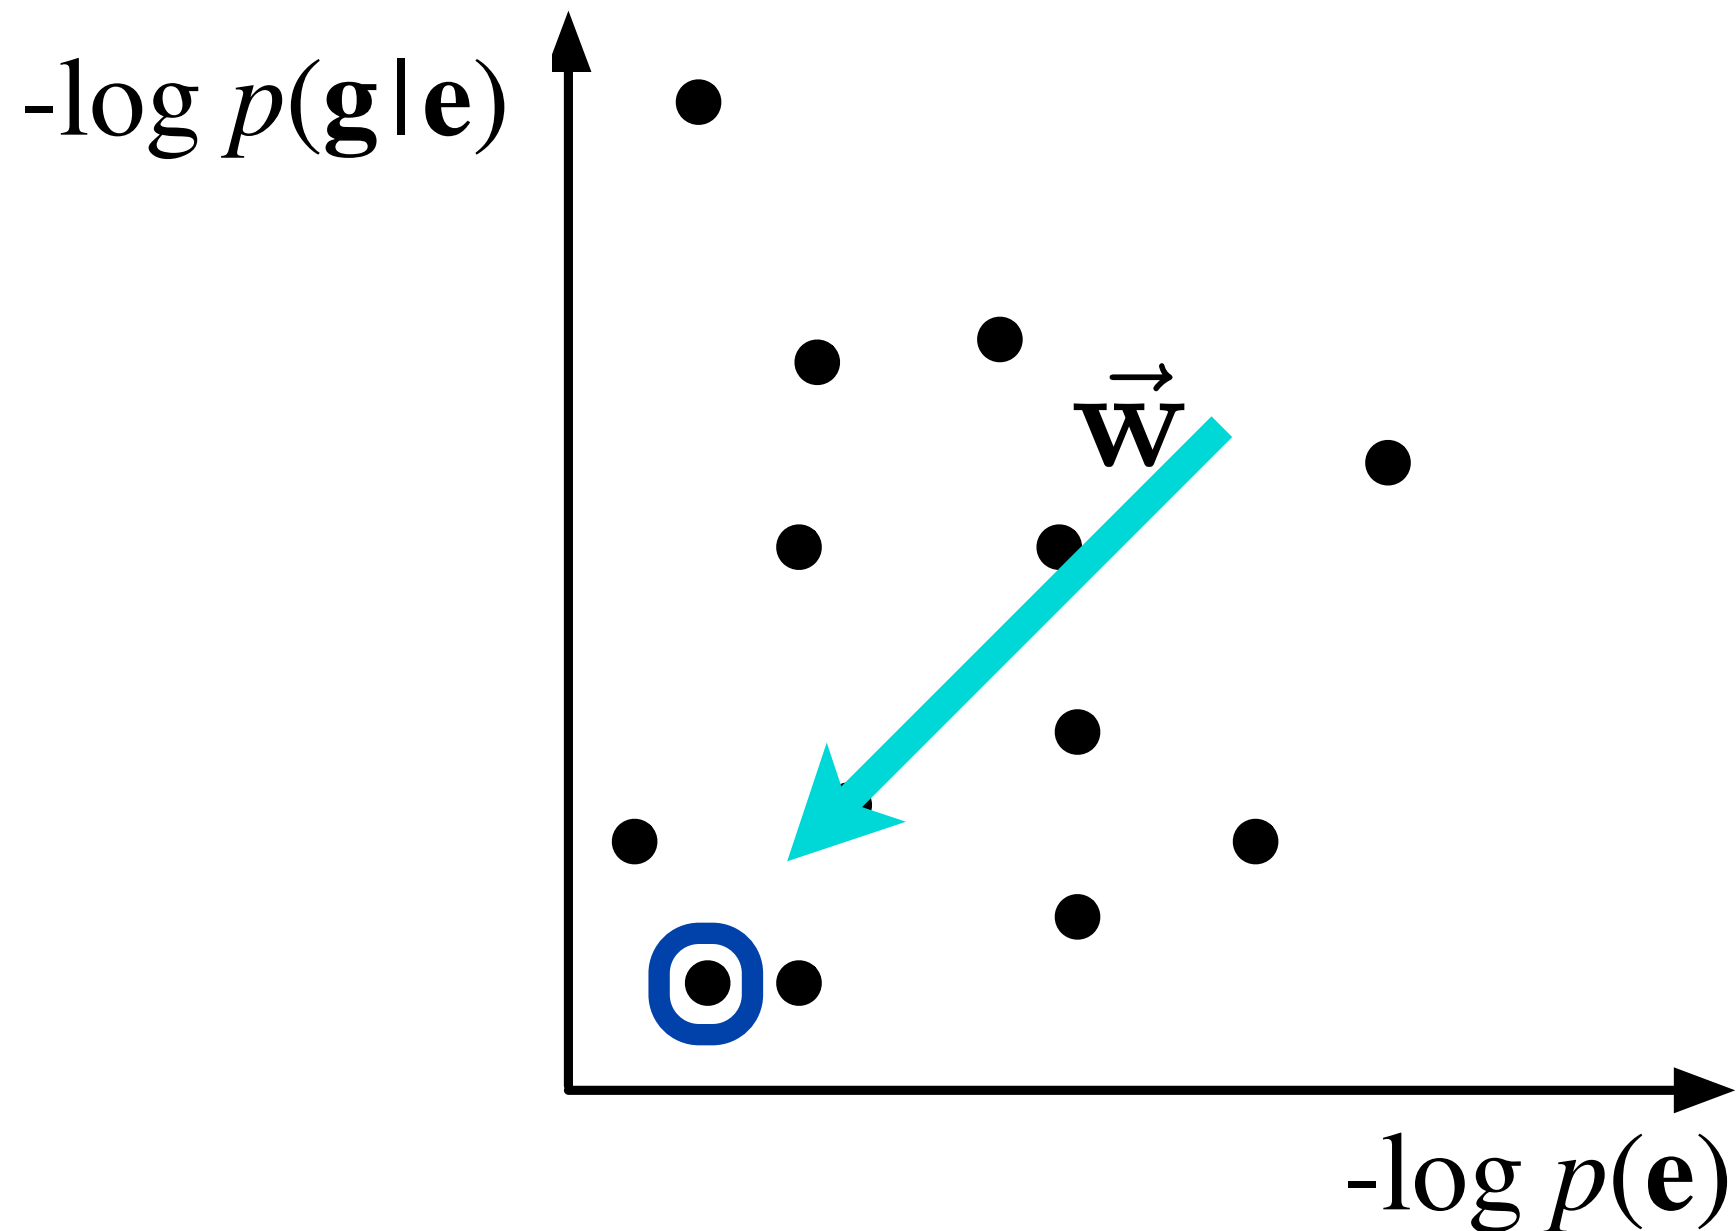
Website: mt-class.org/penn



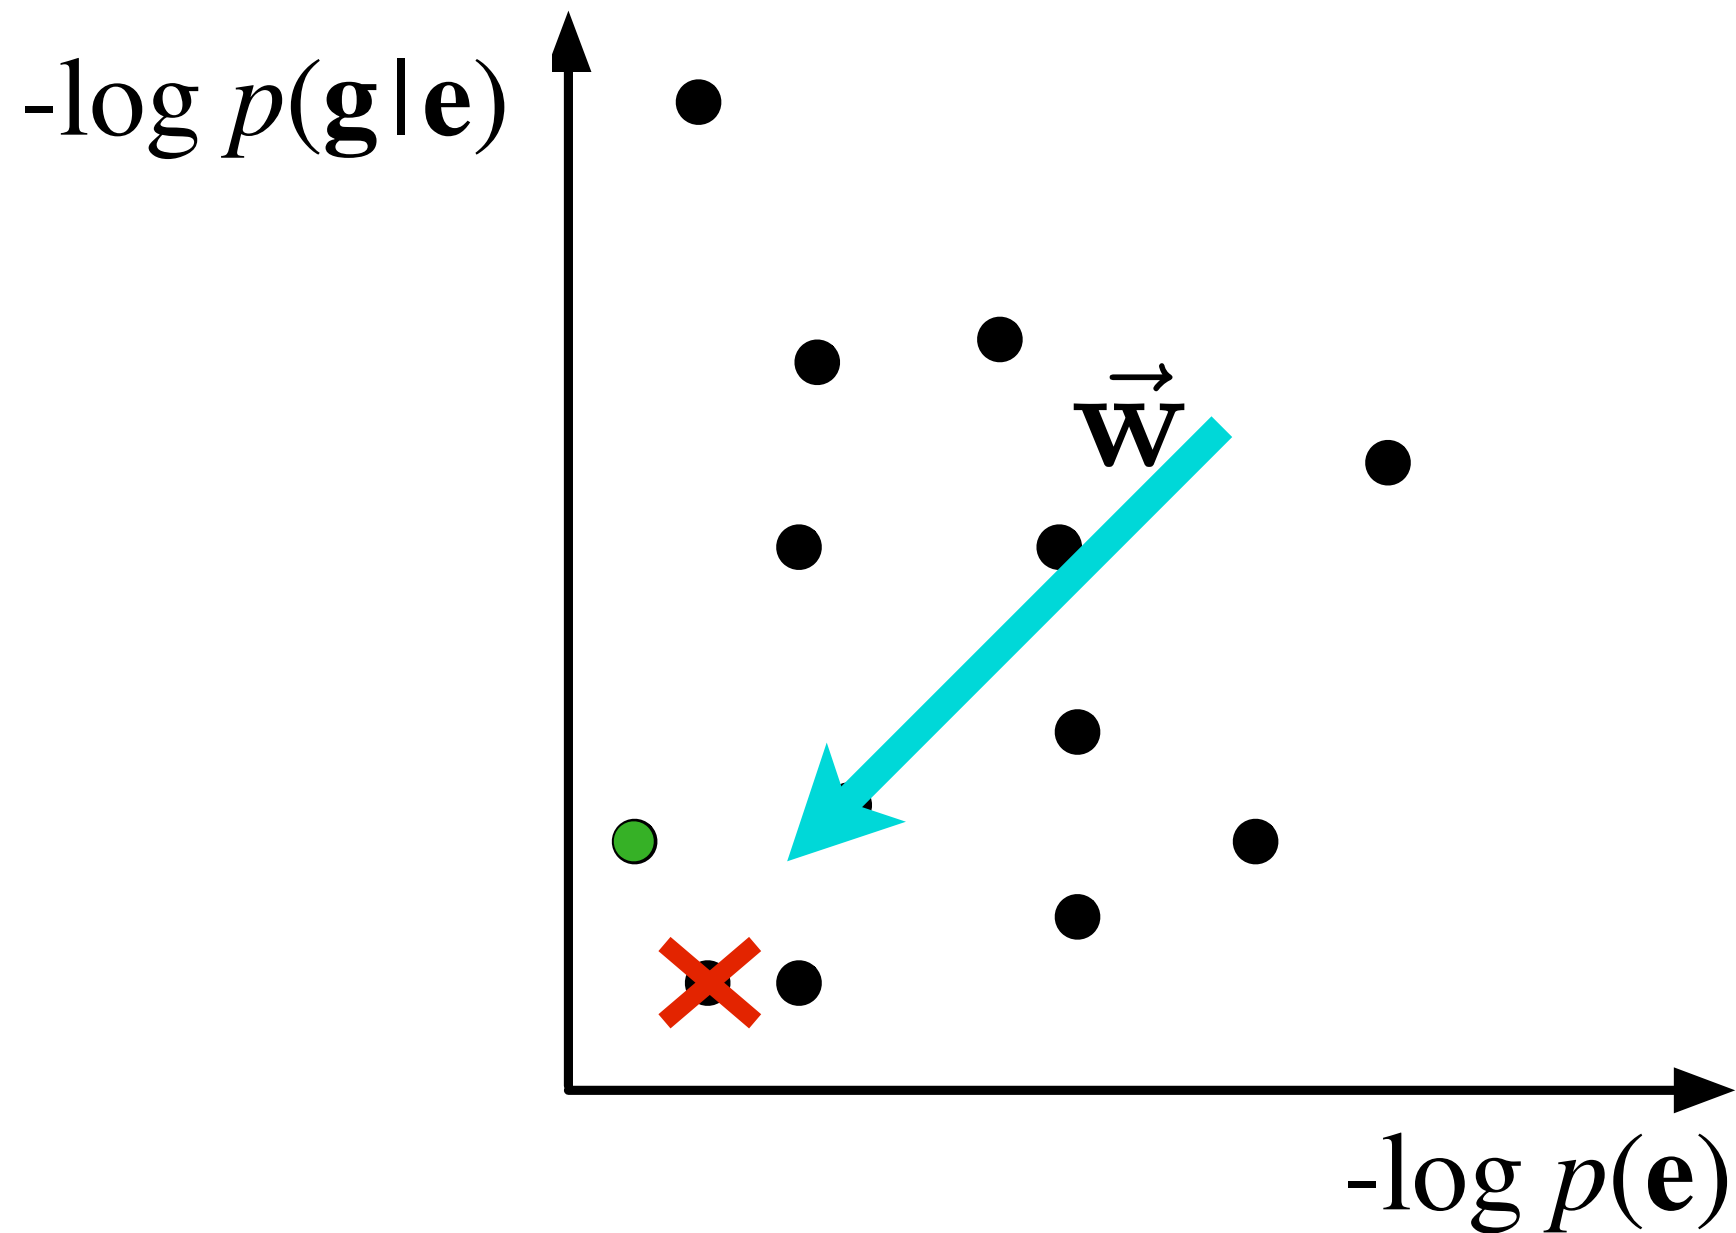
The Noisy Channel



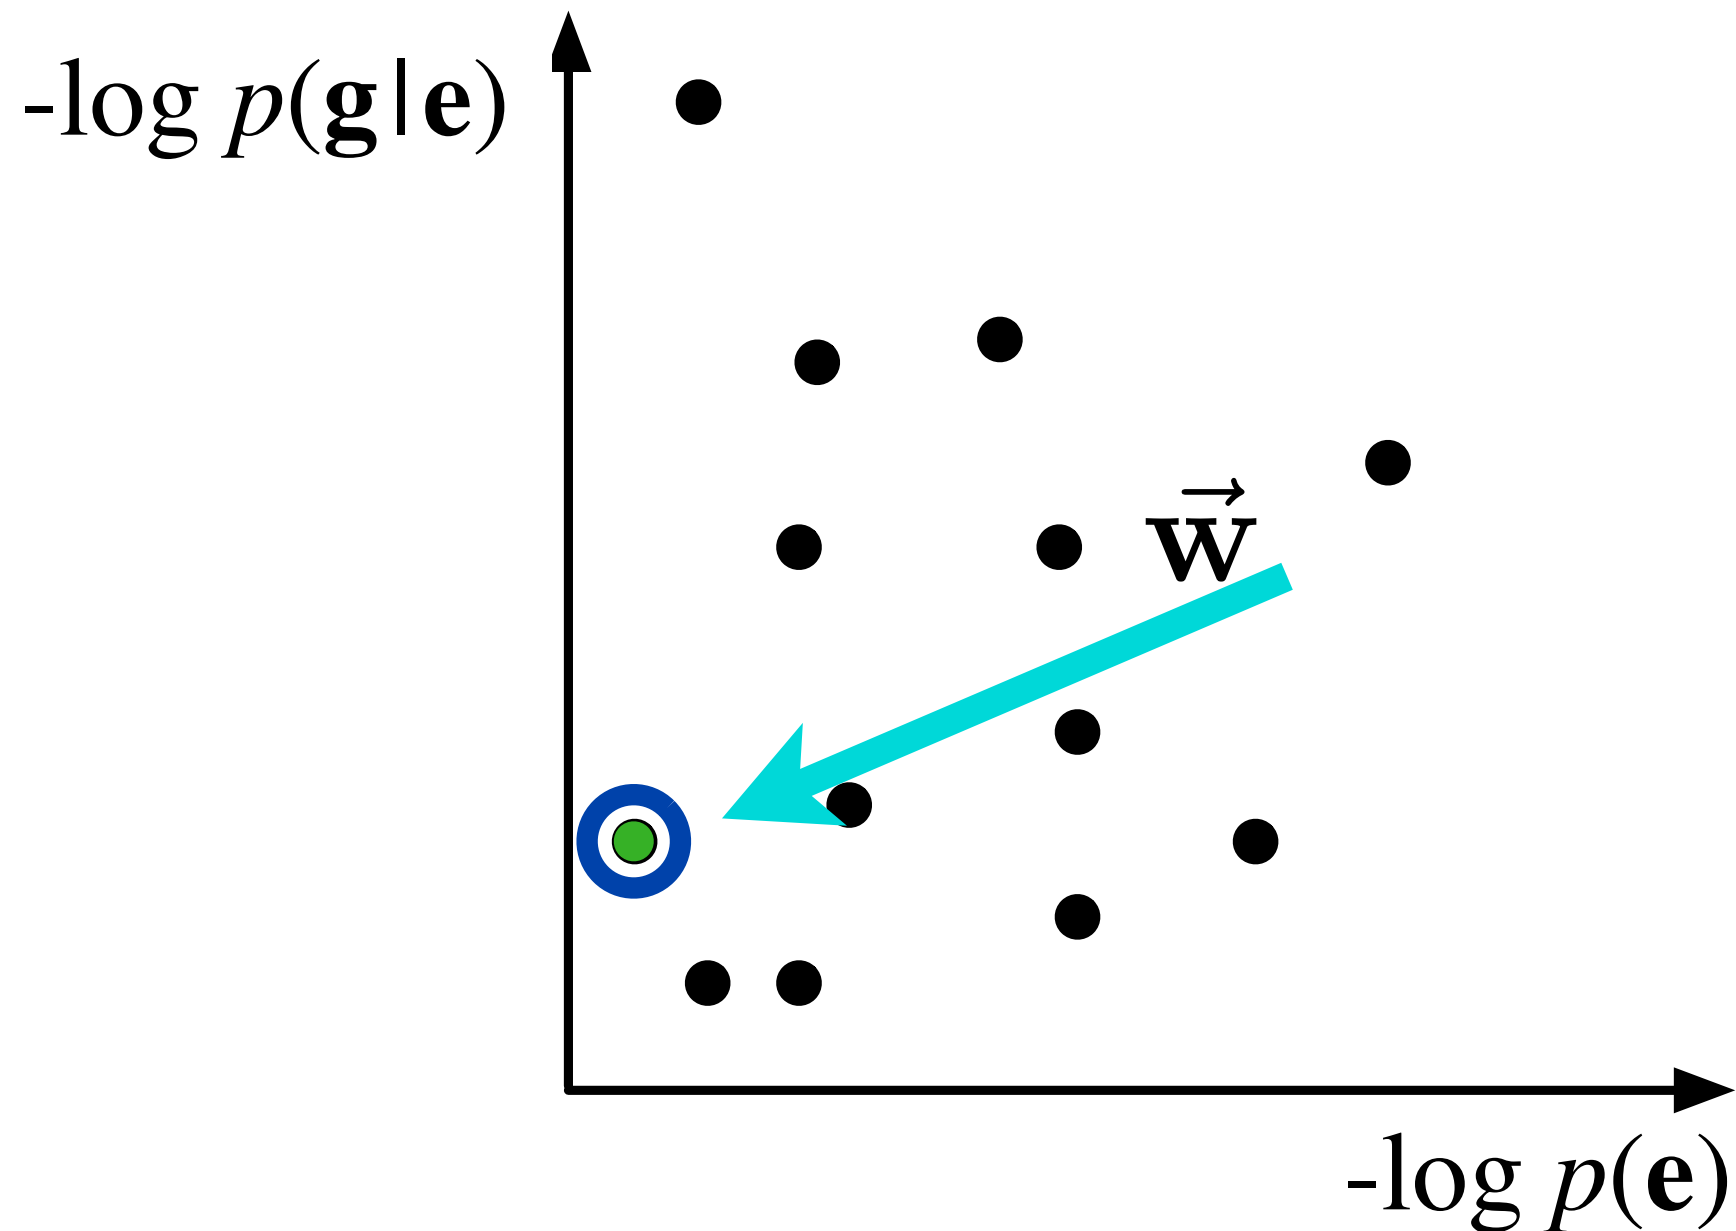
As a Linear Model



As a Linear Model



As a Linear Model

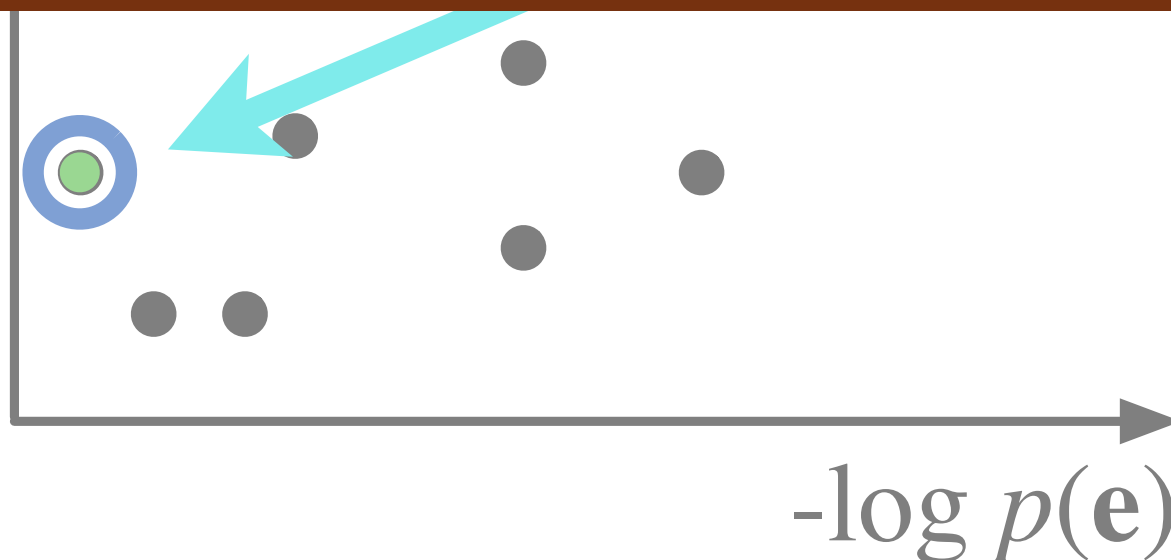


As a Linear Model

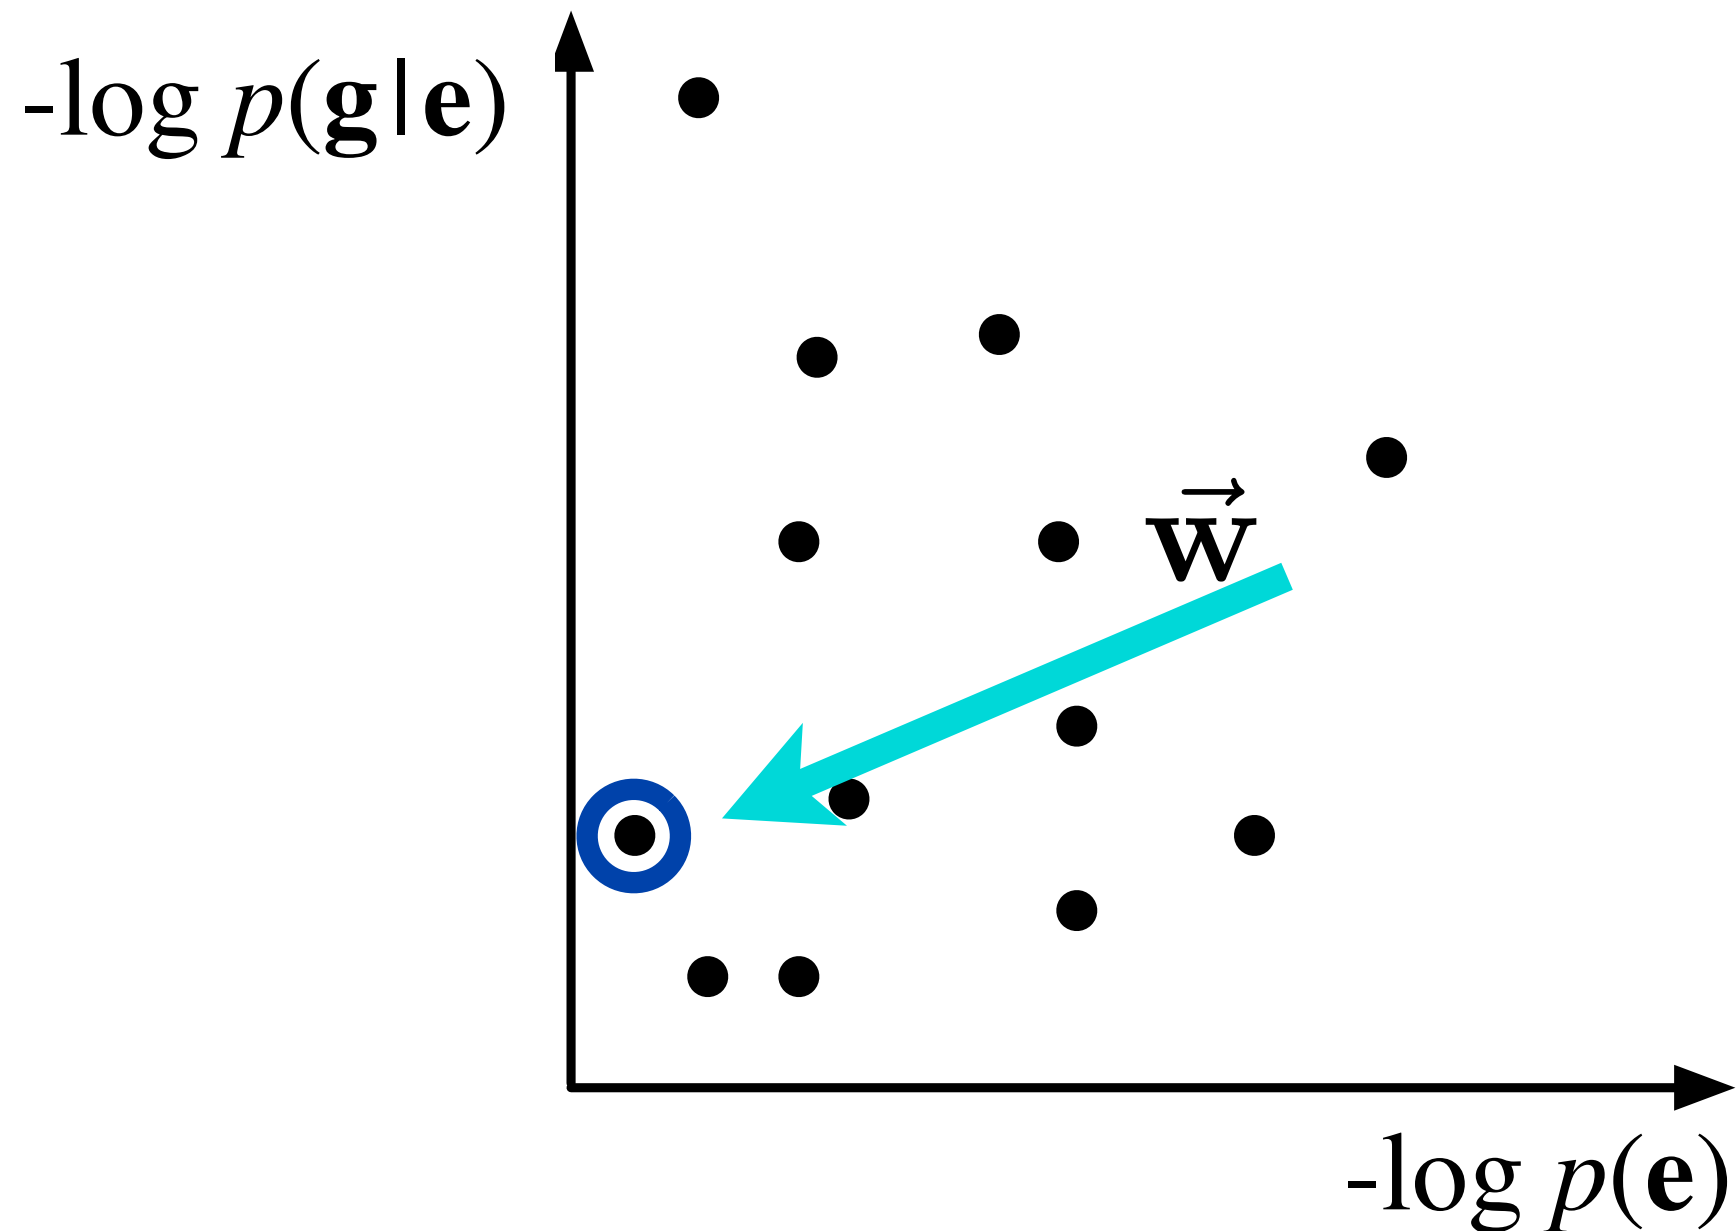
$-\log p(\mathbf{g}|\mathbf{e})$ ↑ •

Improvement 1:

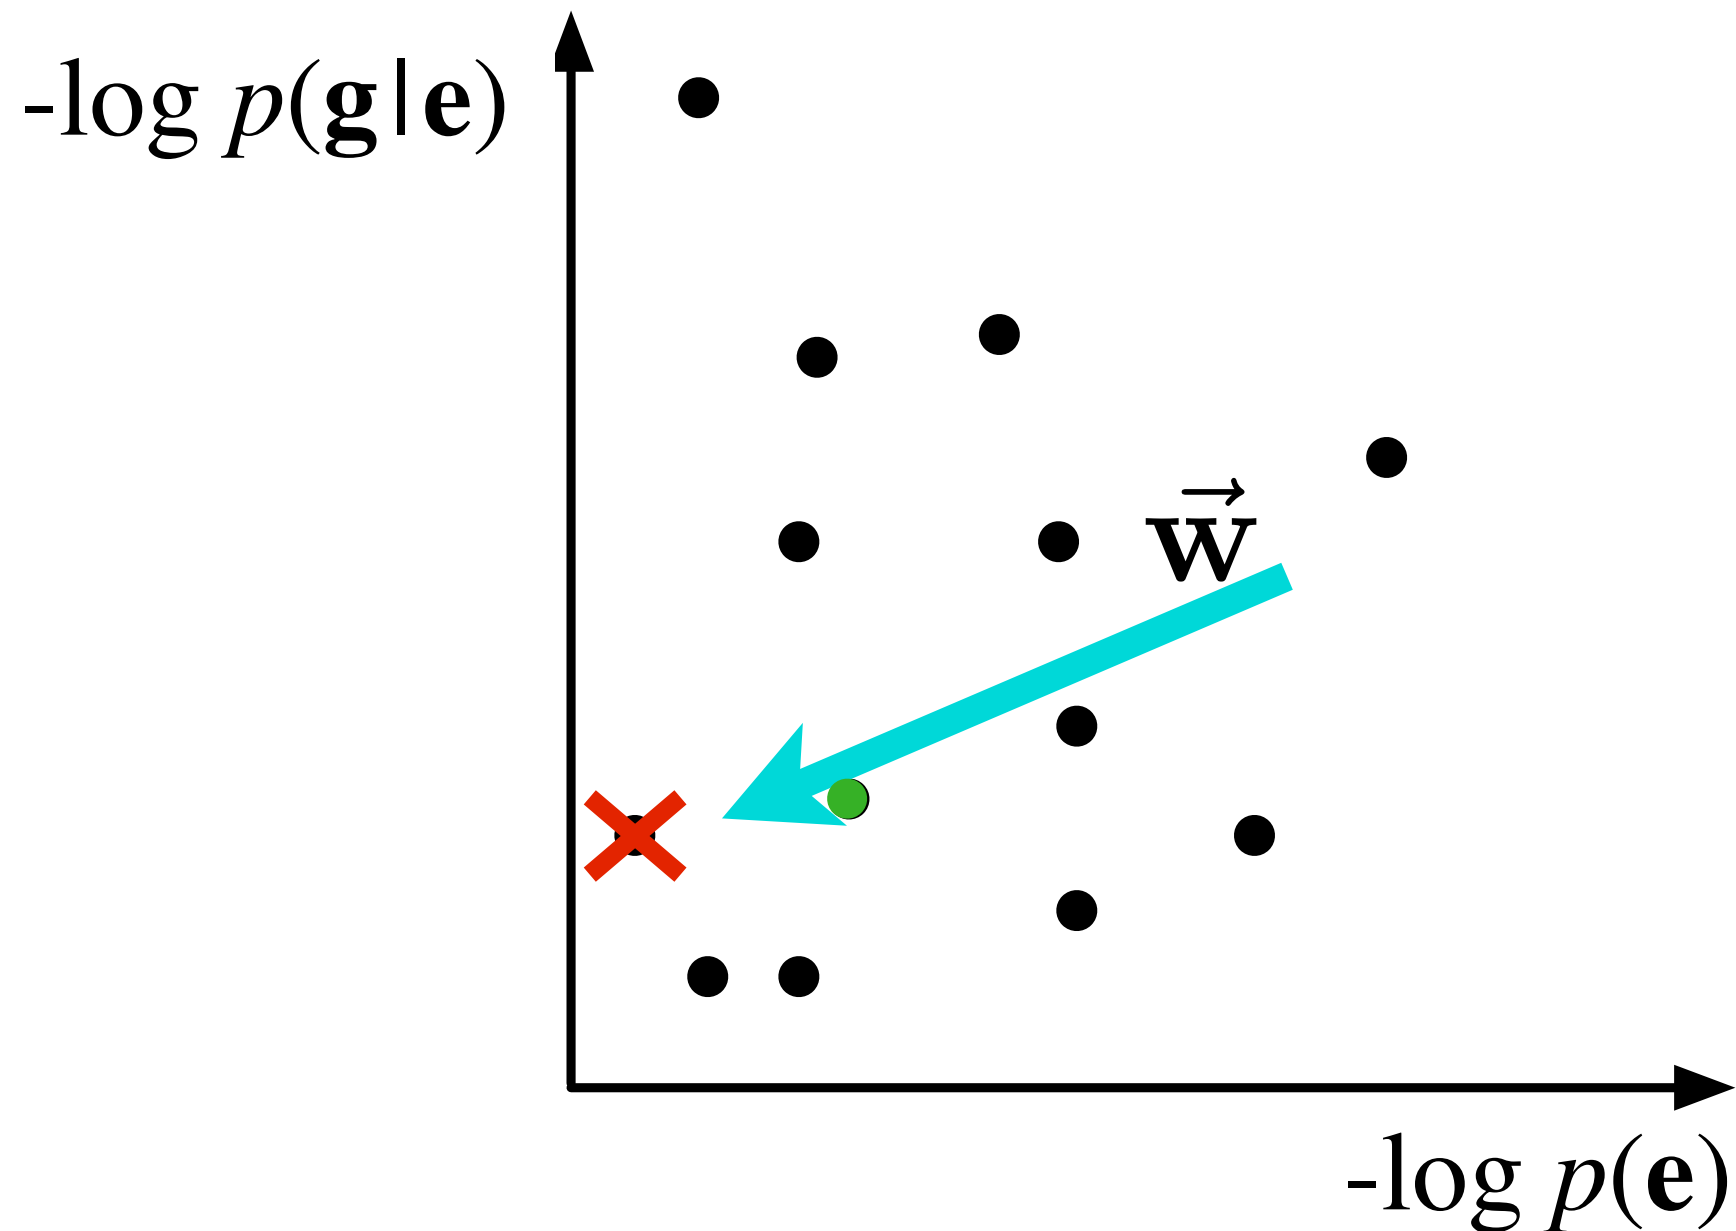
change \vec{w} to find better translations



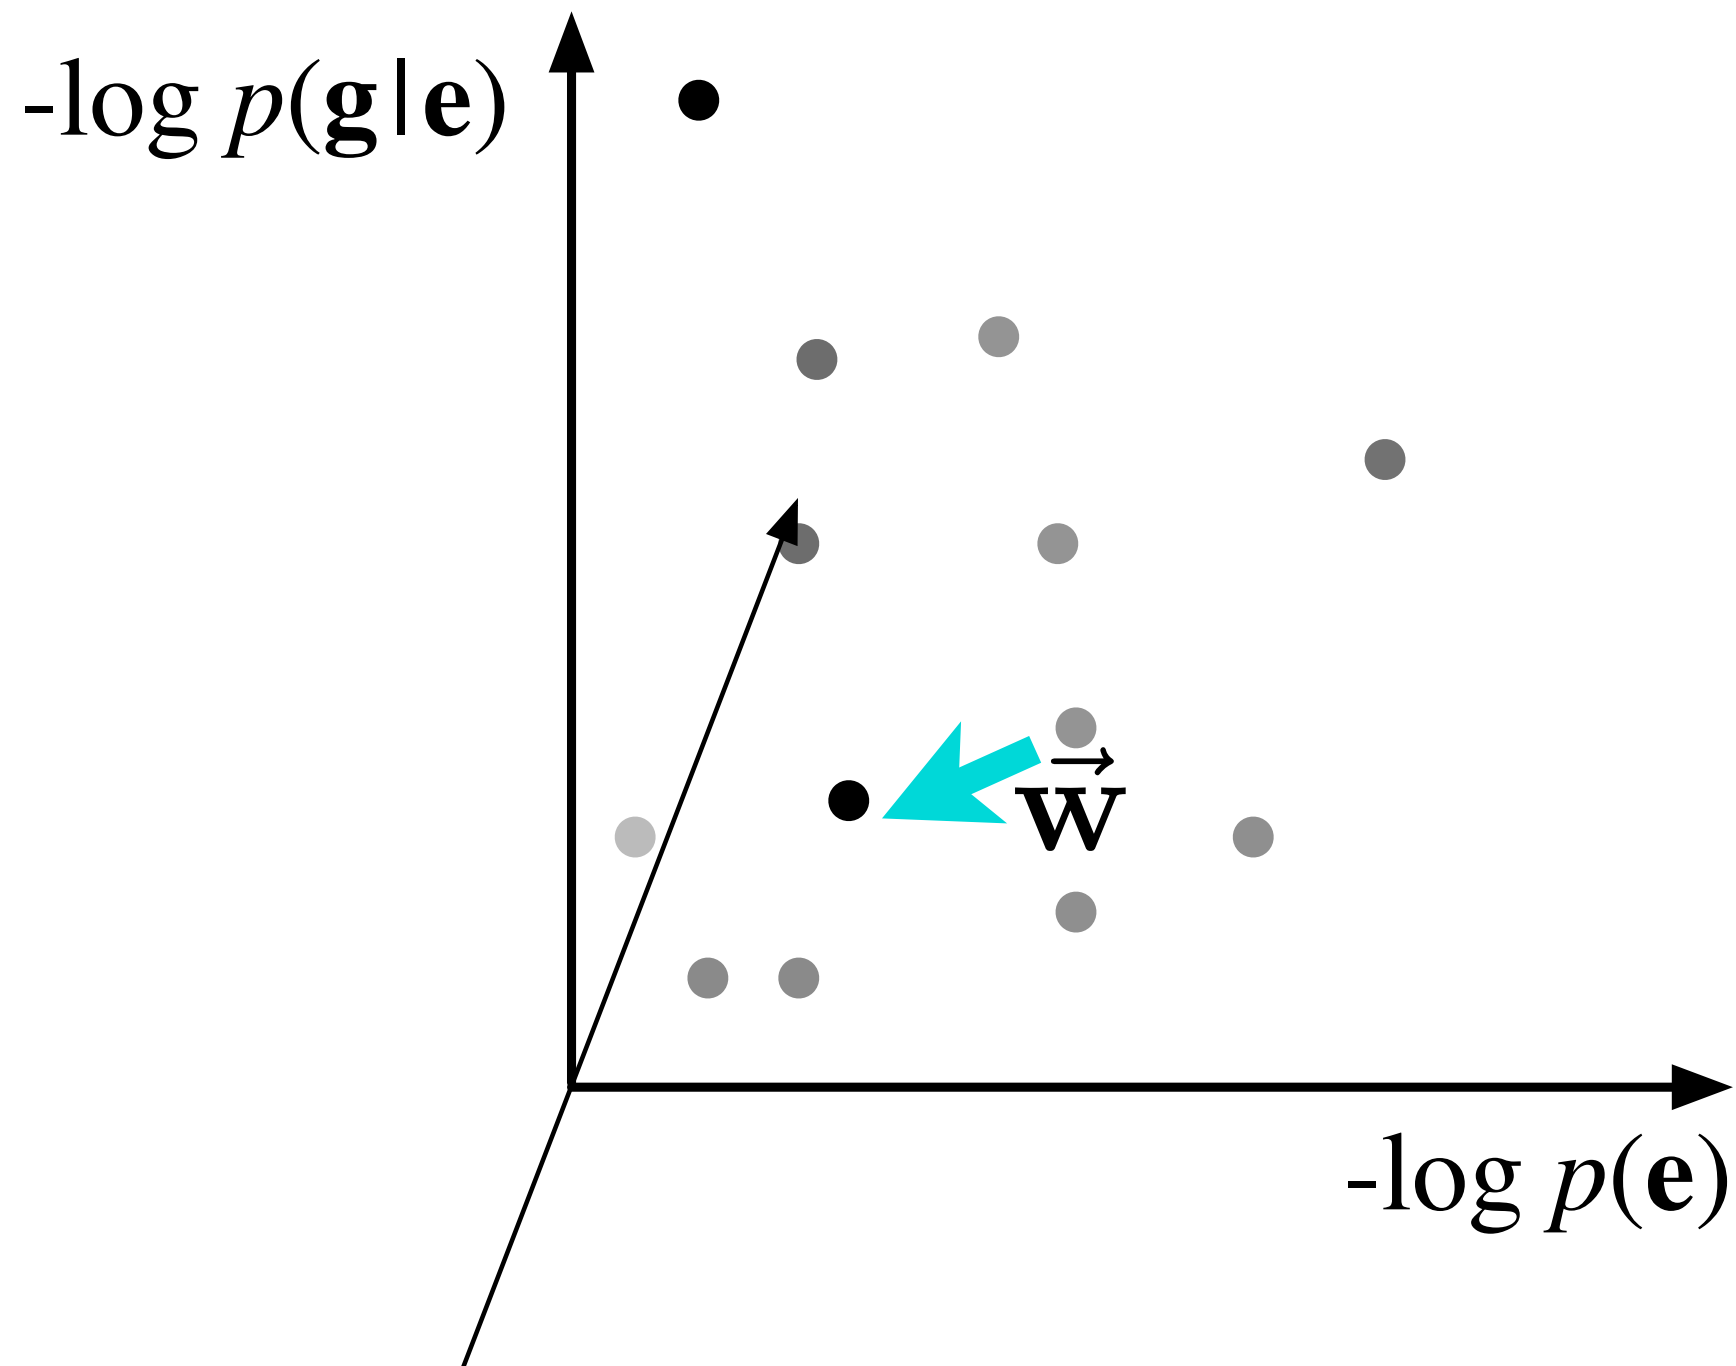
As a Linear Model



As a Linear Model



As a Linear Model

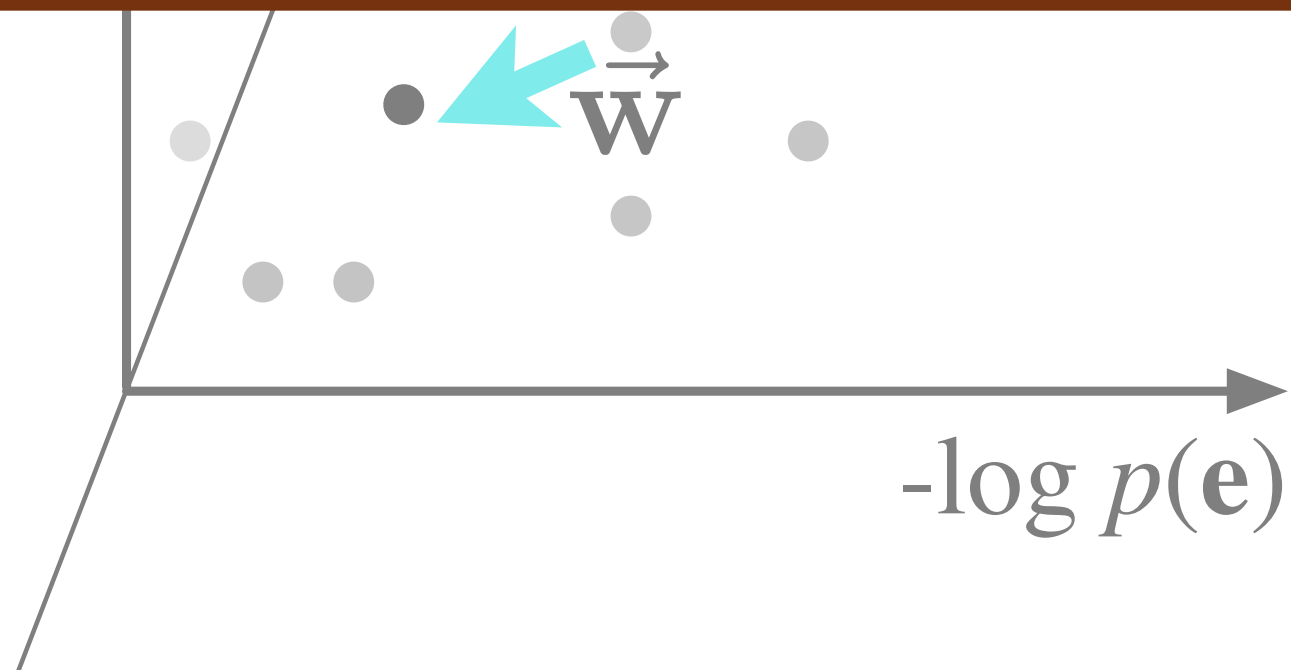


As a Linear Model

$-\log p(\mathbf{g}|\mathbf{e})$ ↑ •

Improvement 2:

Add dimensions to make points separable



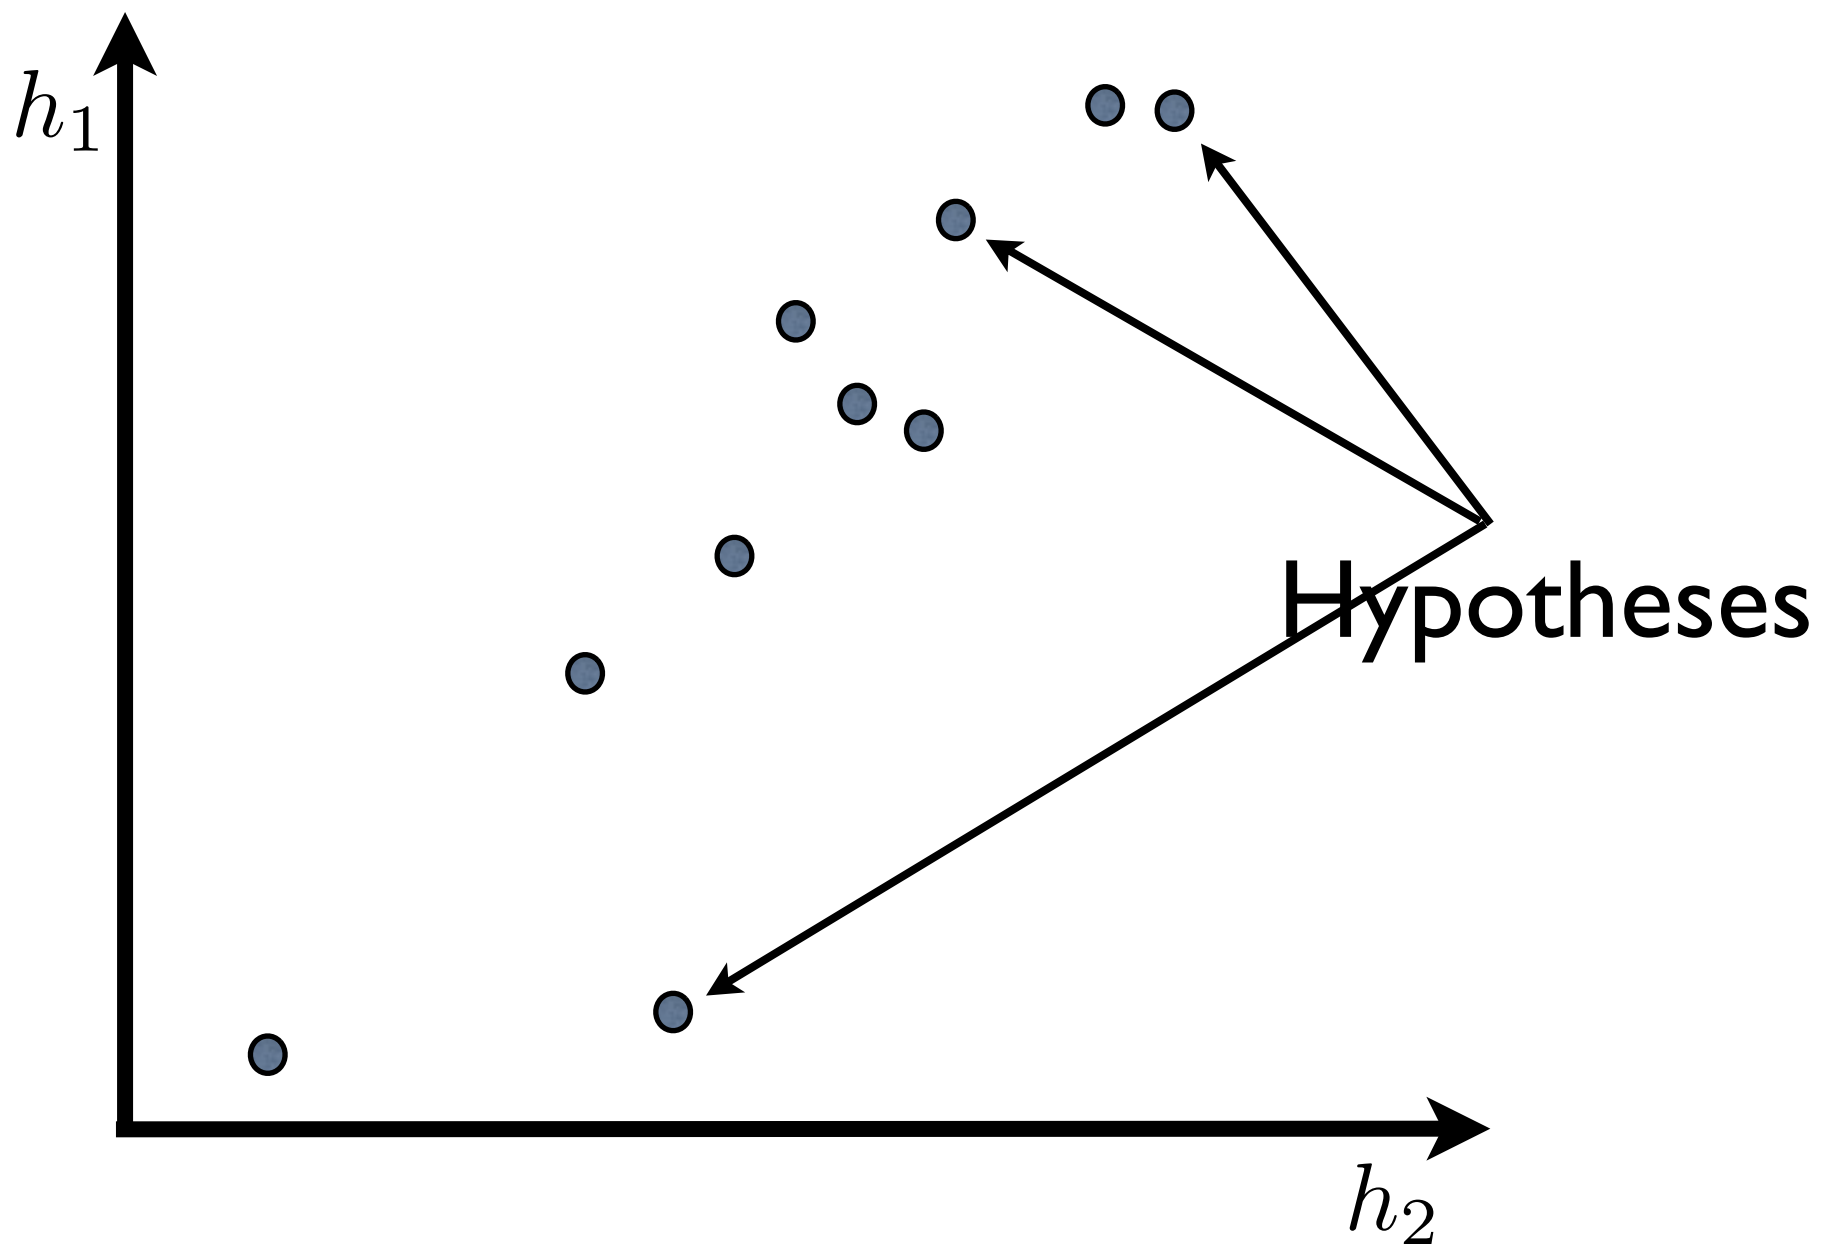
Linear Models

$$\mathbf{e}^* = \arg \max_{\mathbf{e}} \mathbf{w}^\top \mathbf{h}(\mathbf{g}, \mathbf{e})$$

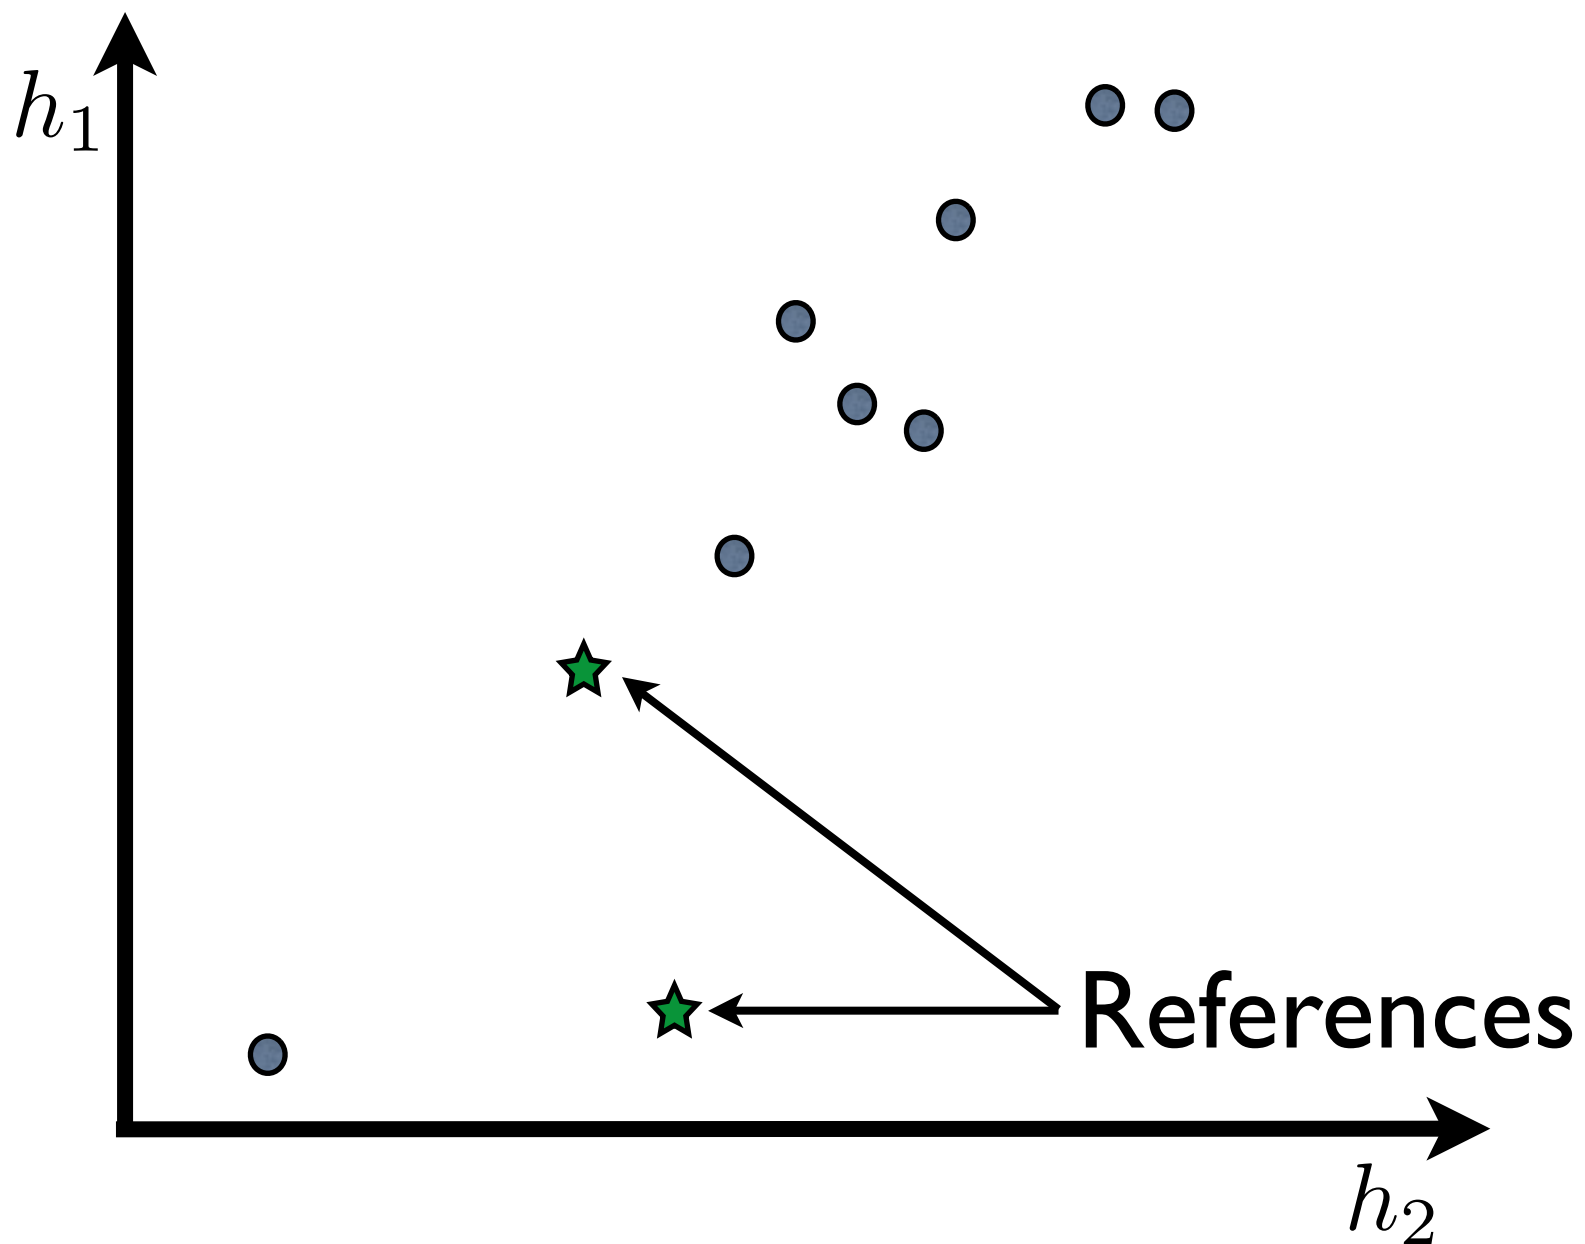
- Improve the modeling capacity of the noisy channel in two ways
 - Reorient the weight vector
 - Add new dimensions (*new features*)
- Questions
 - What features? $\mathbf{h}(\mathbf{g}, \mathbf{e})$
 - How do we set the weights? \mathbf{w}

Parameter Learning

Hypothesis Space



Hypothesis Space



Preliminaries

We assume a **decoder** that computes:

$$\langle \mathbf{e}^*, \mathbf{a}^* \rangle = \arg \max_{\langle \mathbf{e}, \mathbf{a} \rangle} \mathbf{w}^\top \mathbf{h}(\mathbf{g}, \mathbf{e}, \mathbf{a})$$

And ***K*-best lists** of, that is:

$$\{\langle \mathbf{e}_i^*, \mathbf{a}_i^* \rangle\}_{i=1}^{i=K} = \arg i^{\text{th}}\text{-max}_{\langle \mathbf{e}, \mathbf{a} \rangle} \mathbf{w}^\top \mathbf{h}(\mathbf{g}, \mathbf{e}, \mathbf{a})$$

Standard, efficient algorithms exist for this.

Cost-Sensitive Training

- Assume we have a **cost function** that gives a score for how good/bad a translation is

$$\ell(\hat{\mathbf{e}}, \mathcal{E}) \mapsto [0, 1]$$

- Optimize the weight vector by making reference to this function
 - We will talk about two ways to do this

MERT

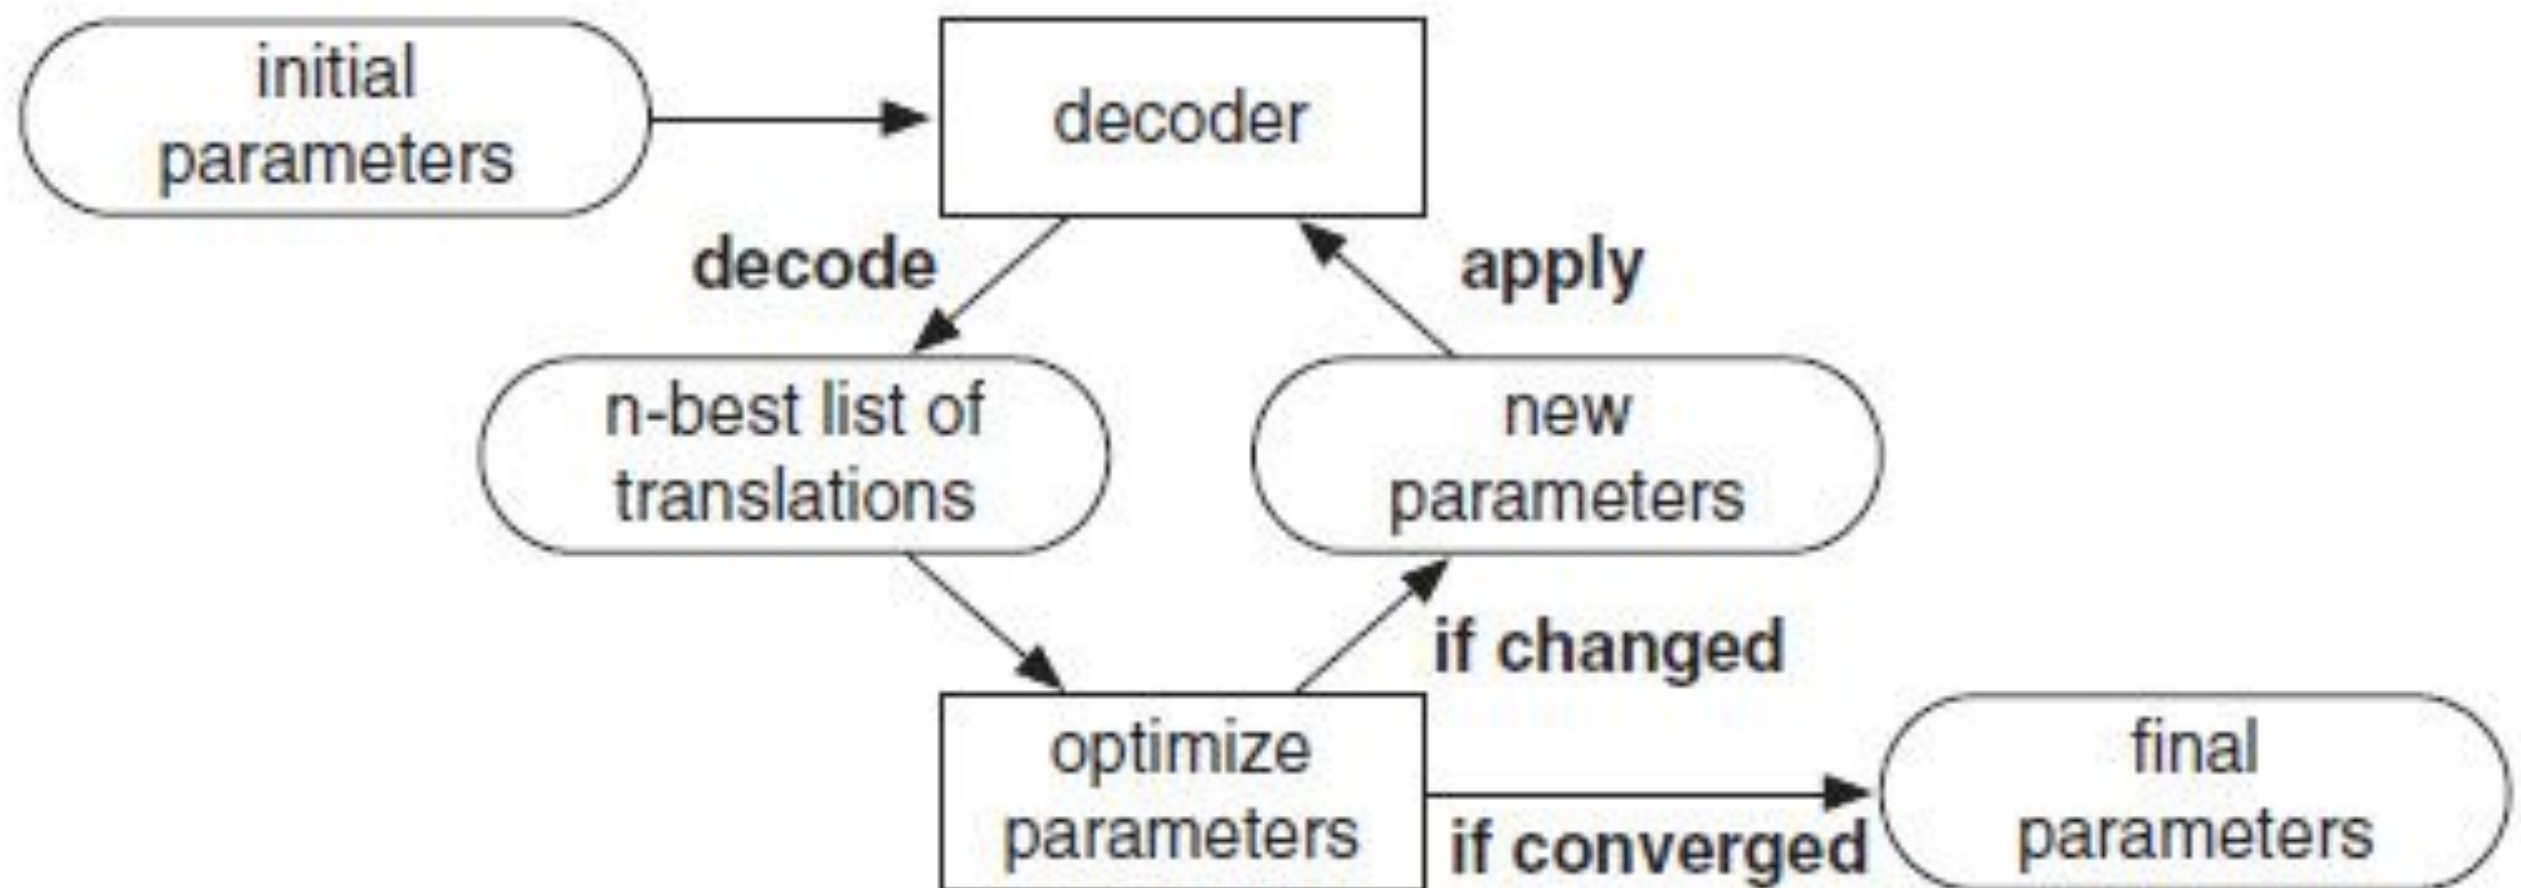


- **Minimum Error Rate Training**
- Directly optimize for an automatic evaluation metric instead of likelihood
- Maximize the BLEU score on a **held out development set**
- Iteratively update the parameters by re-scoring n-best lists and comparing the highest scoring translation to the reference

MERT

- Even with 10-15 features it's not possible to exhaustively search the space of possible feature values
- We need a good heuristic method to search the space
- Another problem: the initial parameters might be so bad that the original n-best list is not a good sample of the translations

Iterative parameter tuning



Powell Search

- Explore a high-dimensional space by finding a better point along one line in the space
- Simplest form: Vary one parameter at a time
- If the optimal value is better than the current value, then change it and move to the next parameter
- Iterate until there are no single parameter updates that increase the score

Powell Search

- Problem: searching for the best value for a single parameter is still daunting
- Parameters are real-valued #s, so they have an infinite number of possible values
- Key insight of MERT: only a small number of threshold values will change the I-best translation
- Only I-best translations change BLEU

Finding the threshold points for 1 sentence

Given weight vector \mathbf{w} , any hypothesis $\langle \mathbf{e}, \mathbf{a} \rangle$ will have a (scalar) score $m = \mathbf{w}^\top \mathbf{h}(\mathbf{g}, \mathbf{e}, \mathbf{a})$

Now pick a **search vector** \mathbf{v} , and consider how the score of this hypothesis will change:

$$\mathbf{w}_{\text{new}} = \mathbf{w} + \gamma \mathbf{v}$$

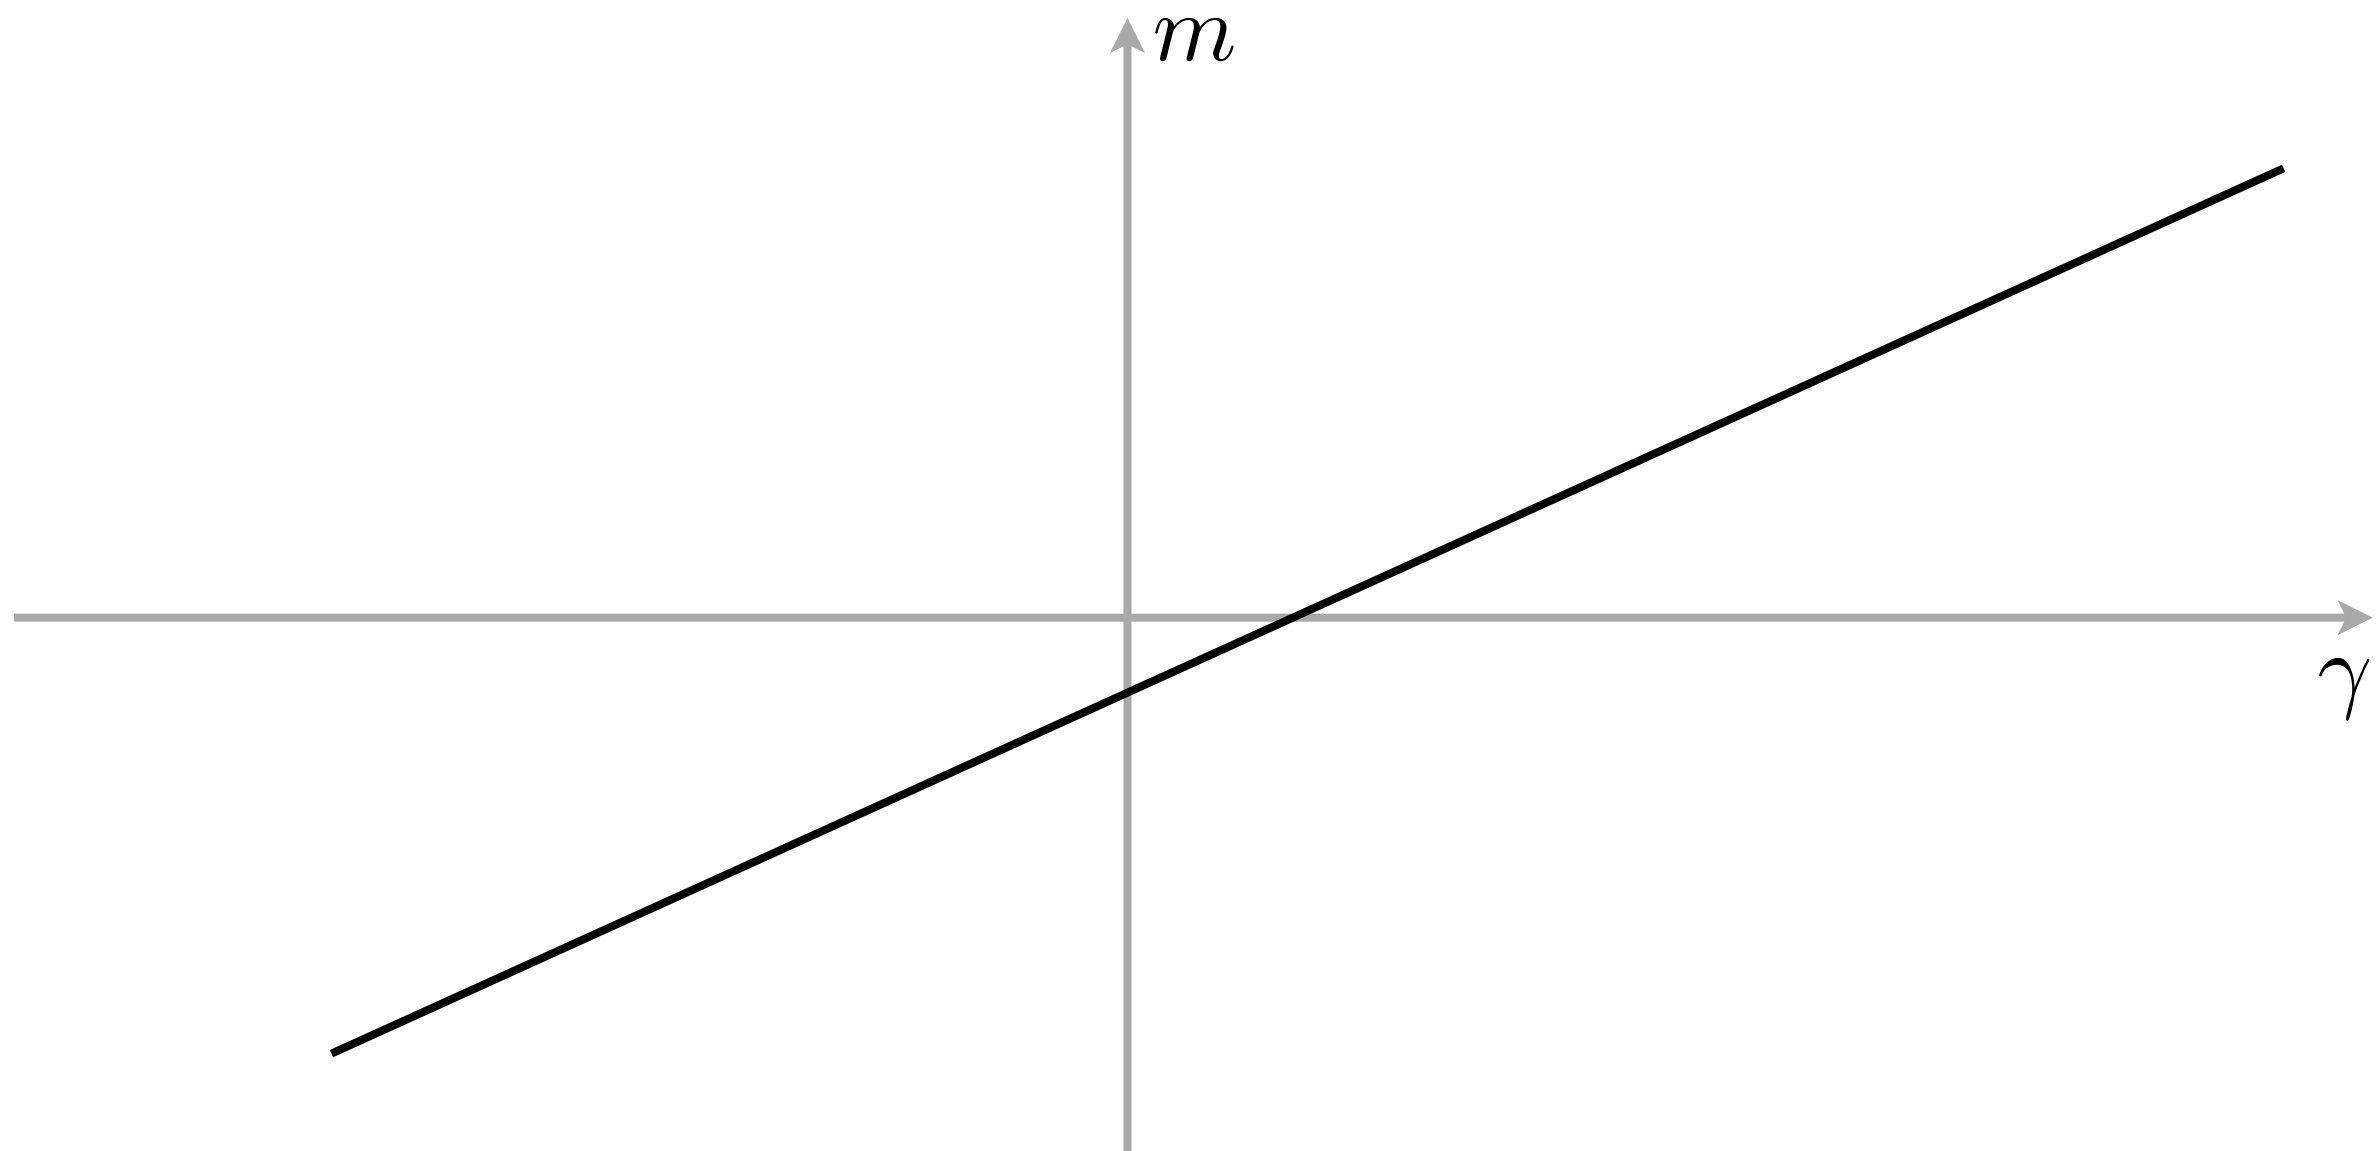
$$m = (\mathbf{w} + \gamma \mathbf{v})^\top \mathbf{h}(\mathbf{g}, \mathbf{e}, \mathbf{a})$$

$$= \underbrace{\mathbf{w}^\top \mathbf{h}(\mathbf{g}, \mathbf{e}, \mathbf{a})}_b + \gamma \underbrace{\mathbf{v}^\top \mathbf{h}(\mathbf{g}, \mathbf{e}, \mathbf{a})}_a$$

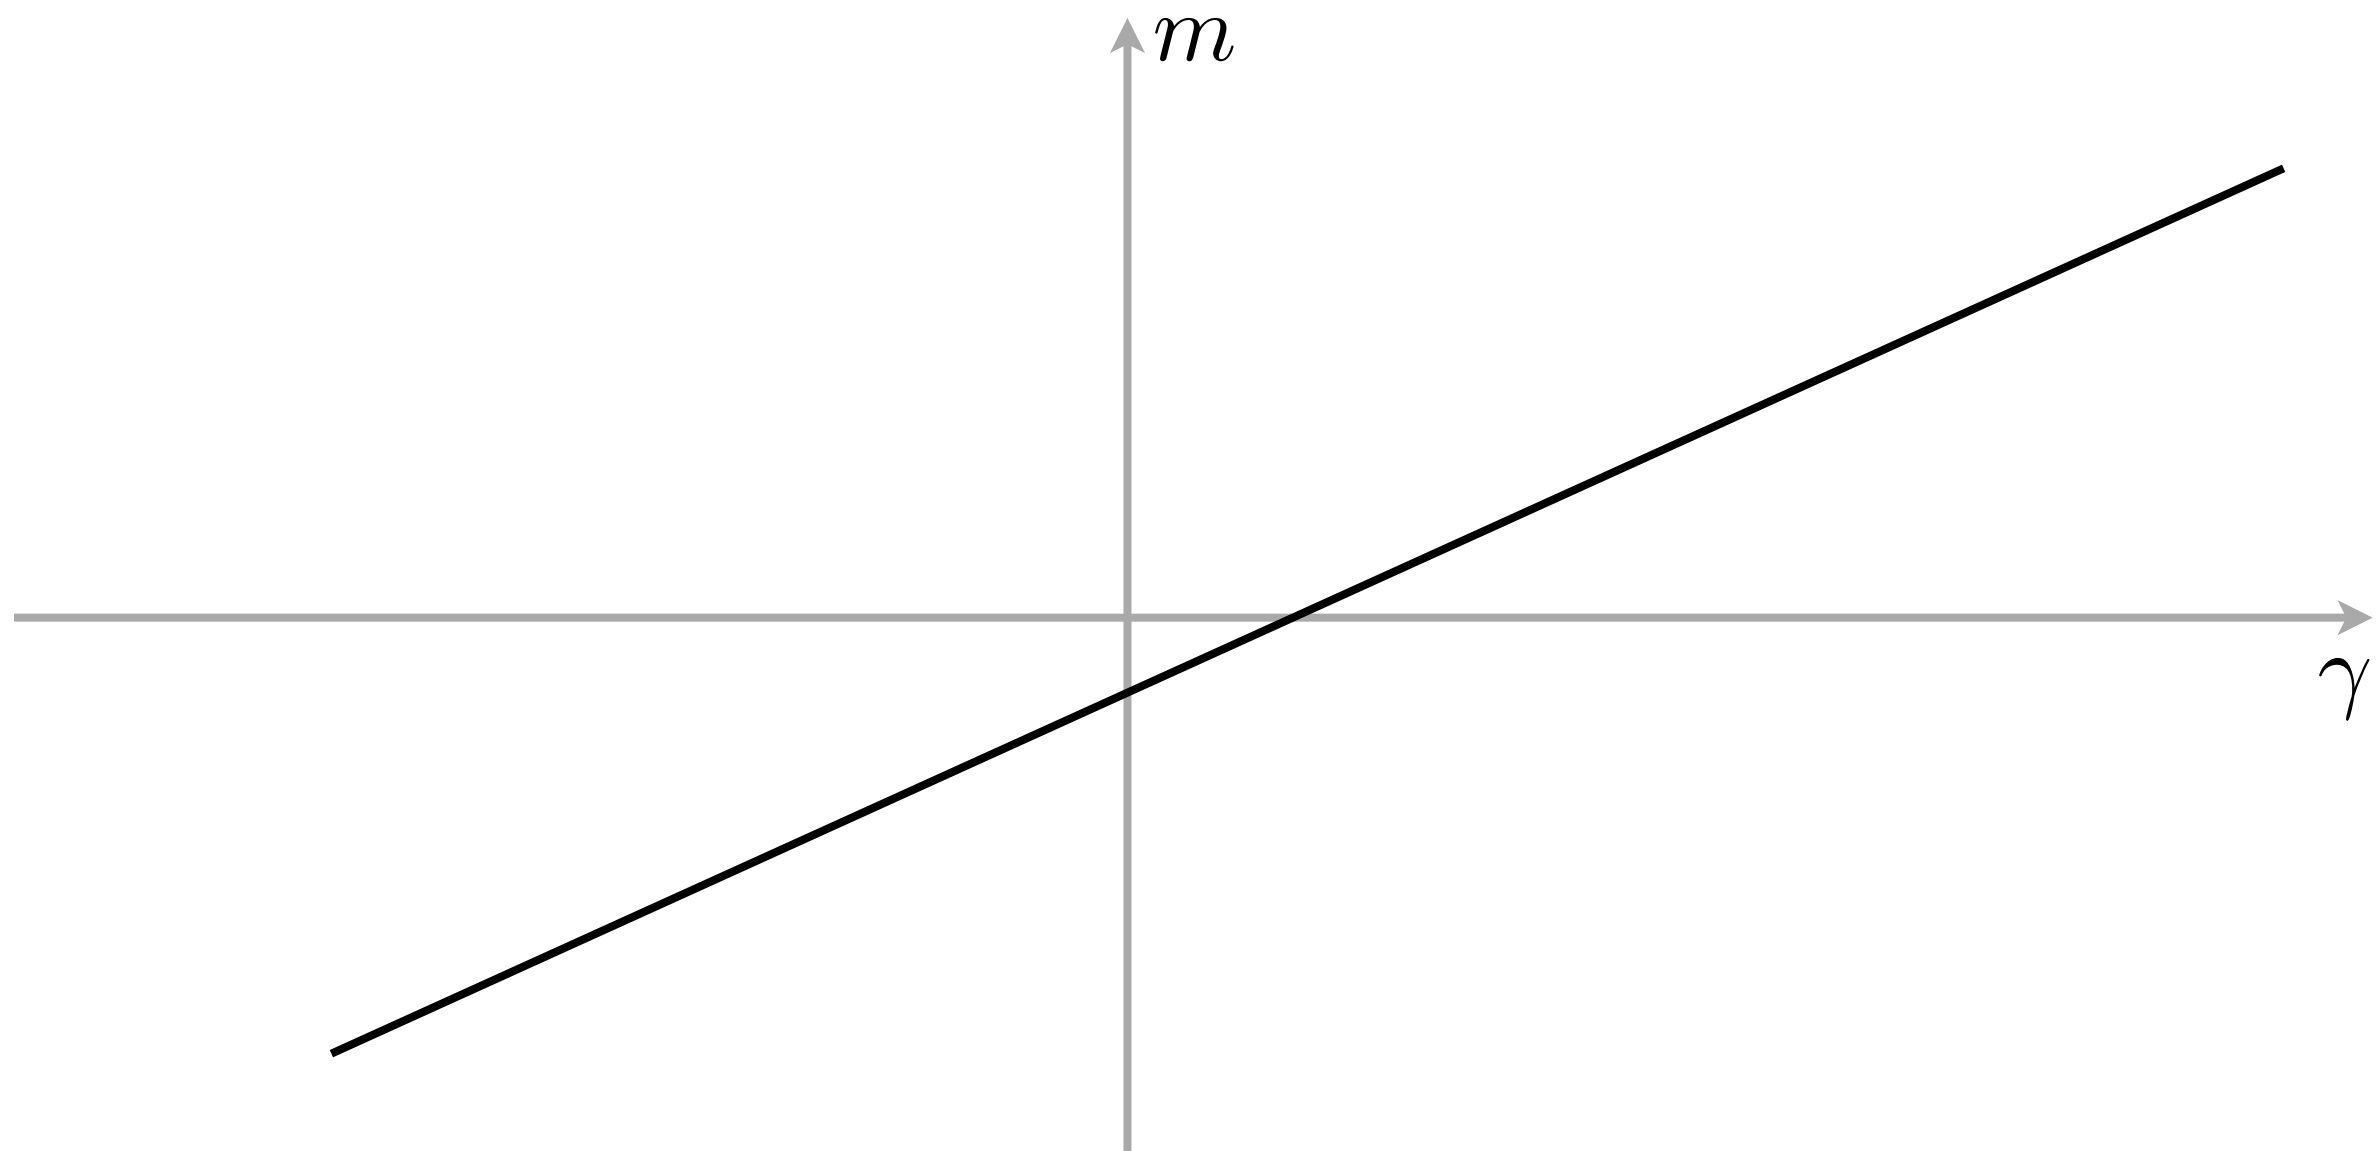
$$m = a\gamma + b$$

Linear function in 2D!

MERT

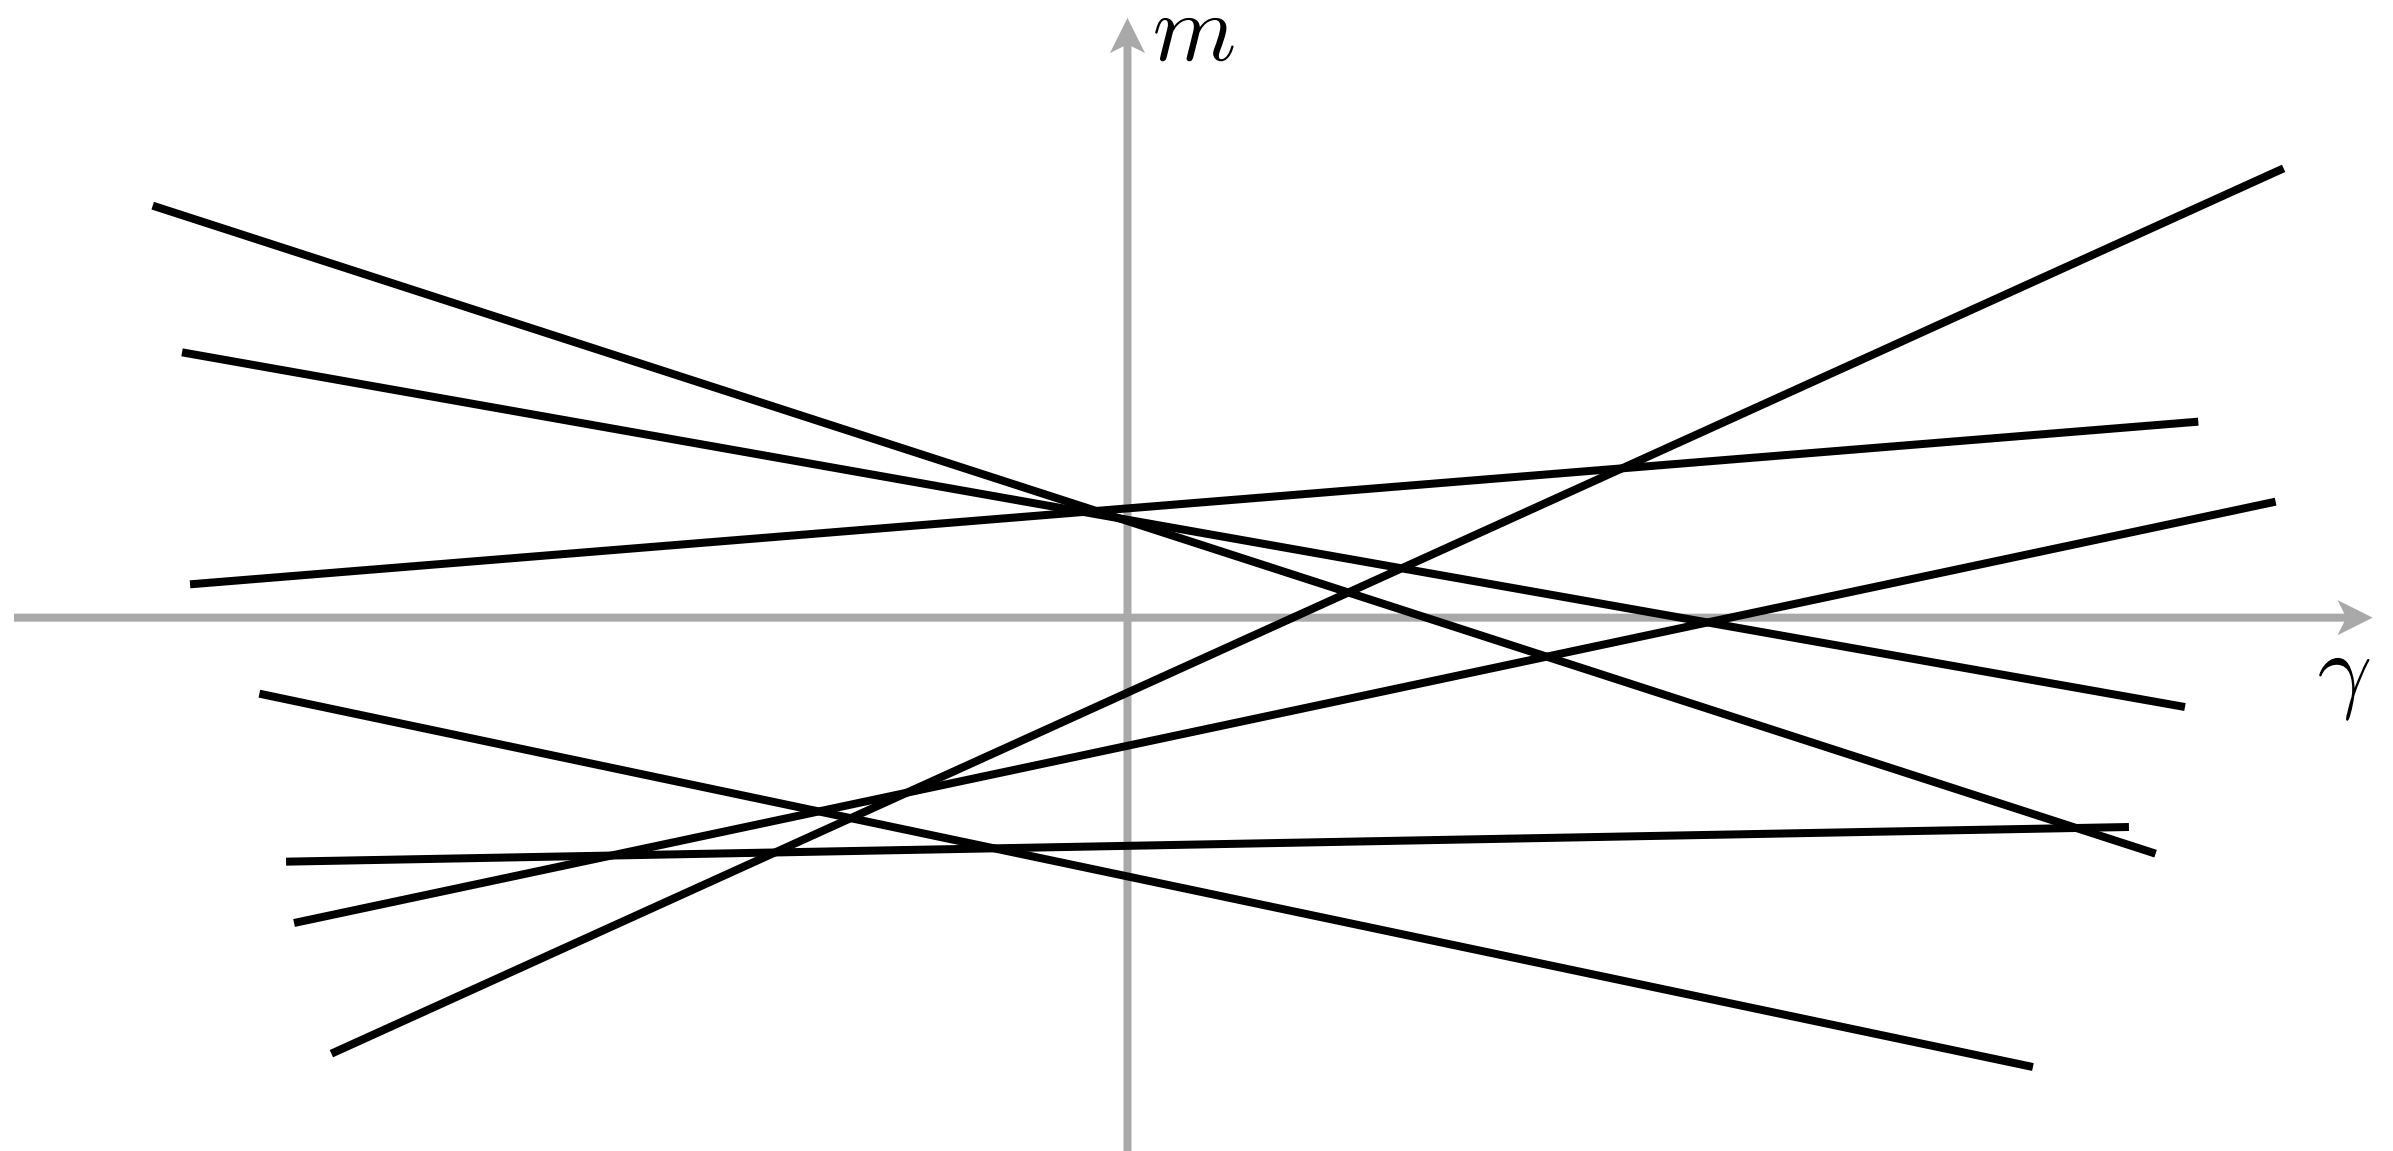


MERT



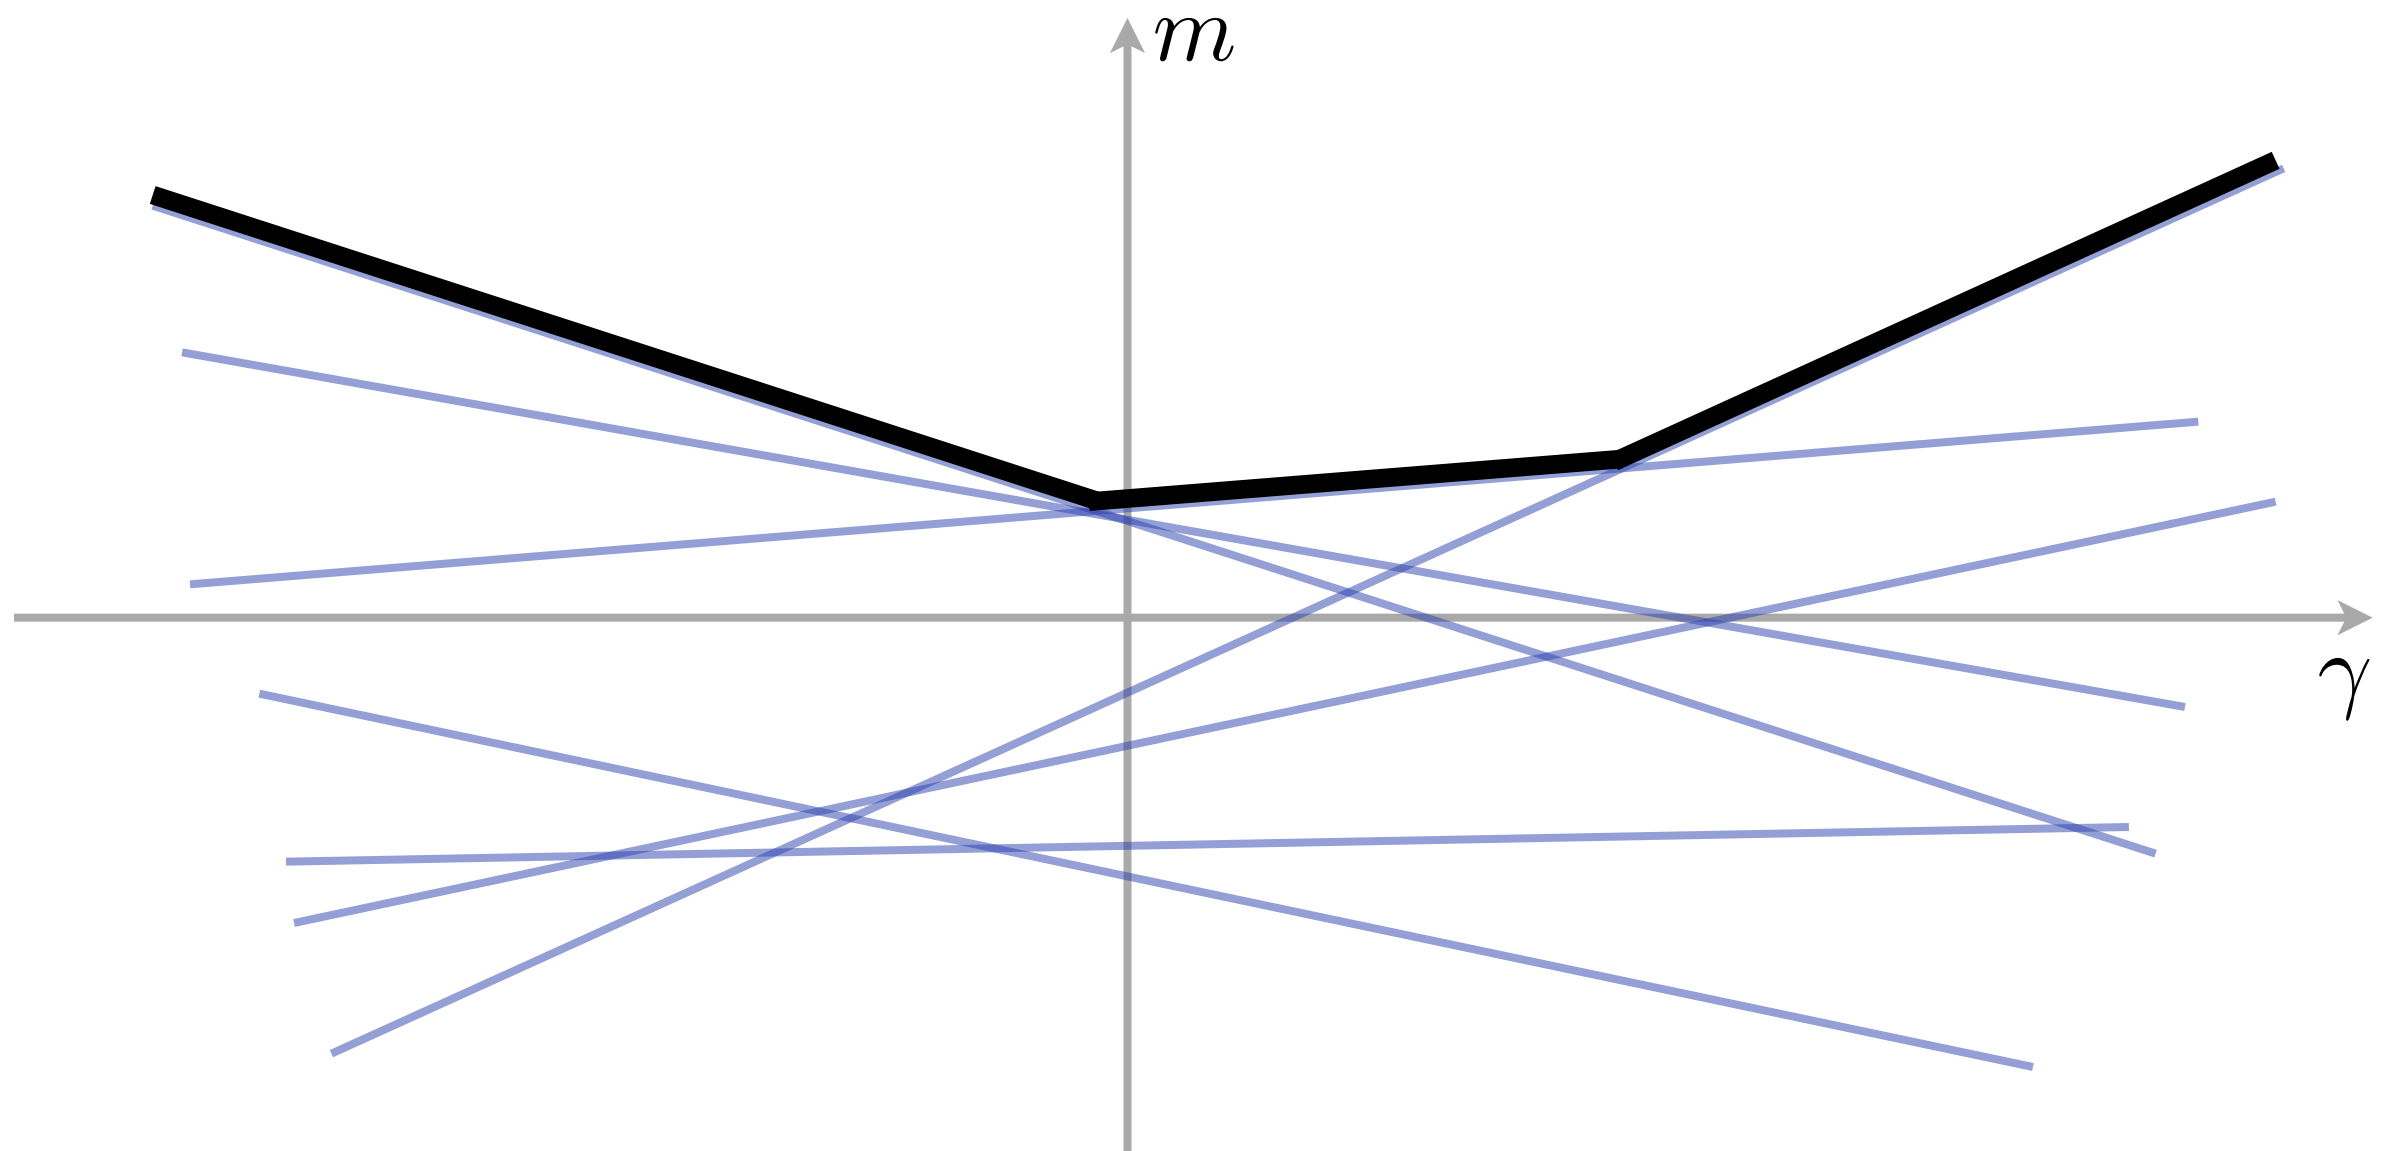
Recall our k-best set $\{\langle \mathbf{e}_i^*, \mathbf{a}_i^* \rangle\}_{i=1}^K$

MERT

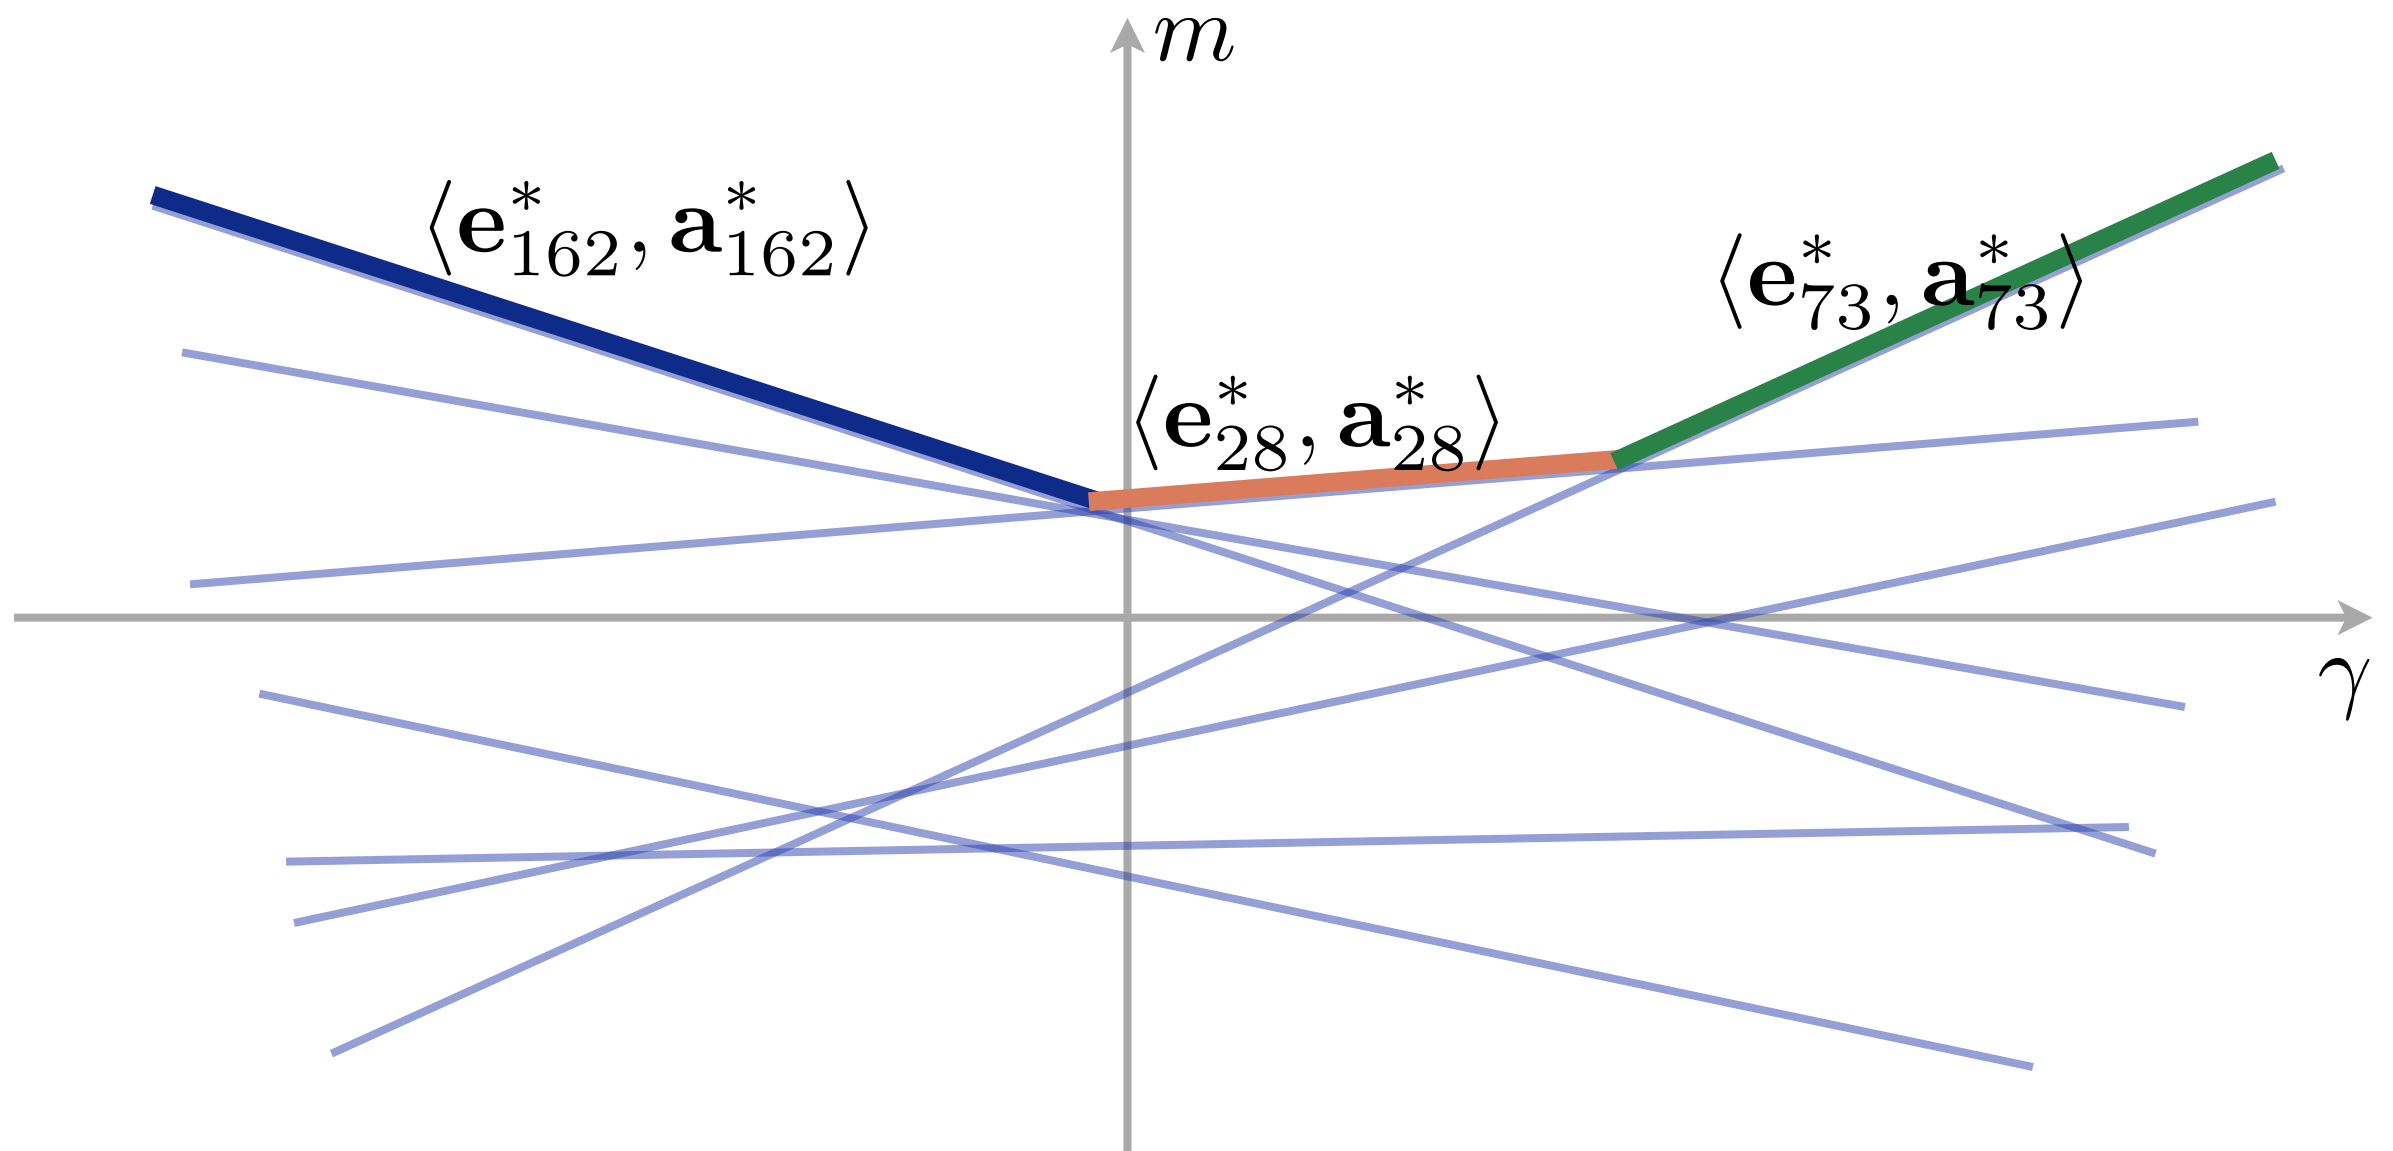


Recall our k-best set $\{\langle \mathbf{e}_i^*, \mathbf{a}_i^* \rangle\}_{i=1}^K$

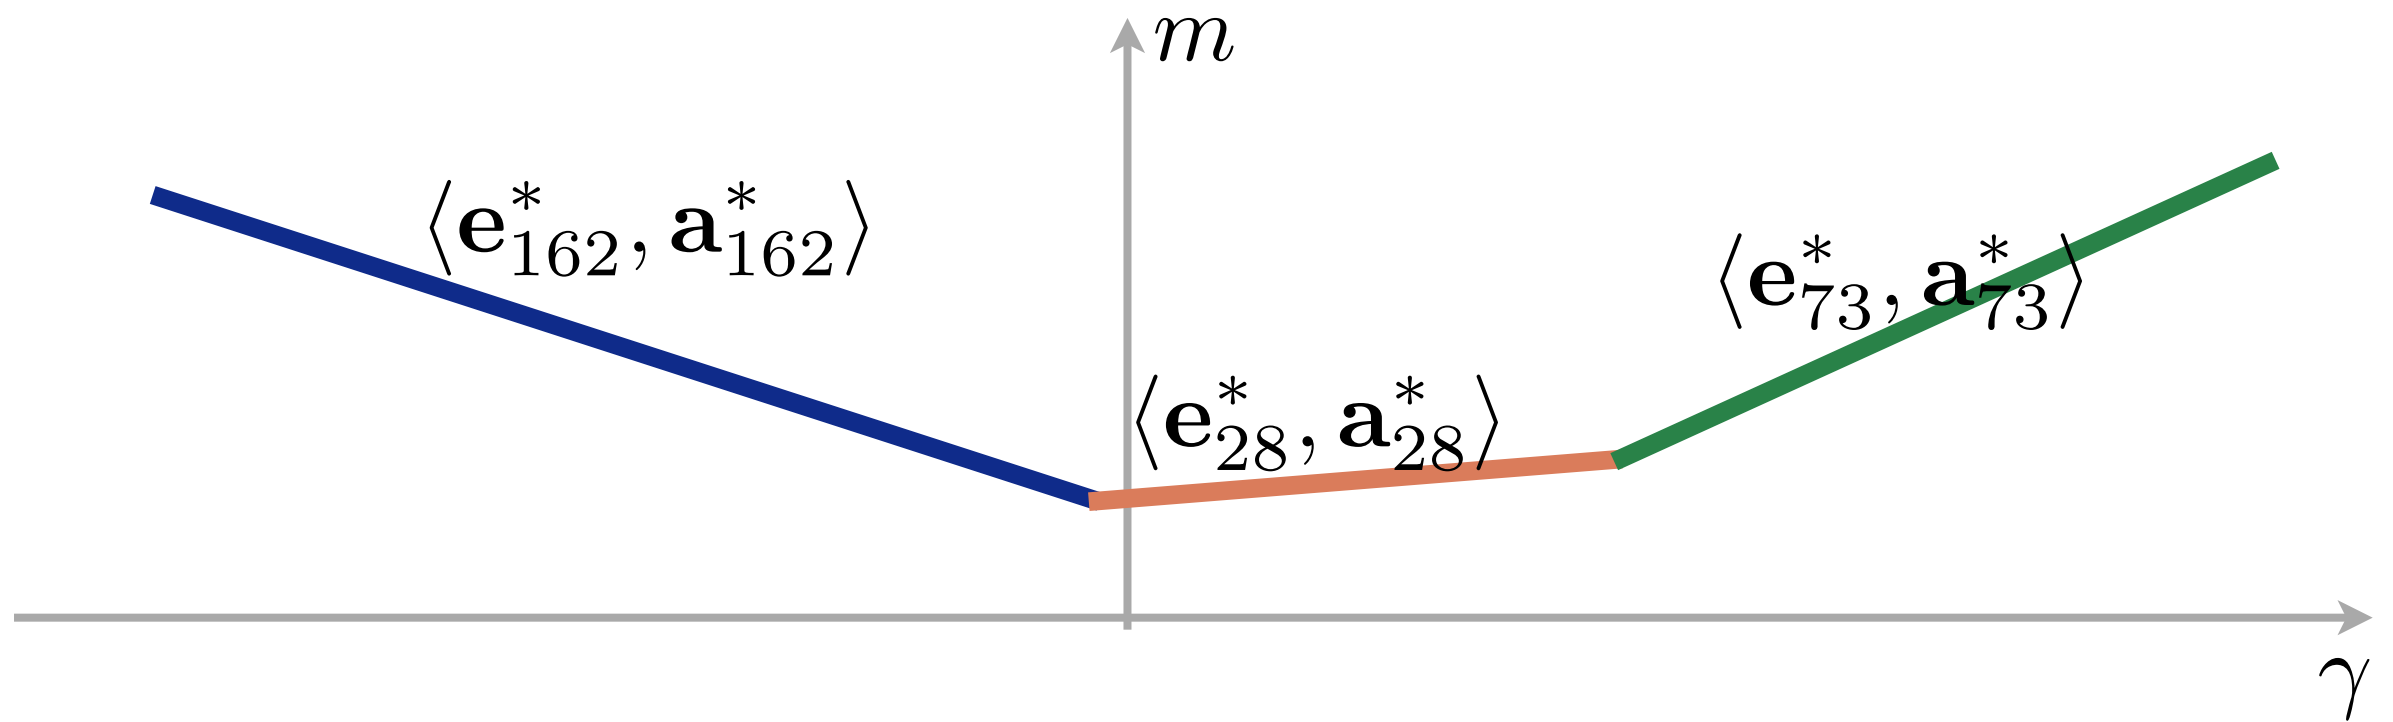
MERT



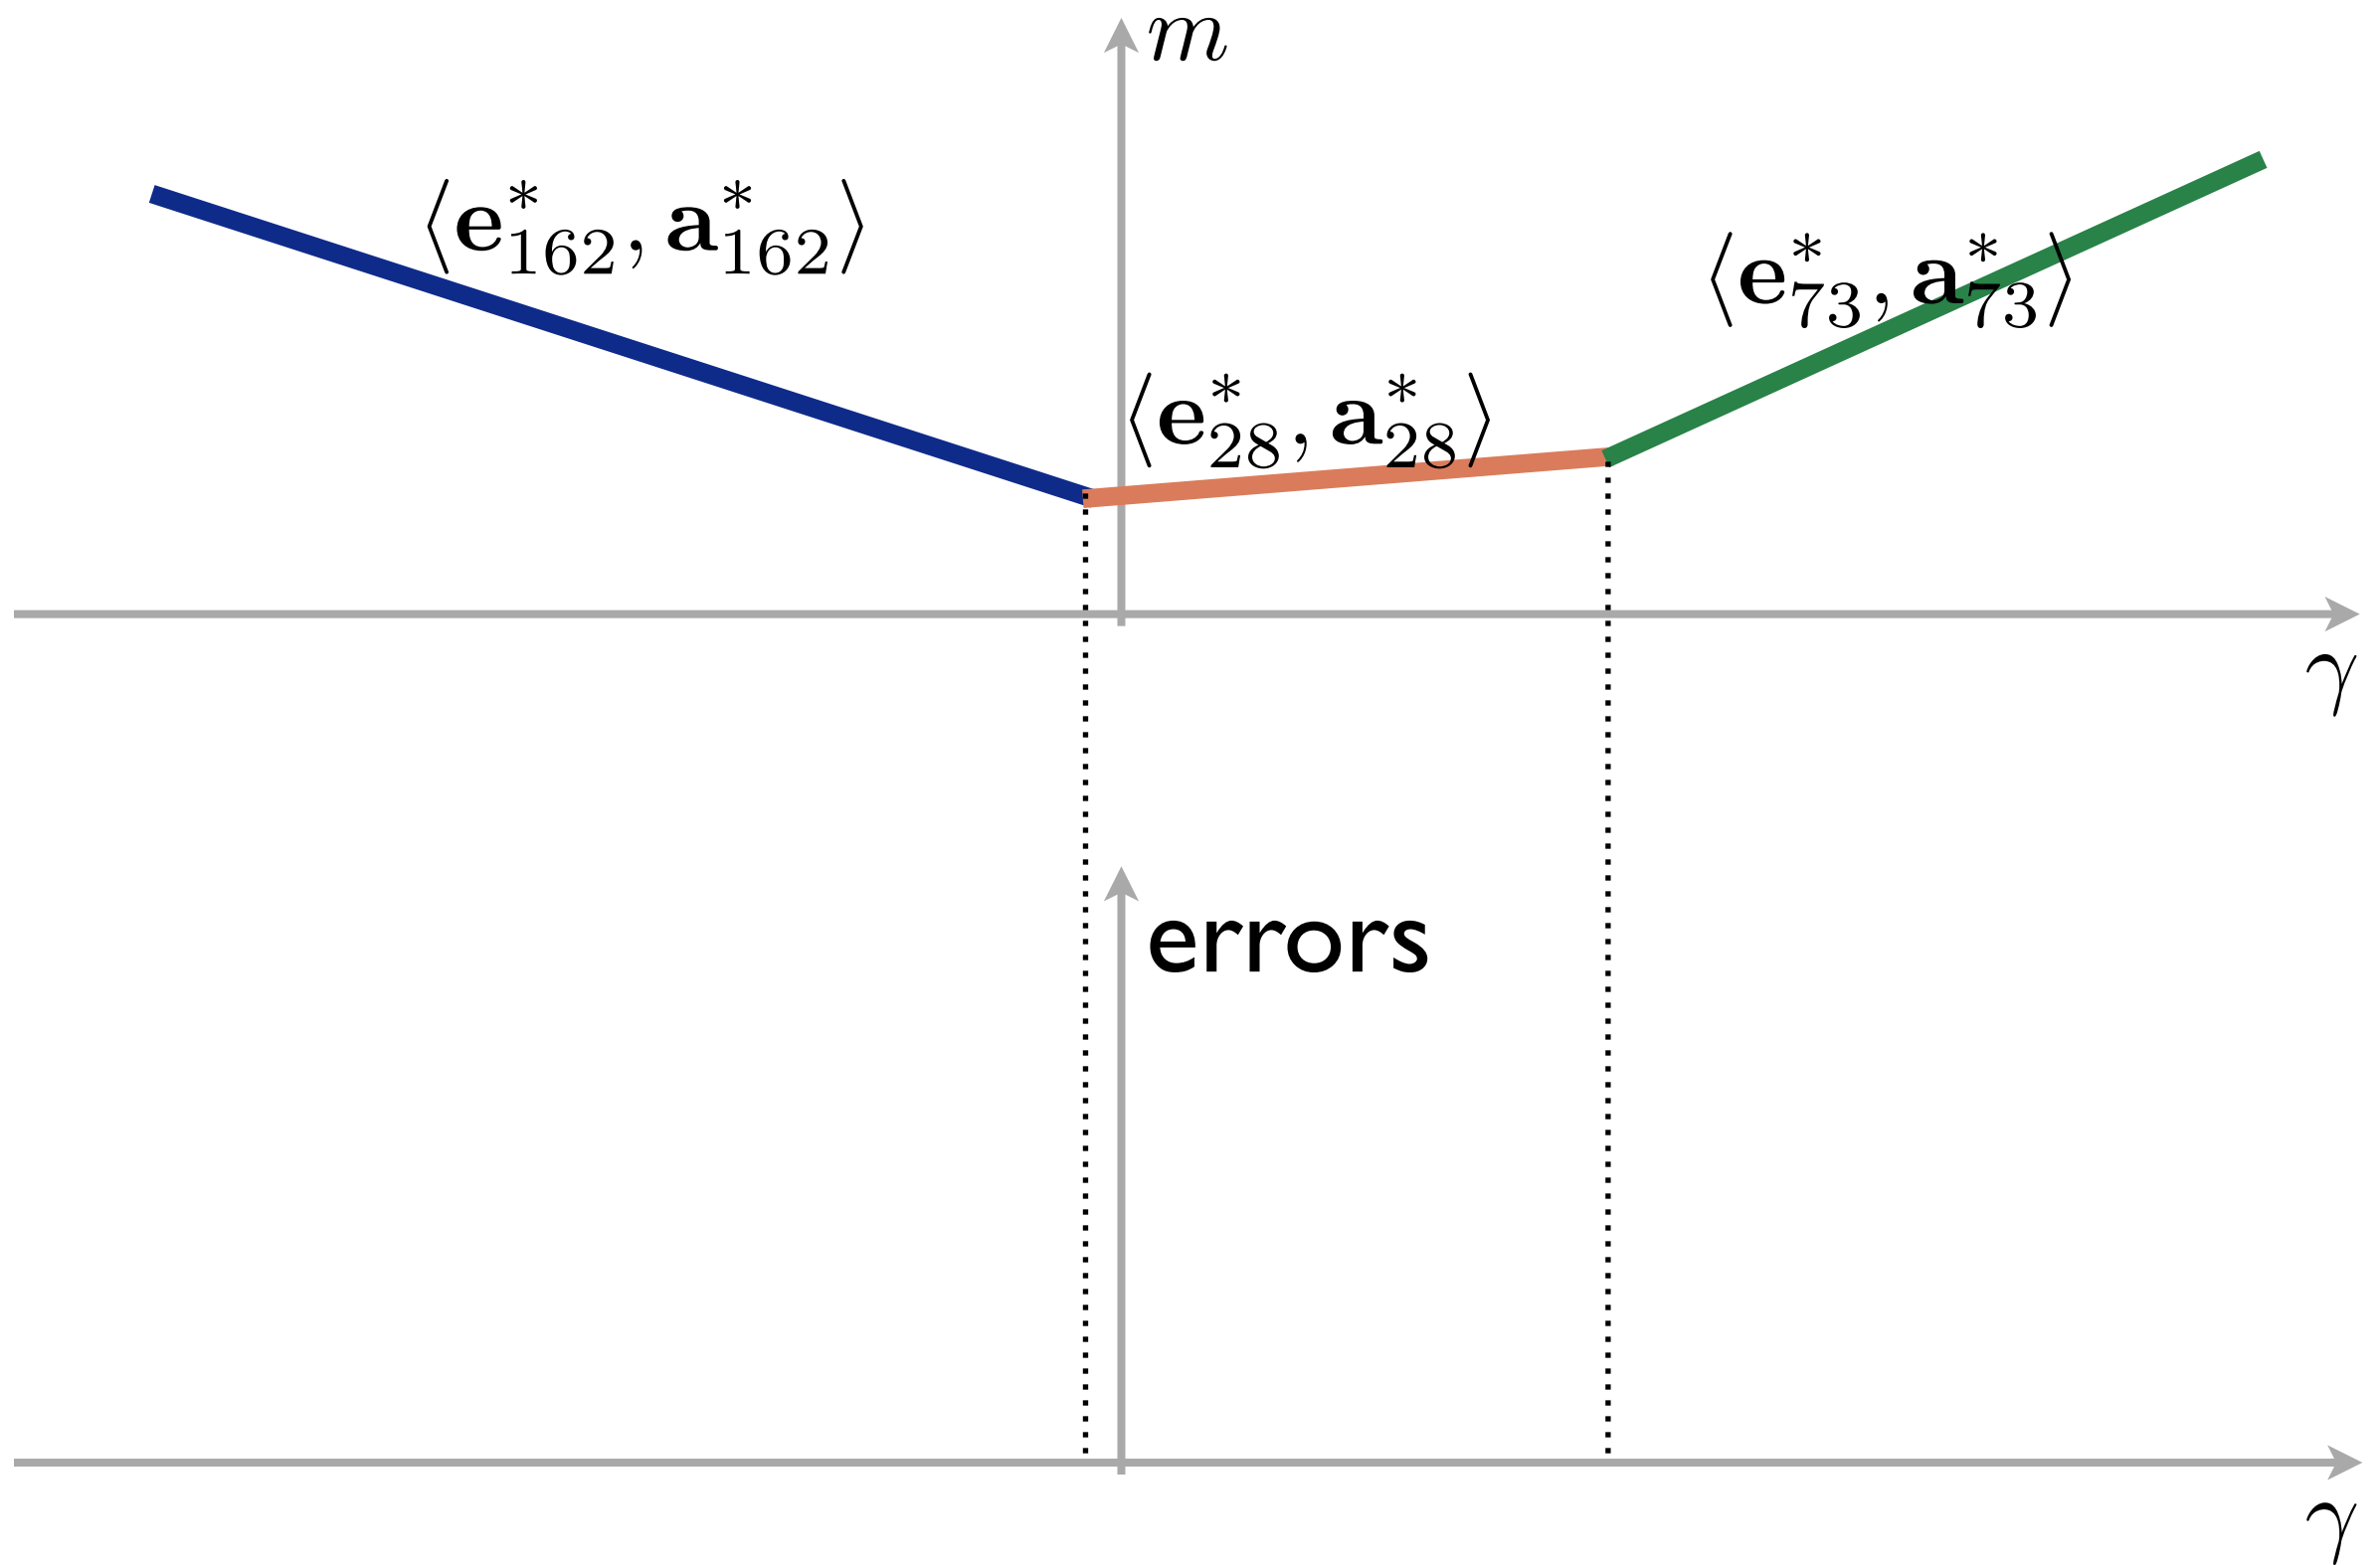
MERT



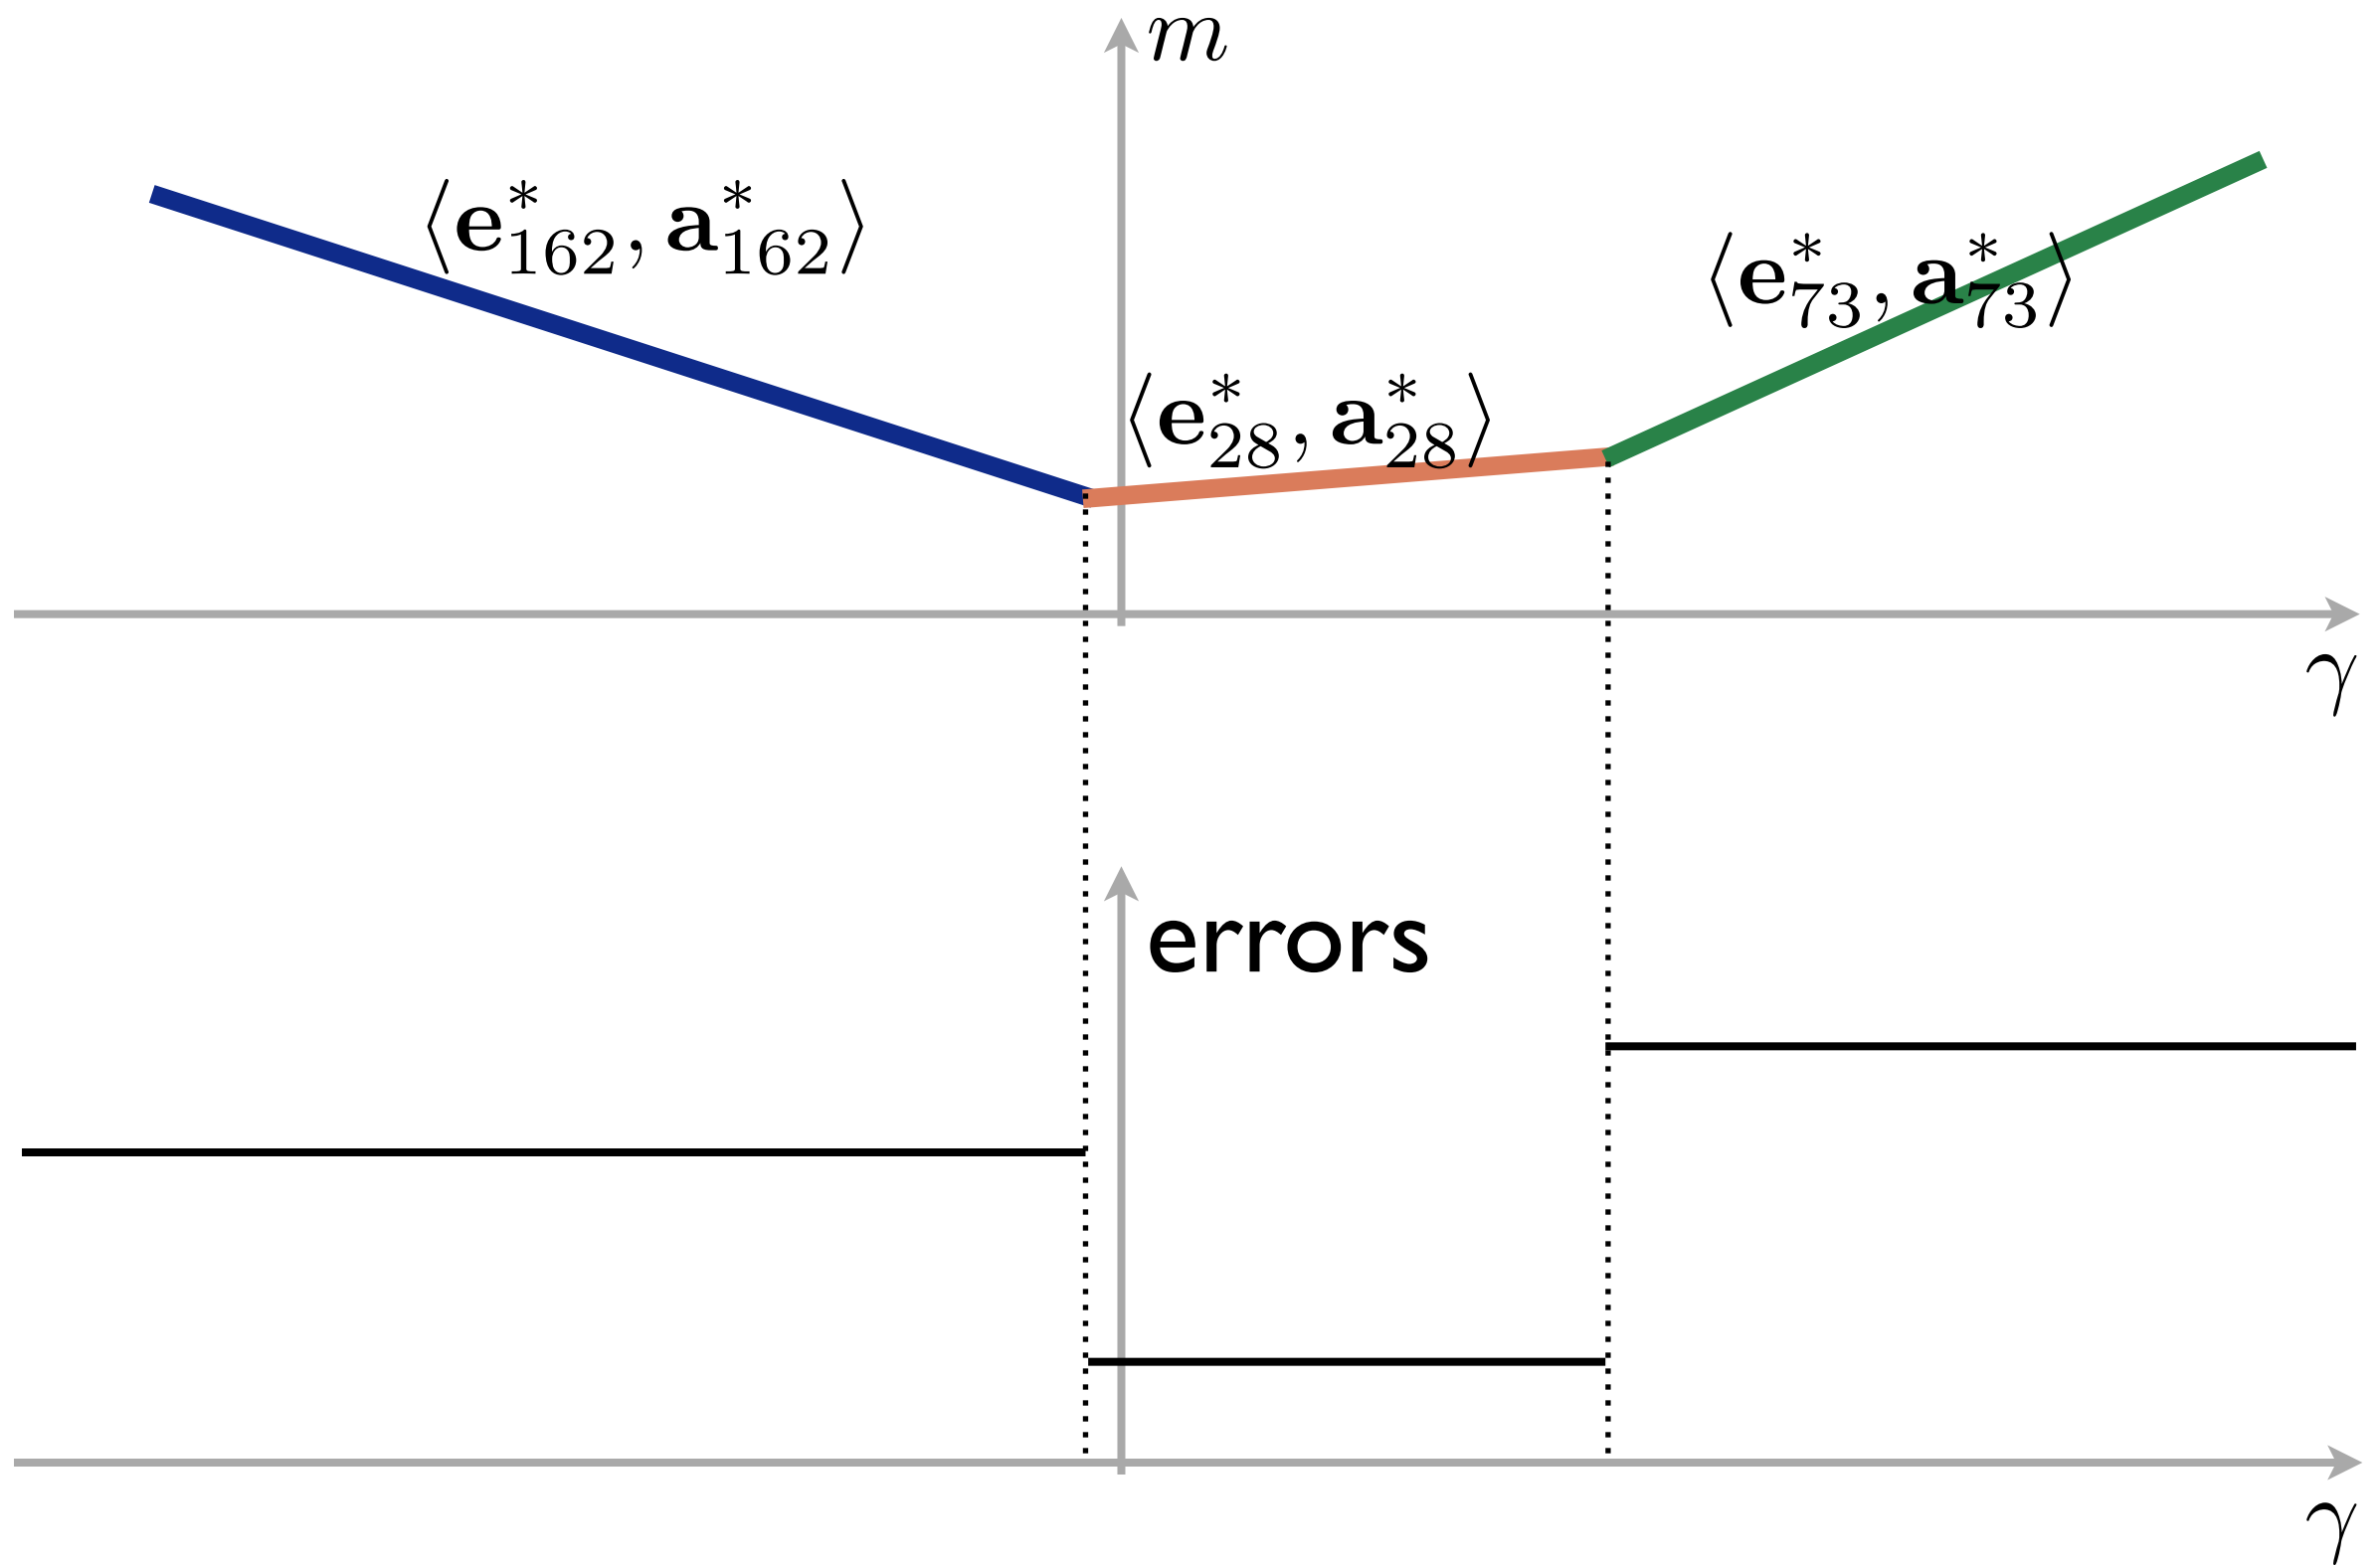
MERT



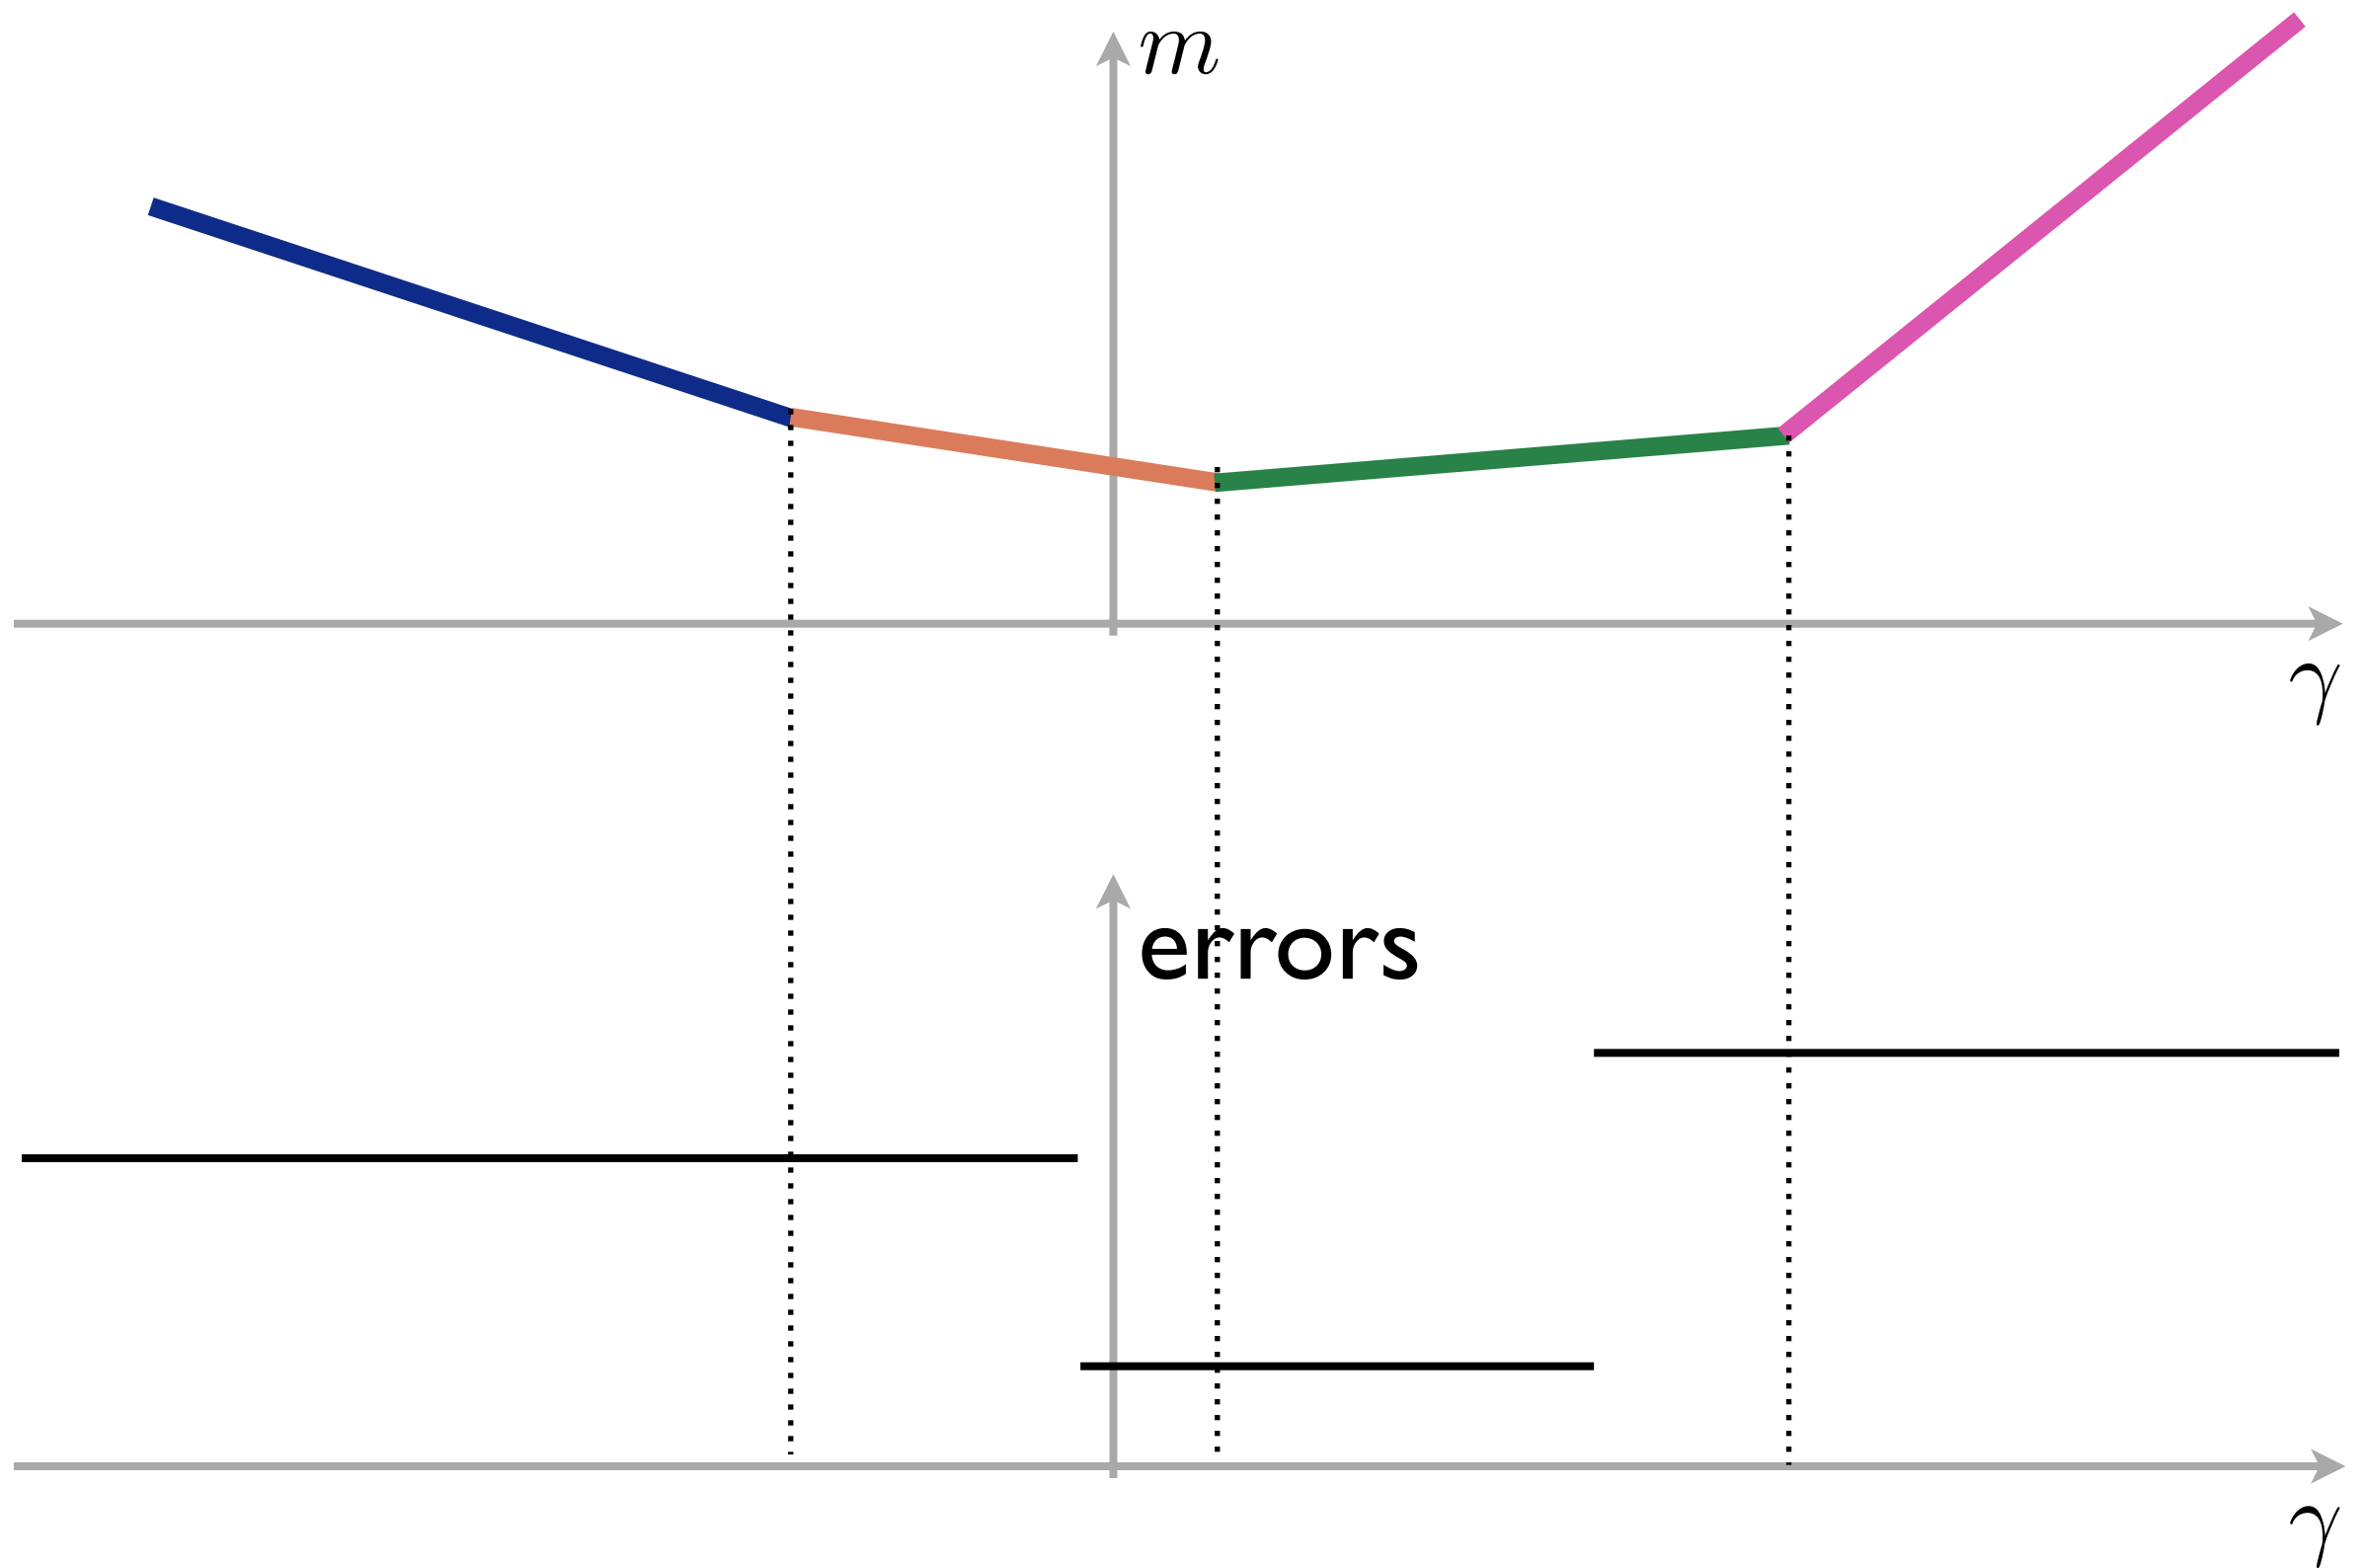
MERT



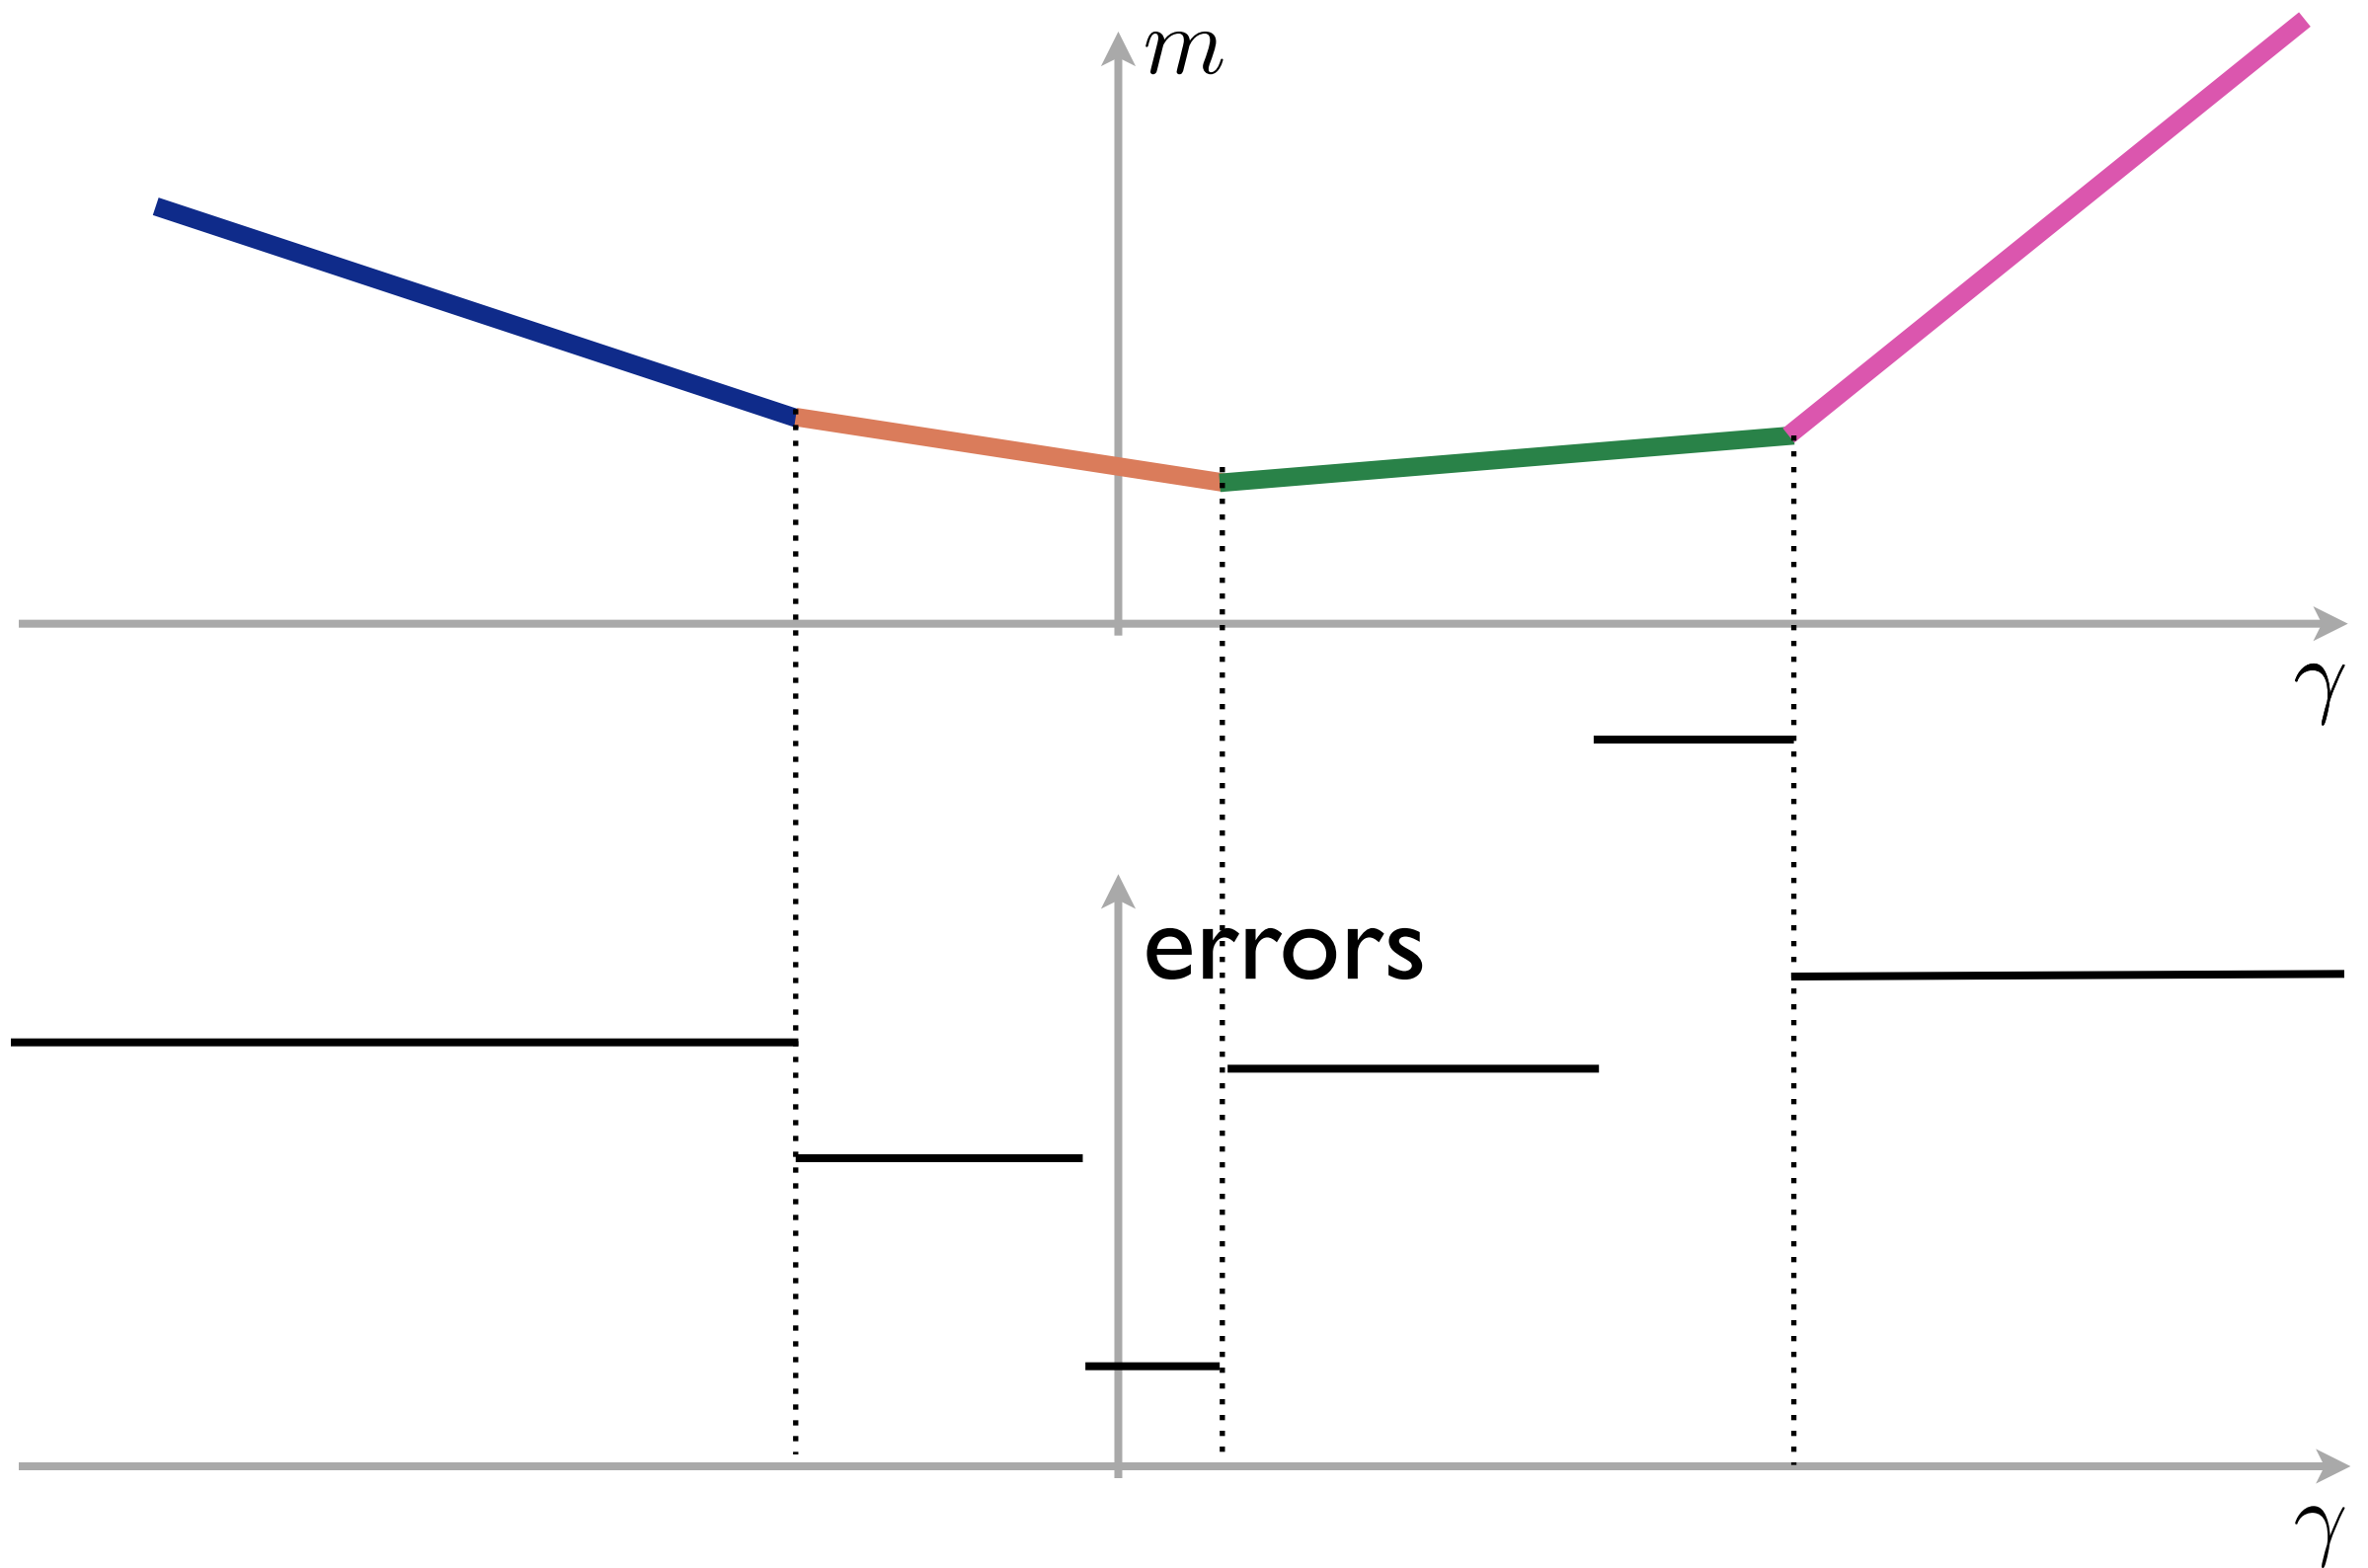
MERT




MERT



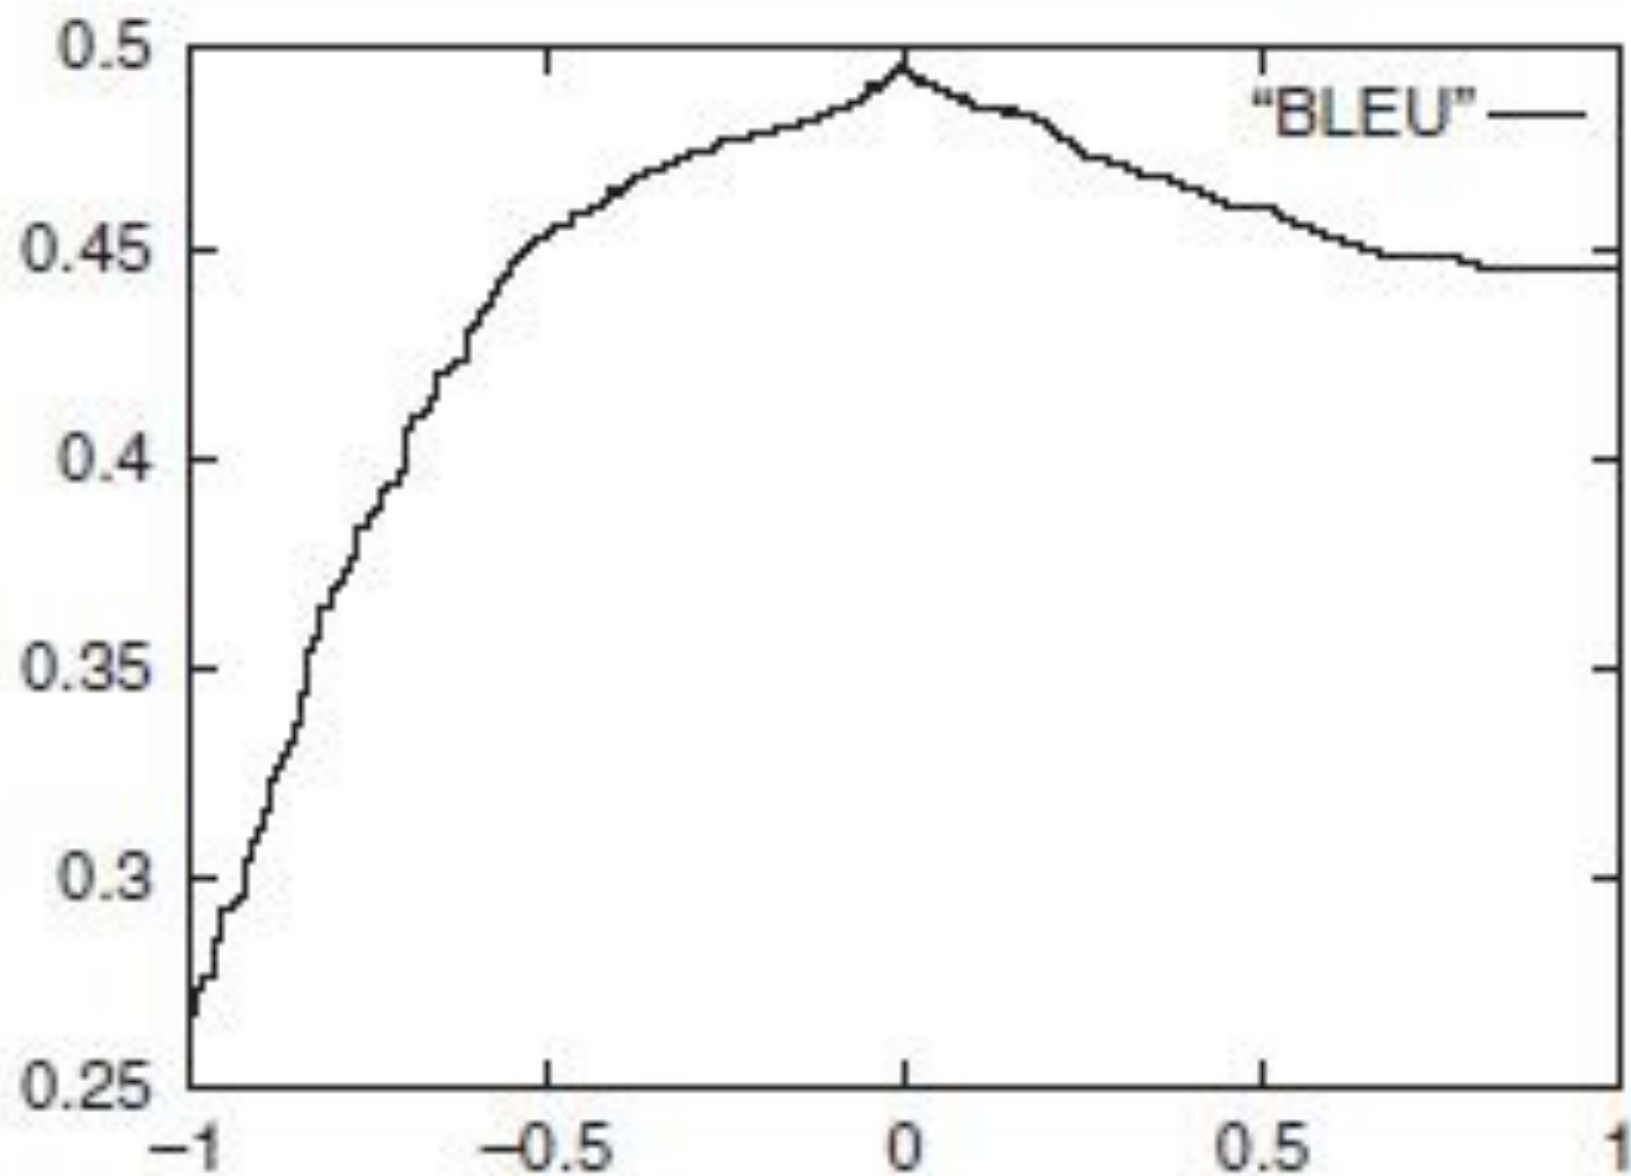
MERT






Let $\mathbf{w}_{\text{new}} = \gamma^* \mathbf{v} + \mathbf{w}$

The effect on BLEU varying one parameter



The effect on BLEU varying one parameter



MERT

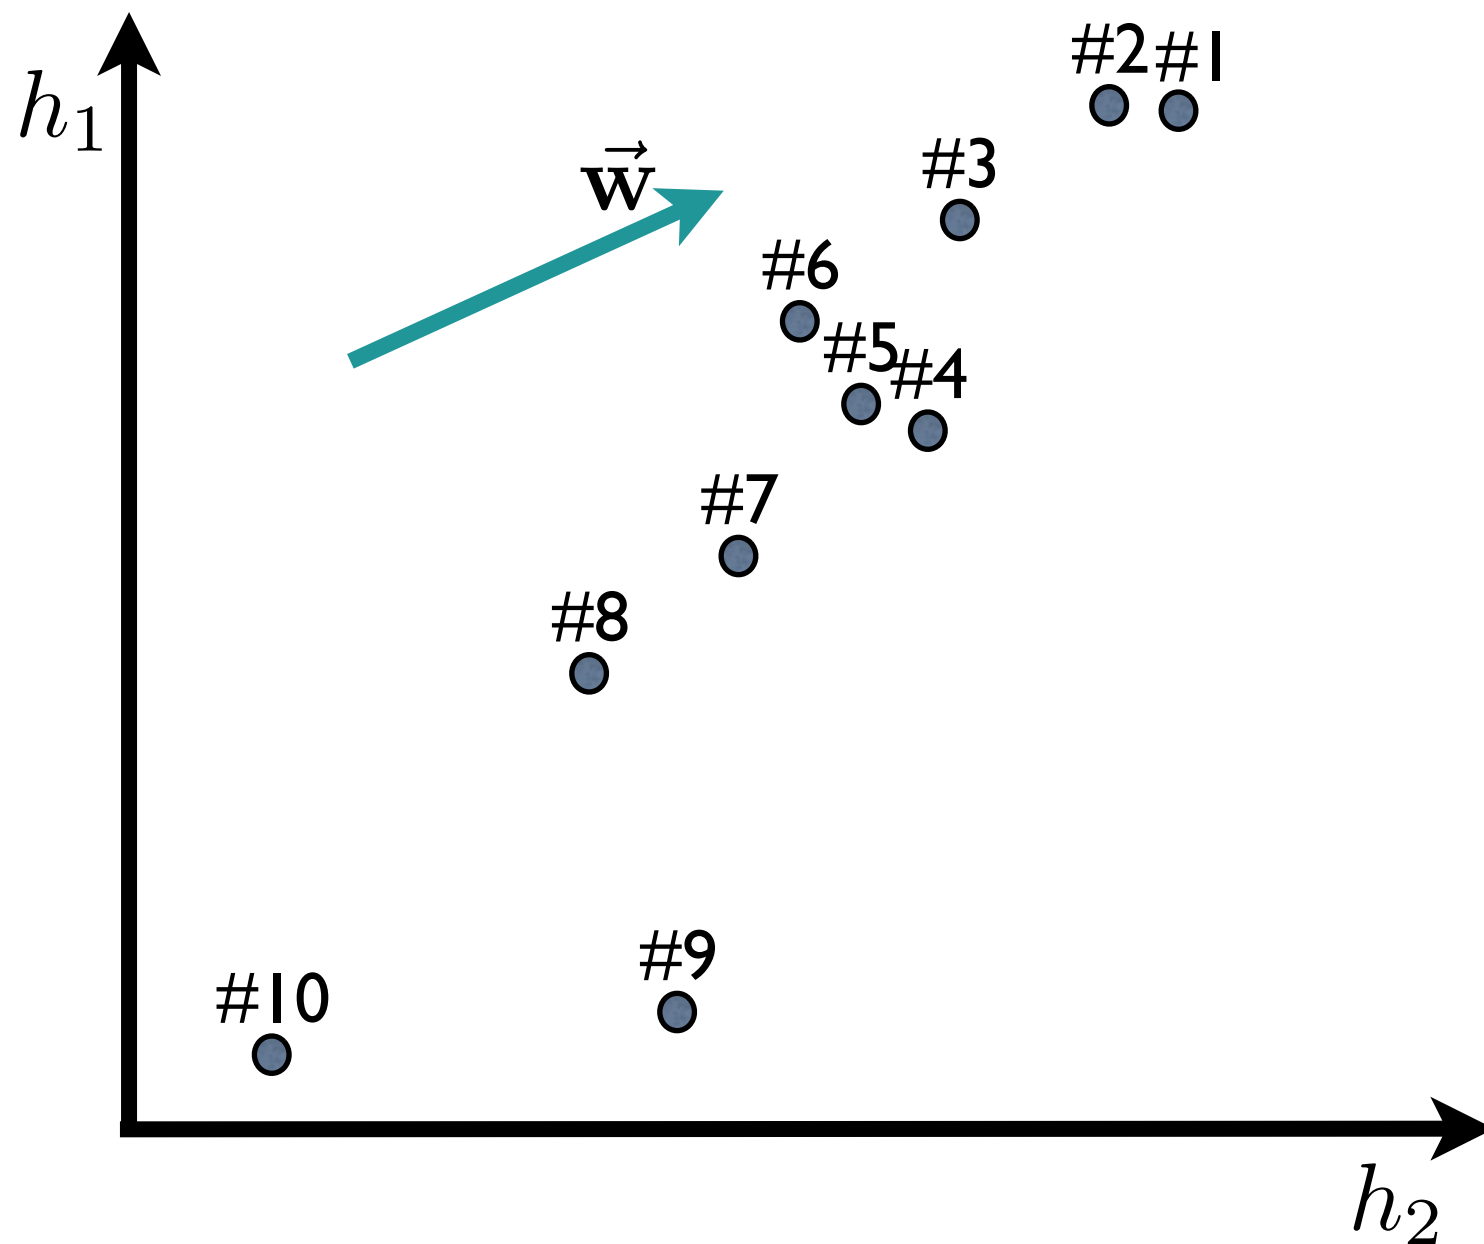
- Minimum error rate training
 - Can maximize or minimize!
- In practice “errors” are sufficient statistics for evaluation metrics (e.g., BLEU, AMBER, TER, etc)
- Downside: MERT can only be used to optimize a small handful of features

Training as Classification





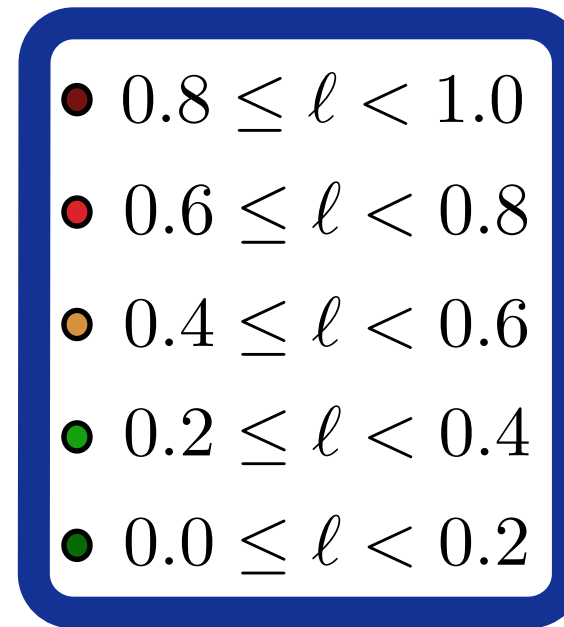
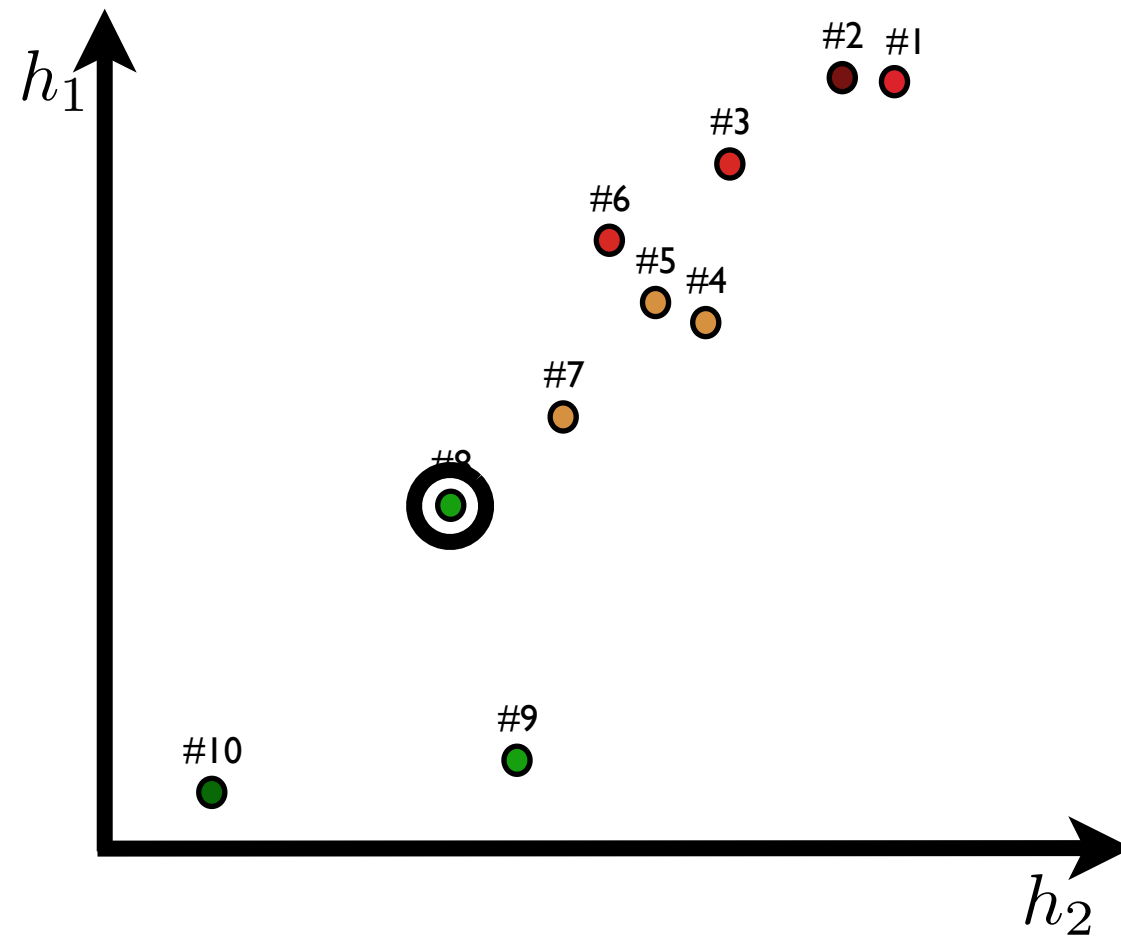
- Pairwise Ranking Optimization
 - Reduce training problem to binary classification with a linear model
- Algorithm
 - For $i=1$ to N
 - Pick random pair of hypotheses (A,B) from K -best list
 - Use cost function to determine if is A or B better
 - Create i th training instance
 - Train binary linear classifier

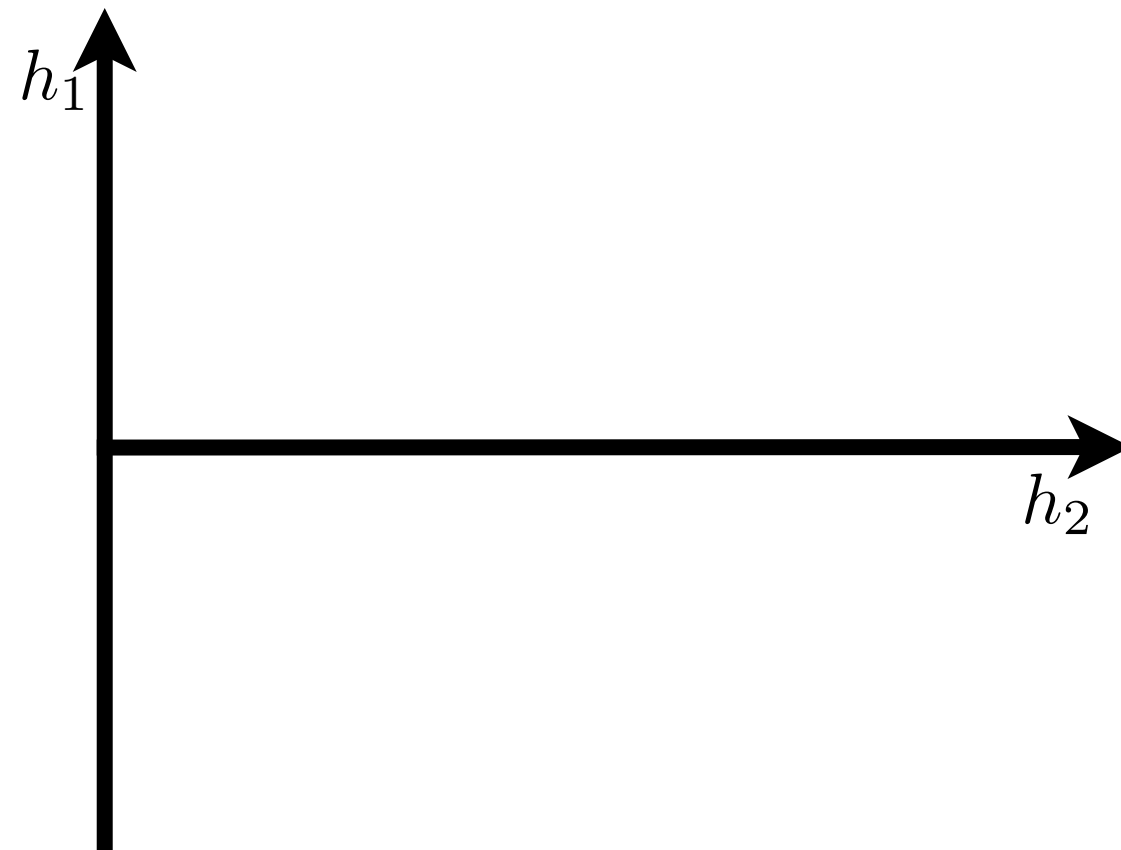
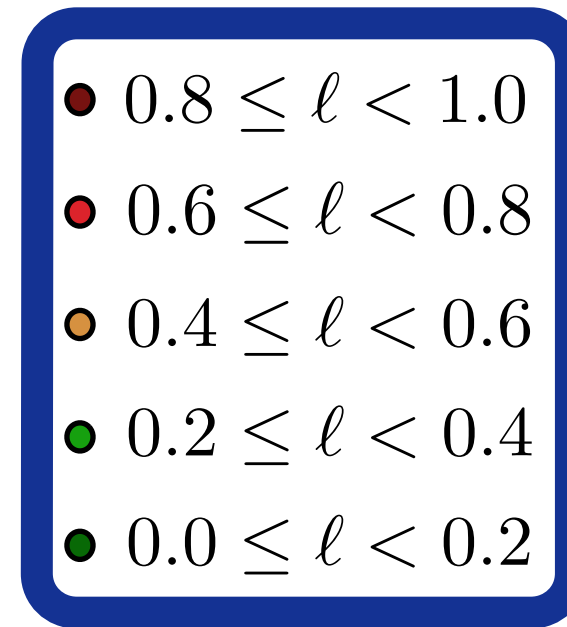
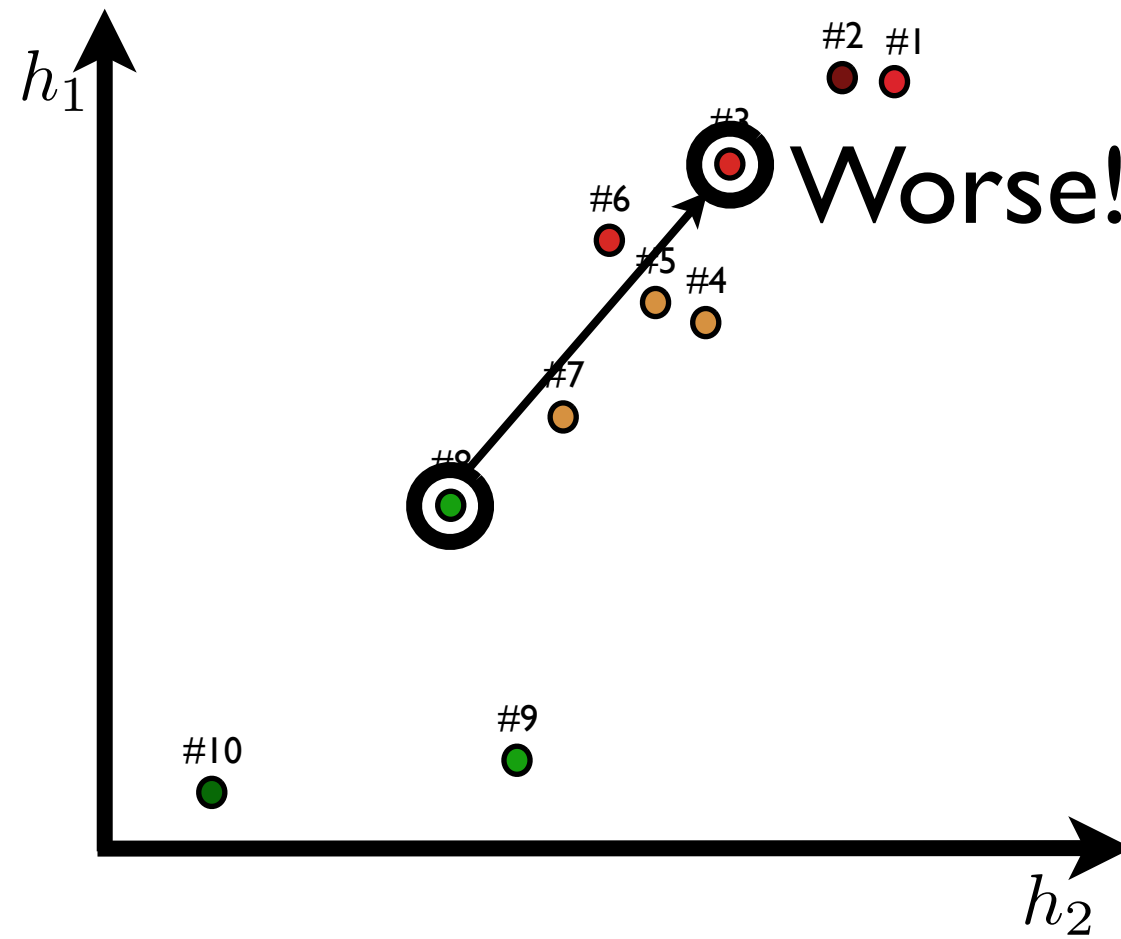
K-Best List Example

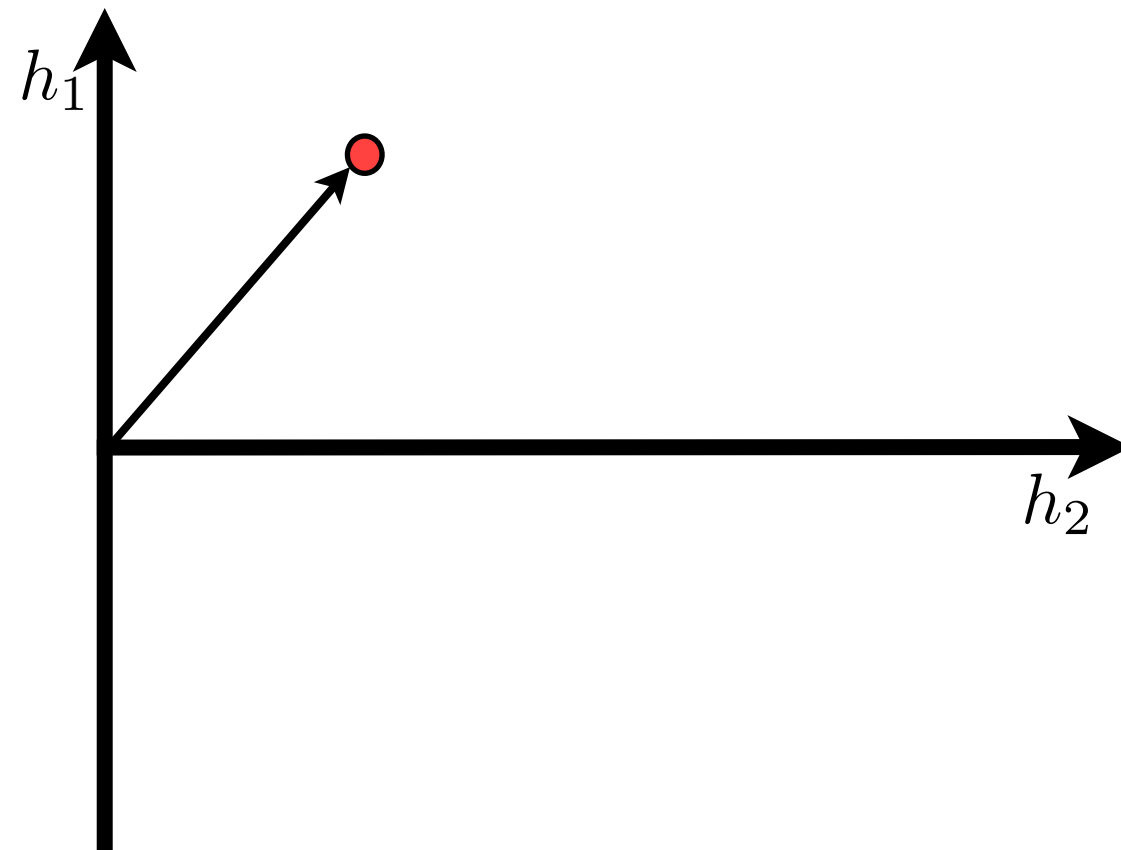
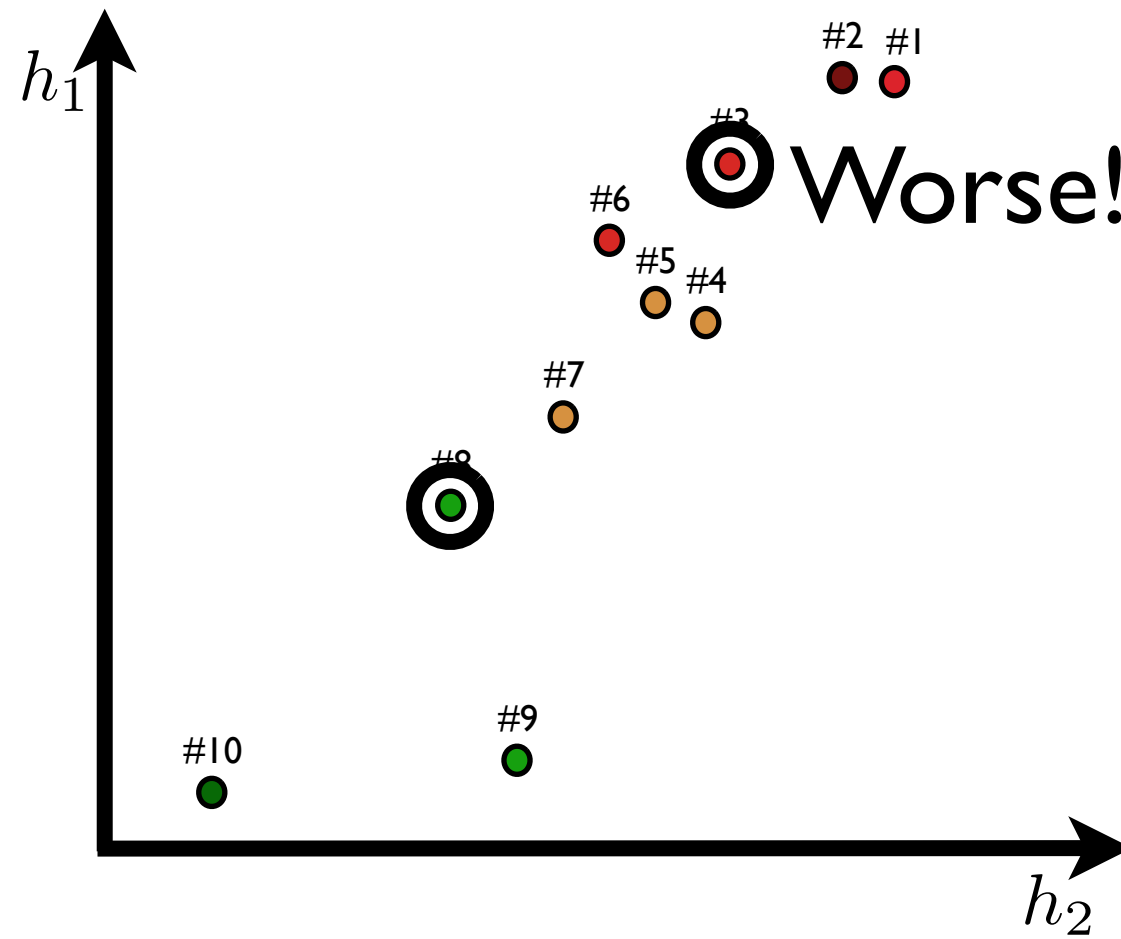





K-Best List Example


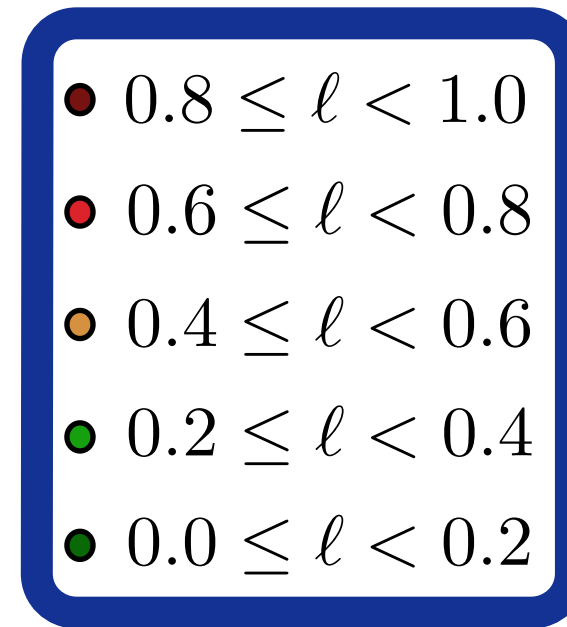



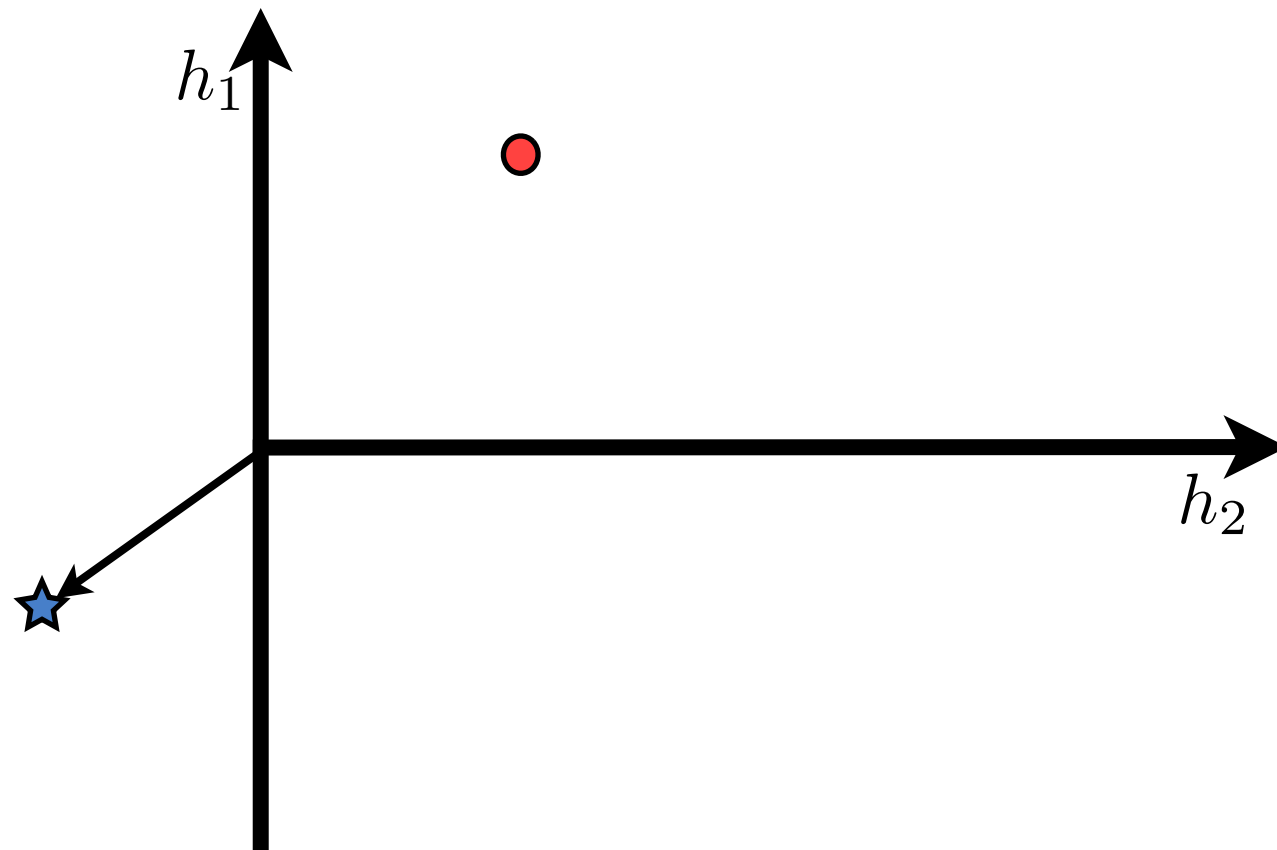
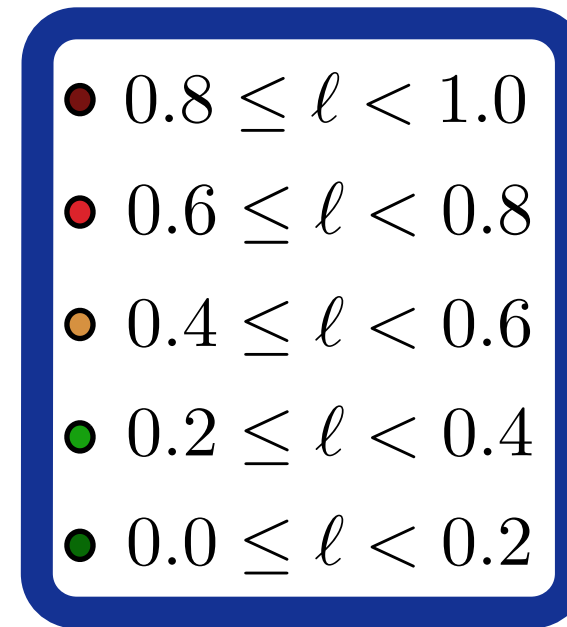
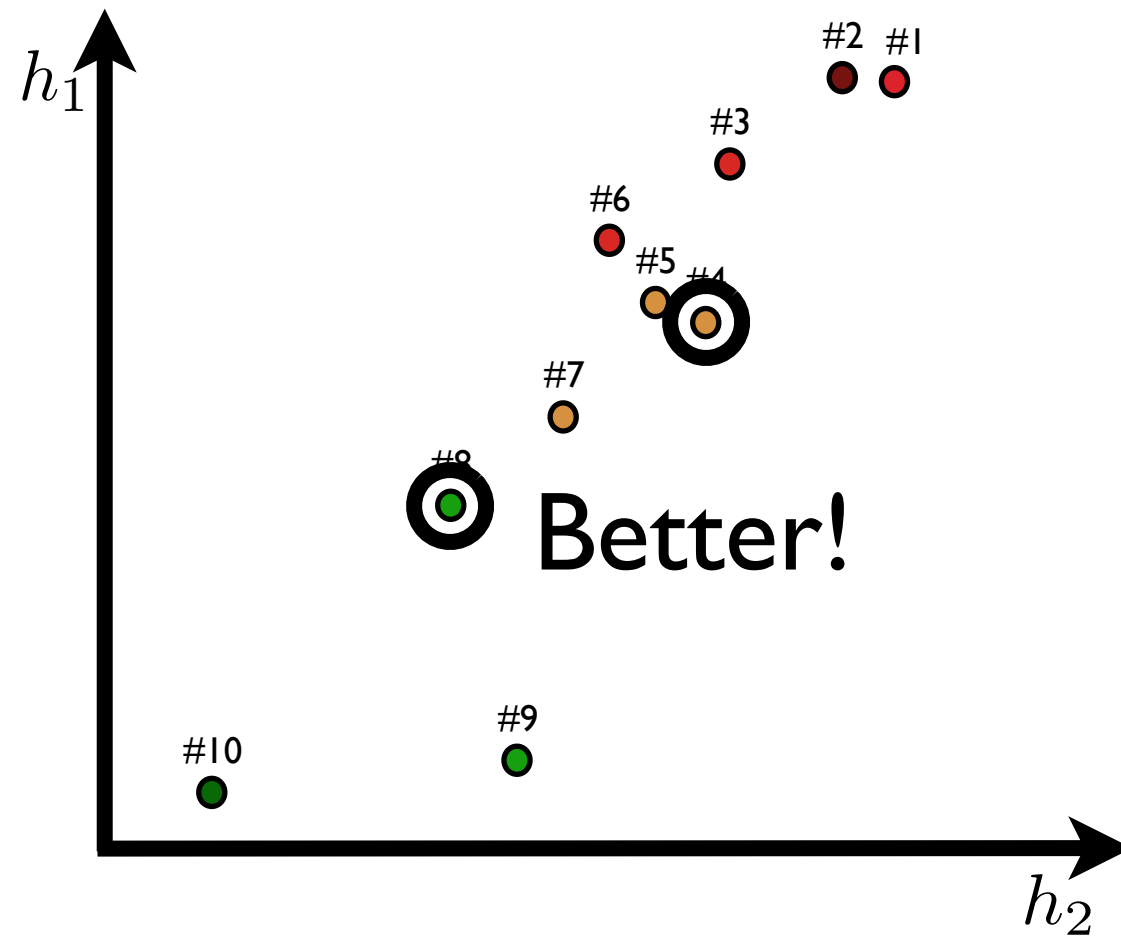


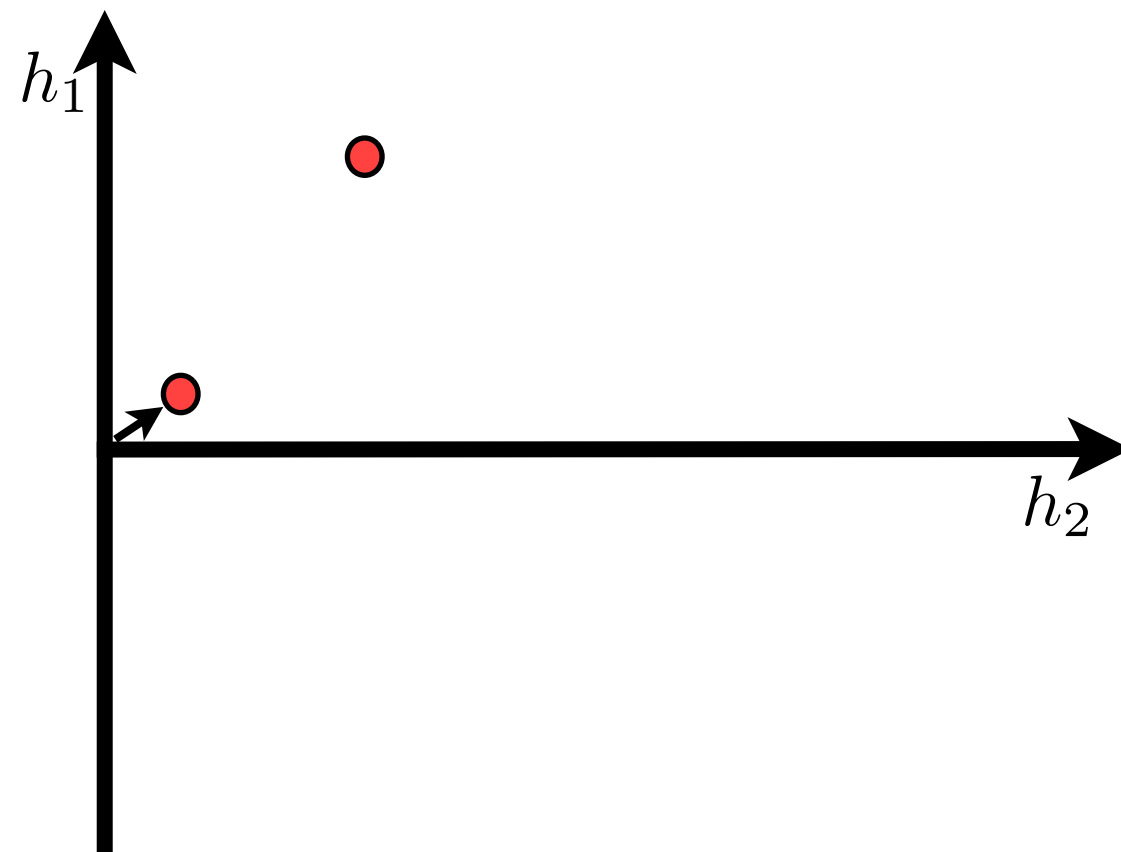
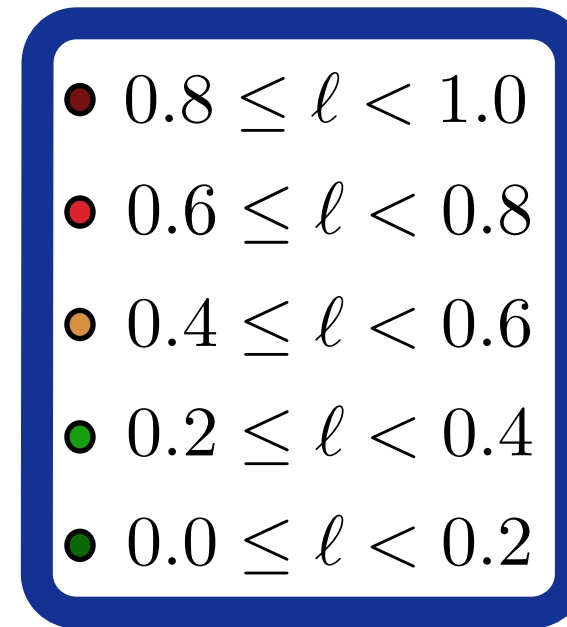




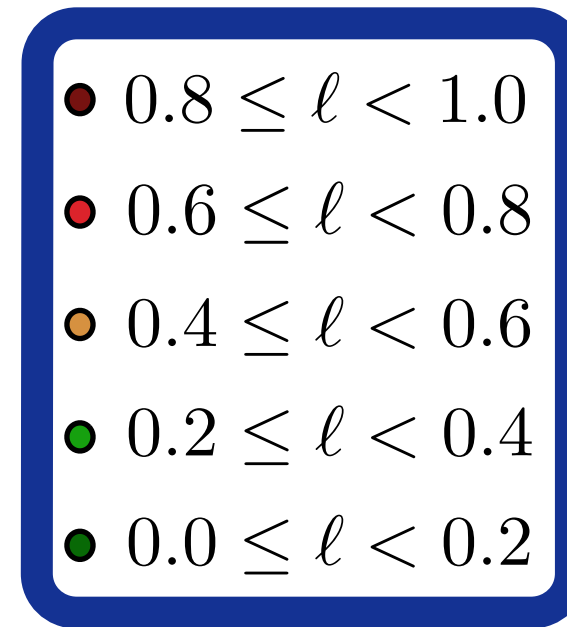
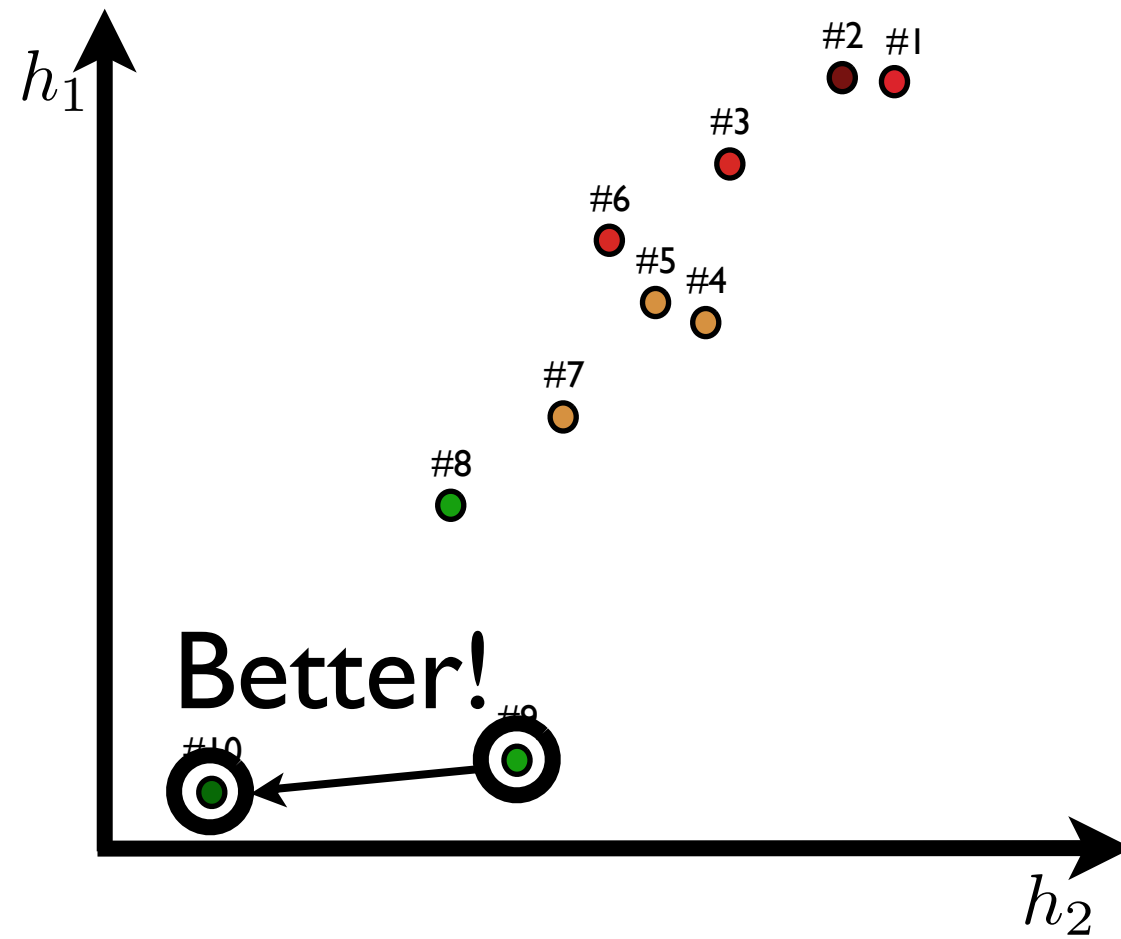



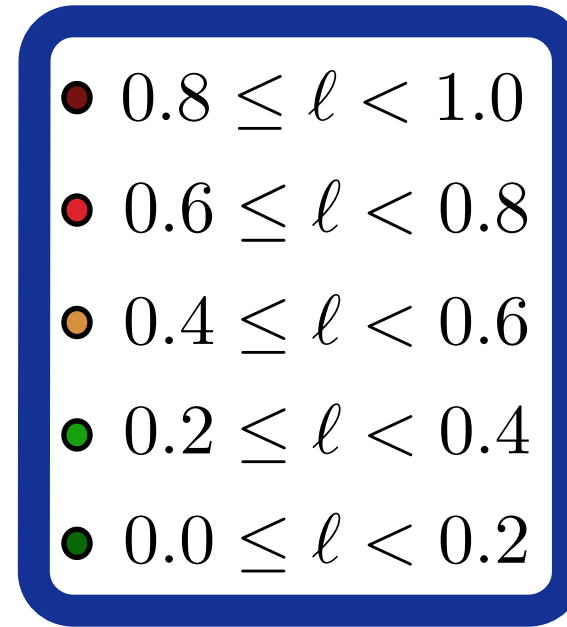
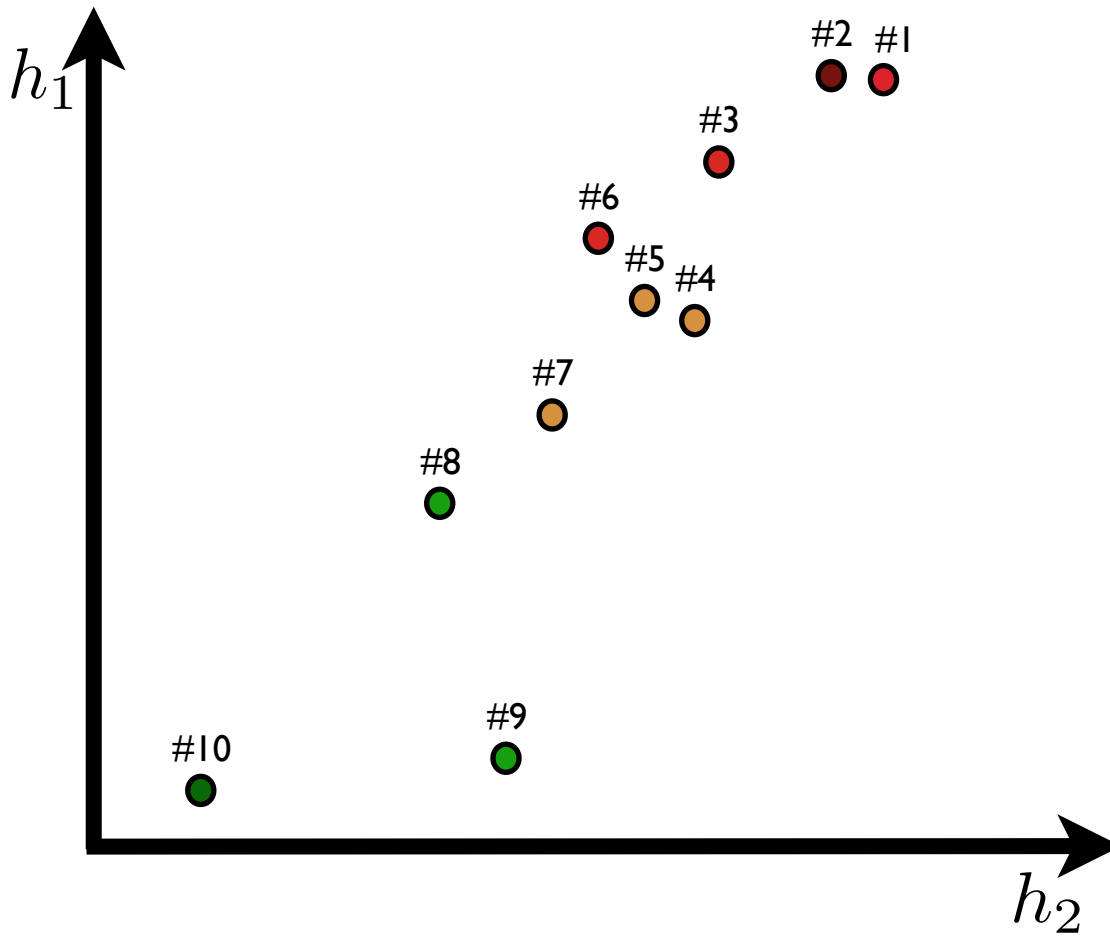







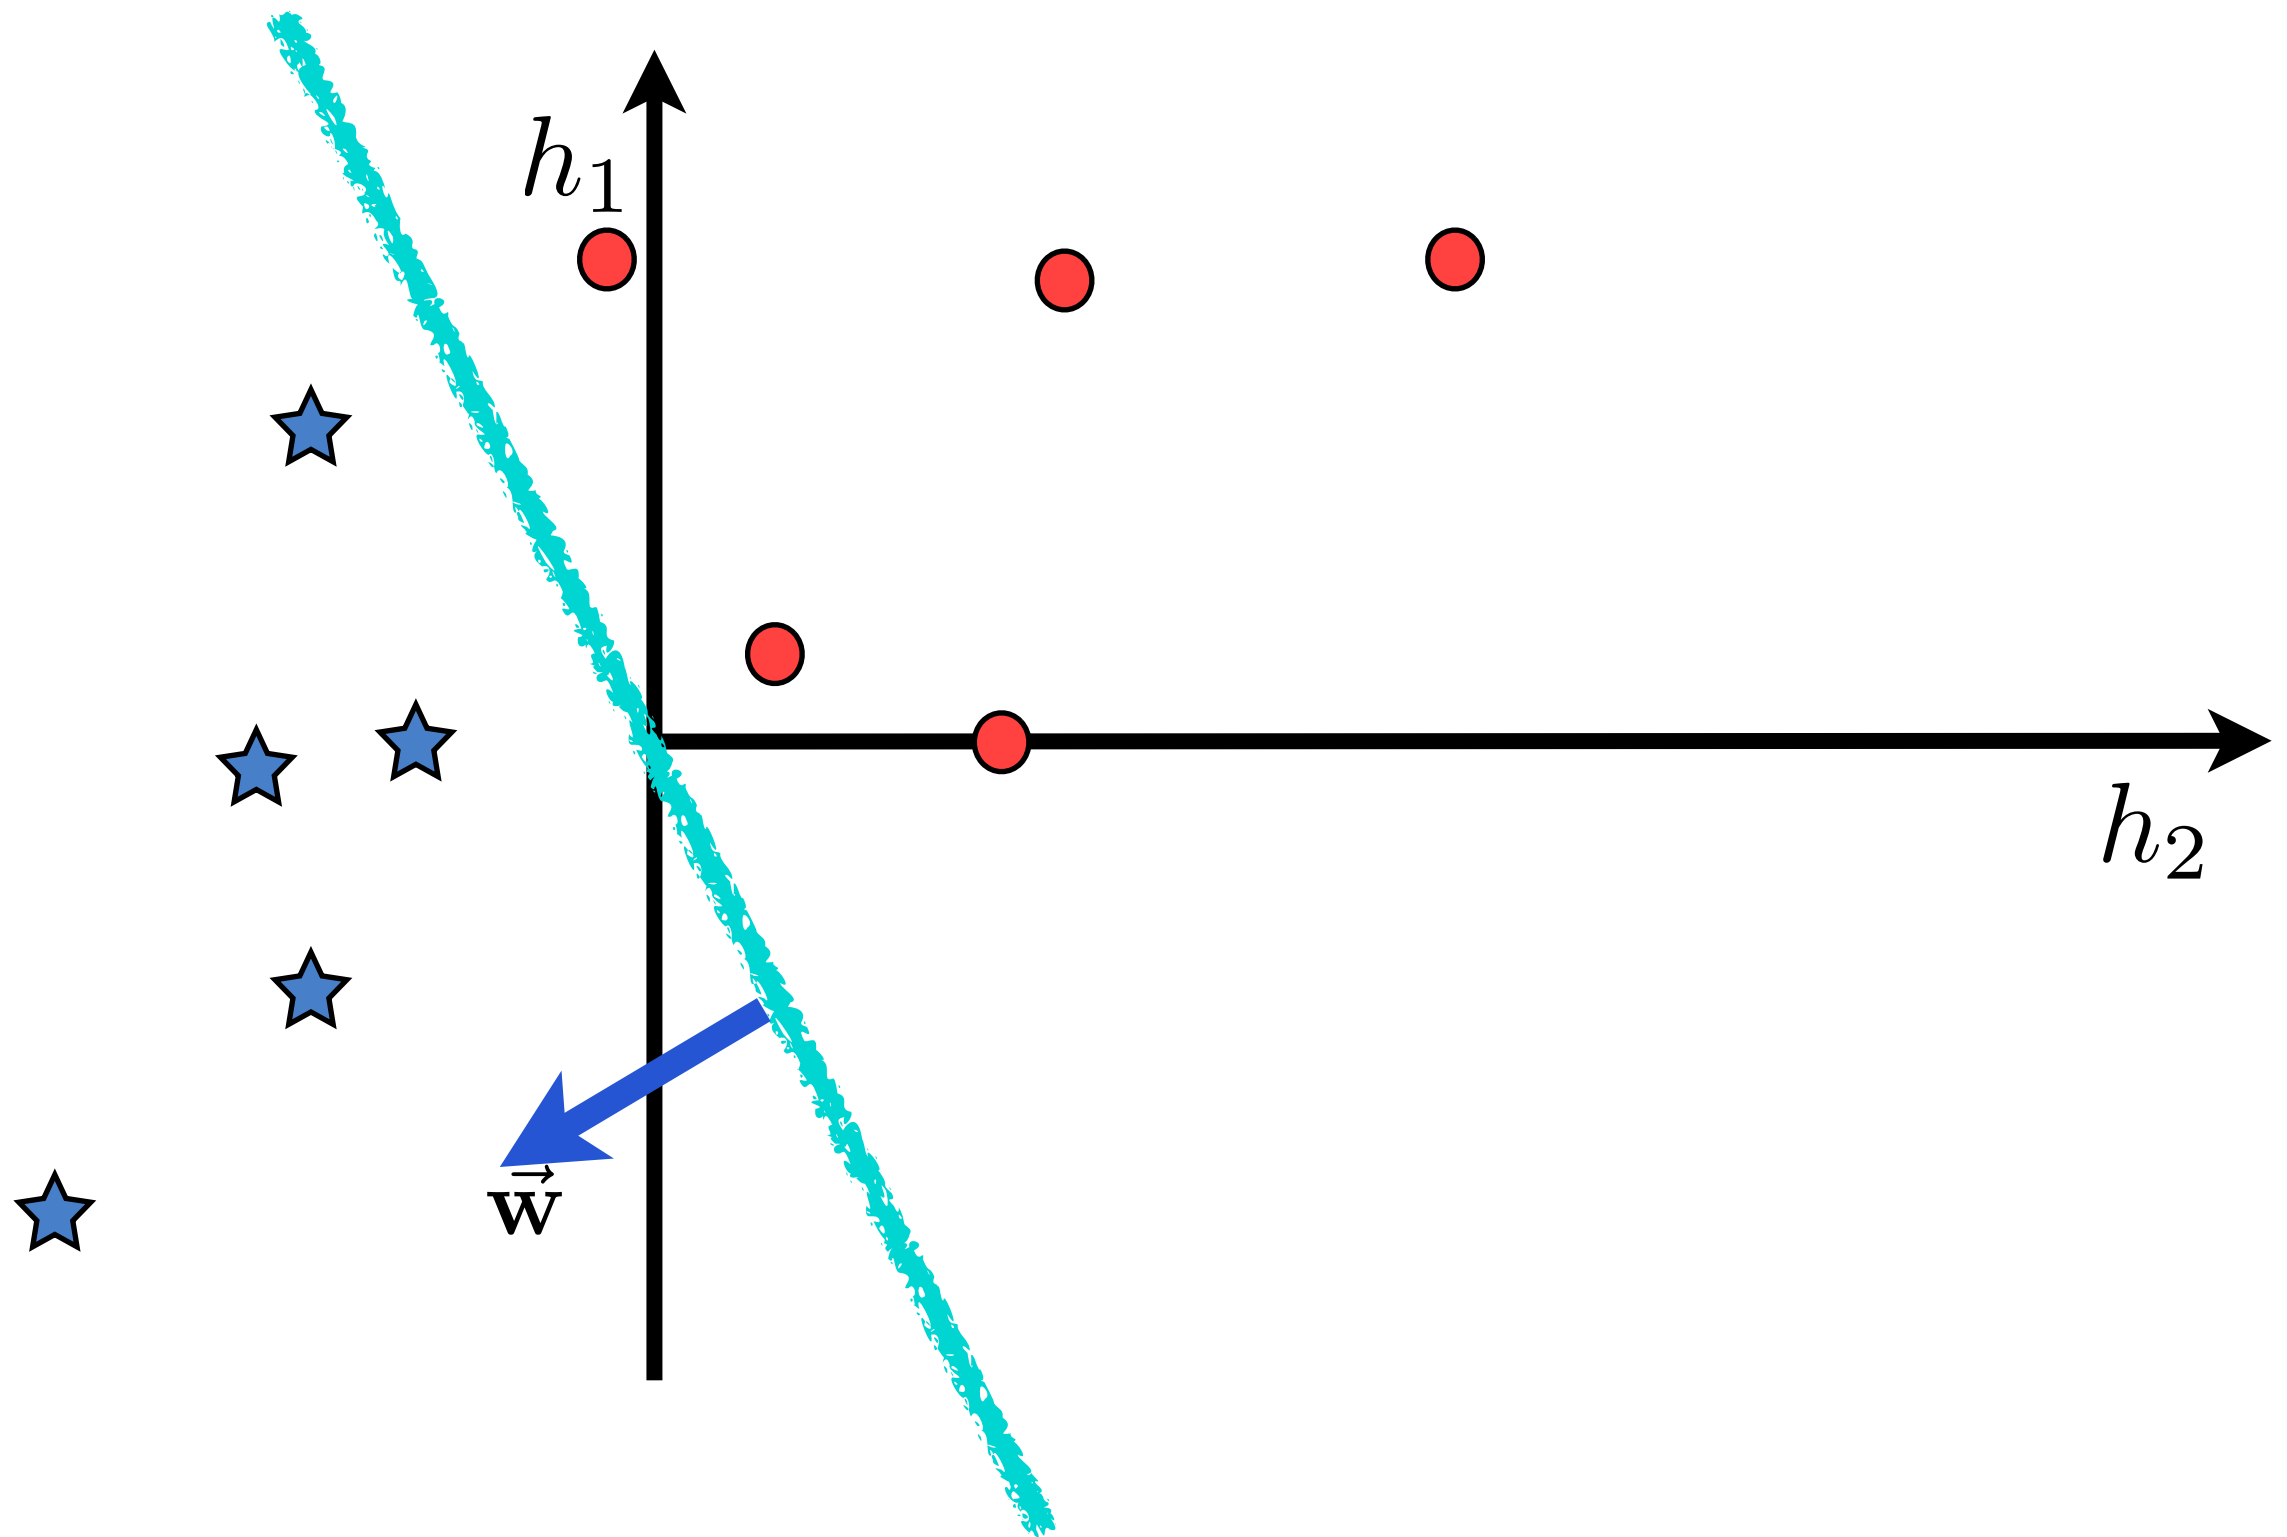









Fit a linear model



Fit a linear model

K-Best List Example

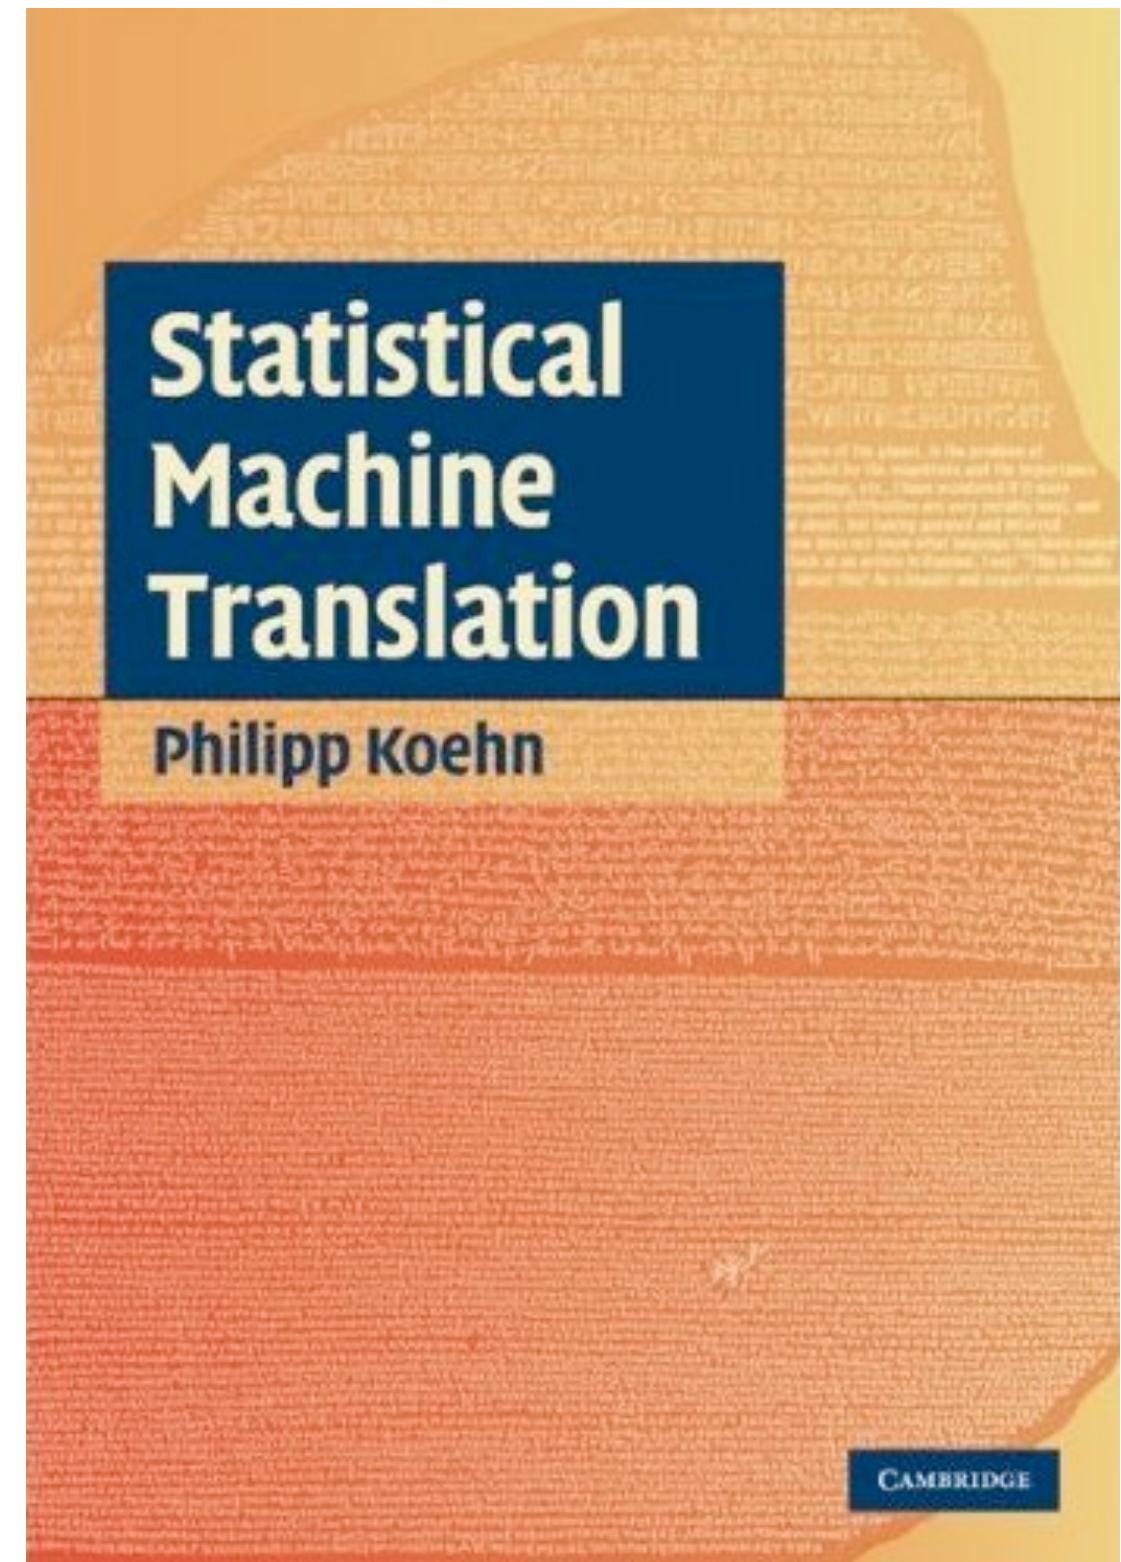


Summary

- Evaluation metrics
 - Figure out how well we're doing
 - Figure out if a feature helps
 - **Train your system**
- What's a great way to improve translation?
 - **Improve evaluation!**

Reading

- Read chapter 9 from the textbook
- HW4 will be a discriminative re-ranking project



Announcements

- HW3 has been released. It is due a week from Thursday.
- Upcoming:
 - Term project (25% of your final grade) and the language research project (10%)
 - These are group projects (2-6 students), where the work scales to the group size
 - Specifications will be posted soon

Term project

- **Problem description** – similar to the descriptions on the HW assignments
- **Data collection** – used to train a model, and evaluate its performance
- **Objective function** – score submissions on a leaderboard
- **Default system** – An implementation of the simplest possible solution
- **Baseline system** – An implementation of a published baseline

Language Research

- Gather monolingual and bilingual data for the language
- Investigate where it is spoken, and what other languages its speakers are exposed to
- Collect information about the syntax and morphology of the language
- Describe its writing system
- Create your own NLP tools for the language (# will vary by team size)